



# a simple way to check returns?

observation: places we return to usually after call instructions

exception: 'tail calls' — we'll ignore this for now

we could check for one...

replace return with:

```
return address ← PopFromStack()  
if DecodeInstruction(return address - 5) == "call thisFunction":  
    goto return address  
else:  
    CRASH
```

# a simple way to check returns?

more practical: `label $ID` instruction with encoding:  
TWO-BYTE-OPCODE FOUR-BYTE-CONSTANT  
(real version: can reuse some sufficiently nop-like instruction)

```
...  
call foo  
label $0xf19279bb // random ID for function foo  
...
```

```
foo:  
...  
pop %rdx // RDX <- return address  
cmp $0xf19279bb, 2(%rdx)  
jne CRASH  
jmp *%rdx
```

# looks like canaries? (1)

what attacks does this stop that canaries don't?

# looks like canaries? (1)

what attacks does this stop that canaries don't?

ID does not need to be secret!

assuming non-executable writeable memory, no!  
attacker can't write new places for return to go

# looks like canaries? (1)

what attacks does this stop that canaries don't?

ID does not need to be secret!

- assuming non-executable writeable memory, no!
- attacker can't write new places for return to go

avoids “stack pivoting” attacks

- attacker can't make stack pointer point to wrong part of stack...
- and expect it to return differently

## looks like canaries? (2)

what attacks does this NOT stop that canaries don't?

example: SortList can be called from Innocent,  
then return from Dangerous

assumption: attacker can overwrite return address at right time  
(running on another core? problem with sortFunc1?)

```
void Innocent() {  
    ...  
    SortList(someList1,  
             sortFunc1);  
    Use(someList1);  
    ...  
}
```

```
void Dangerous() {  
    ...  
    SortList(someList2,  
             sortFunc2);  
    UseDangerously(someList2);  
    ...  
}
```

# checking a VTable call

```
class A { public:  
    virtual void bar() { ... }  
};  
class B : public A { public:  
    void bar() { ... }  
};  
void example(A *obj) {  
    obj->bar();  
}
```

example:

```
// rax <- vtable address  
movq (%rdi), %rax  
// rdx <- first vtable entry  
movq (%rax), %rax  
// call using vtable entry  
call *%rax
```



# checking a VTable call

```
class A { public:  
    virtual void bar() { ... }  
};  
class B : public A { public:  
    void bar() { ... }  
};  
void example(A *obj) {  
    obj->bar();  
}
```

example:

```
// rax <- vtable address  
movq (%rdi), %rax  
// rdx <- first vtable entry  
movq (%rax), %rax  
// call using vtable entry  
call *%rax
```

example uses VTable to call method  
target for memory corruption attacks  
just like return addresses  
so, apply same strategy

# checking a VTable call

```
class A { public:  
    virtual void bar() { ... }  
};  
class B : public A { public:  
    void bar() { ... }  
};  
void example(A *obj) {  
    obj->bar();  
}
```

example:

```
// rax <- vtable address  
movq (%rdi), %rax  
// rdx <- first vtable entry  
movq (%rax), %rax  
// call using vtable entry  
call *%rax
```

```
A::bar():  
    label $0xe0c5df0b  
...  
B::bar():  
    label $0xe0c5df0b  
...
```

example:

```
movq (%rdi), %rax  
movq (%rax), %rax  
cmpq $0xe0c5df0b, 2(%rax)  
jne CRASH  
call *%rax  
...
```

# checking a VTable return

```
A::bar():  
    label $0xe0c5df0b  
    ...  
    pop %rdx // RDX <- return address  
    cmp $0x64a0cfe3, 2(%rdx)  
    jne CRASH  
    jmp *%rdx  
B::bar():  
    label $0xe0c5df0b  
    ...  
    pop %rdx // RDX <- return address  
    cmp $0x64a0cfe3, 2(%rdx)  
    jne CRASH  
    jmp *%rdx
```

```
example:  
    movq (%rdi), %rax  
    movq (%rax), %rax  
    cmpq $0xe0c5df0b, 2(%rax)  
    jne CRASH  
    call *%rax  
    label $0x64a0cfe3  
    ret
```

if we want to use this label-checking on the return

# calls through function pointers

```
typedef int (*CompareFnType)(const char*, const char*)
void SortFunction(const char **items, CompareFnType compare)
{
    ...
    (*compare)(a, b);
    ...
}
```

here: call through explicitly passed function pointer

want to do the same thing we did for VTable calls

- all the compare functions have the same label

- all the returns from compare functions have the same label

yes, if we can somehow label all the compare functions

# calls through function pointers

```
typedef int (*CompareFnType)(const char*, const char*)
void SortFunction(const char **items, CompareFnType compare)
{
    ...
    (*compare)(a, b);
    ...
}
```

here: call through explicitly passed function pointer

want to do the same thing we did for VTable calls

- all the compare functions have the same label

- all the returns from compare functions have the same label

yes, *if we can somehow label all the compare functions*

# CFI overhead

Abadi et al's 2004 paper:

- used label-based approach

- 0-45% time overhead on SPECcpu2000 benchmarks

- best: compression program

- worst: chess engine

Tice et al's 2014 paper (clang-style impl, sometimes in GCC, sometimes in Clang)

- could separately enable different parts

- in tests on SPECcpu 2006 benchmarks:

- 0-10% slowdown for VTable dereference checks

  - but 20% without tuning

- 0-6% for other indirect call checking

## looks like canaries? (2)

what attacks does this NOT stop that canaries don't?

example: SortList can be called from Innocent,  
then return from Dangerous

assumption: attacker can overwrite return address at right time  
(running on another core? problem with sortFunc1?)

```
void Innocent() {  
    ...  
    SortList(someList1,  
             sortFunc1);  
    Use(someList1);  
    ...  
}
```

```
void Dangerous() {  
    ...  
    SortList(someList2,  
             sortFunc2);  
    UseDangerously(someList2);  
    ...  
}
```

# concept: labels and control flow graph




Figure 1: Example program fragment and an outline of its CFG and CFI instrumentation.

## control flow graph

nodes = blocks of code

edges = *potential jump/call*

assigning labels: every in-edge needs to check same label at source




# library-crossing CFGs

main.c

```
#include <png.h>
void ReadImageFromNetwork(
    png_structp libpng_handle,
    unsigned char *bytes,
    size_t size
) { ... }

int main() {
    /* init libpng */
    png_structp libpng_handle = ...;
    /* tell libpng how to read image data */
    png_set_read_fn(
        libpng_handle, ...,
        ReadImageFromNetwork
    )
    ...
    /* extract "header"
       information from image */
    png_get_IHDR(libpng_handle, ...)
    ...
}
```



# CFGs will be imprecise

```
FunctionPtr p = functionA;
Example() {
    while (true) {
        ...
        if (SomethingComplicated()) {
            (*p)();
        } else if (SomethngElseComplicated()) {
            foo();
        }
        ...
    }
}
foo() {
    ...
    if (AnotherComplexThing()) {
        p = functionB;
    }
}
```

# finding possible function pointer values?

given call using function pointers

how do we find the **legitimate** possible values?

one high-level idea:

```
for each fptr constant X:
    PossibleValues[X] = {X}
for each fptr variable X:
    PossibleValues[X] = empty set
until PossibleValues stops changing:
    for each fptr assignment LHS=RHS:
        for each fptr variable/constant Y
            that RHS could evaluate to:
                PossibleValues[LHS] = Union(
                    PossibleValues[LHS],
                    PossibleValues[Y]
                )
```

# finding possible function pointer values?

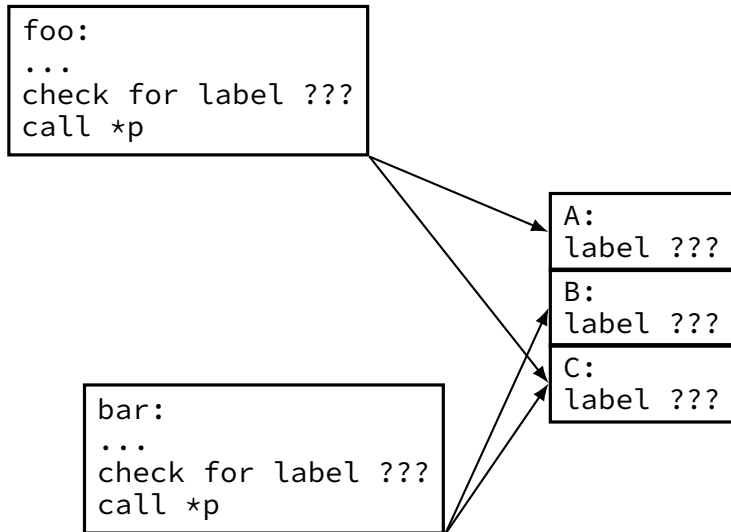
given call using function pointers

how do we find the **legitimate** possible values?

one high-level idea:

```
for each fptr constant X:
    PossibleValues[X] = {X}
for each fptr variable X:
    PossibleValues[X] = empty set
until PossibleValues stops changing:
    for each fptr assignment LHS=RHS:
        for each fptr variable/constant Y
            that RHS could evaluate to:
                PossibleValues[LHS] = Union(
                    PossibleValues[LHS],
                    PossibleValues[Y]
                )
```

# labels aren't enough?



# labels aren't enough?

```
foo:  
...  
check for label ???  
call *p
```

```
bar:  
...  
check for label ???  
call *p
```

A:  
label ???

B:  
label ???

C:  
label ???

two possible fixes:

make checks scan  
for multiple labels  
(more overhead)

allow foo to call B  
and bar to call A  
(easier to attack)

# clang's CFI implementation

<https://clang.llvm.org/docs/ControlFlowIntegrity.html>

also <https://www.usenix.org/conference/usenixsecurity14/technical-sessions/presentation/tice>

only checks calls via VTables or function pointers

stable implementation requires libraries compiled with support

label information is placed in separate data structure

looked up using function or VTable addresses

trick: keep functions in one region of memory

# clang's CFI implementation

<https://clang.llvm.org/docs/ControlFlowIntegrity.html>

also <https://www.usenix.org/conference/usenixsecurity14/technical-sessions/presentation/tice>

only checks calls via VTables or function pointers

stable implementation requires libraries compiled with support

label information is placed in separate data structure

*looked up using function* or VTable addresses

trick: keep functions in one region of memory



# clang idea for CFI indirect calls

```
start_funcs_with_two_string_args:  
    .align 8  
compare_alpha:  
    jmp real_compare_alpha  
    .align 8  
run_command_with_arg:  
    jmp real_run_command_with_arg  
    .align 8  
print_two_strings:  
    jmp real_print_two_strings  
    .align 8  
move_file:  
    jmp real_move_file  
    .align 8  
compare_reverse_alpha:  
    jmp real_compare_reverse_alpha  
end_funcs_with_two_sting_args:
```

functions of same type  
placed together

every func's address  
is multiple of 8

# clang idea for CFI indirect calls

```
start_funcs_with_two_string_args:
    .align 8
compare_alpha:
    jmp real_compare_alpha
    .align 8
run_command_with_arg:
    jmp real_run_command_with_arg
    .align 8
print_two_strings:
    jmp real_print_two_strings
    .align 8
move_file:
    jmp real_move_file
    .align 8
compare_reverse_alpha:
    jmp real_compare_reverse_alpha
end_funcs_with_two_sting_args:
```

check pseudocode:  
round fptr to multiple of 8  
**if** fptr < start or fptr > end:  
    CRASH  
allowed  $\leftarrow$  [1,0,0,0,1]  
    *'mask' for compare funcs*  
offset  $\leftarrow$  fptr - start  
**if** bit (offset/8) of allowed  
    is not set:  
        CRASH

# clang idea for VTables

check VTable element address instead of function address

otherwise

- place all VTables for related classes together

- check start/end address for VTables

- bit mask indicating which VTable entries are okay for call

# CFI prevents?

```
class Foo { public:  
    virtual void f() { }  
};  
class Bar : public Foo { public:  
    virtual void f() { g(1); }  
};  
class Quux : public Foo { public:  
    virtual void f() { }  
};  
void g(int x) {  
    if (x == 0) { danger(); }  
}  
int h(int x) { return 0; }  
int (*ptr)(int) = &h;
```

with clang's CFI, which likely can end up calling `danger()` if an attacker can first write to arbitrary memory locations?

- A. `(*ptr)(1);`
- B. `(*ptr)(0);`
- C. `Foo *q = attacker_controlled(); q->f()`

# CFI prevents?

```
class Foo { public:  
    virtual void f() { }  
};  
class Bar : public Foo { public:  
    virtual void f() { g(1); }  
};  
class Quux : public Foo { public:  
    virtual void f() { }  
};  
void g(int x) {  
    if (x == 0) { danger(); }  
}  
int h(int x) { return 0; }  
int (*ptr)(int) = &h;
```

with clang's CFI, which likely can end up calling `danger()` if an attacker can first write to arbitrary memory locations?

- A. `(*ptr)(1);`
- B. `(*ptr)(0);` *if compiler thinks ptr set to g ever, yes; otherwise, no*
- C. `Foo *q = attacker_controlled(); q->f()` *can only call*