



# executable encoding: security issues

- malware wants to modify executables

  - hide in 'normal' programs (viruses)

  - evade program analysis tools by doing weird things

- memory vulnerabilities require understanding memory layout

  - going from "write memory address X" to "run arbitrary code"

  - where are pointers to functions (we could change?)


  - where is code (vulnerabilities might want to execute)

  - where is data (vulnerabilities might want to read/change)

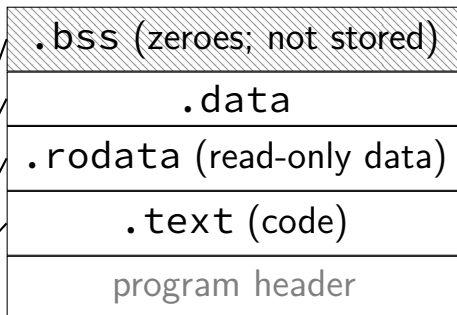
- ways we can load+layout programs that are harder to exploit?

# memory v. disk

(virtual) memory



program on disk



↑  
higher  
addresses  
(and offsets)

# ELF (executable and linking format)

Linux (and some others) executable/object file format

<b>header:</b> machine type, file type, etc.
<b>program header:</b> “ <i>segments</i> ” to load (also, some other information)
<b>segment 1 data</b>
<b>segment 2 data</b>
<b>section header:</b> list of “ <i>sections</i> ”(mostly for linker)

# segments versus sections?

note: ELF terminology; may not be true elsewhere!

sections — *object files* (and usually executables), used by *linker*

- have information on intended purpose

- linkers combine these to create executables

- linkers might omit unneeded sections

segments — executables, used to actually load program

- program loader is *dumb* — doesn't know what segments are for

# section headers

Sections:

Idx	Name	Size	VMA	LMA	File off	Algn
0	.note.ABI-tag	00000020	0000000000400190	0000000000400190	00000190	2**2
	CONTENTS, ALLOC, LOAD, READONLY, DATA					
1	.note.gnu.build-id	00000024	00000000004001b0	00000000004001b0	000001b0	2**2
	CONTENTS, ALLOC, LOAD, READONLY, DATA					
2	.rela.plt	00000210	00000000004001d8	00000000004001d8	000001d8	2**3
	CONTENTS, ALLOC, LOAD, READONLY, DATA					
3	.init	0000001a	00000000004003e8	00000000004003e8	000003e8	2**2
	CONTENTS, ALLOC, LOAD, READONLY, CODE					
4	.plt	00000160	0000000000400410	0000000000400410	00000410	2**4
	CONTENTS, ALLOC, LOAD, READONLY, CODE					
5	.text	0017ff1d	0000000000400570	0000000000400570	00000570	2**4
	CONTENTS, ALLOC, LOAD, READONLY, CODE					
6	__libc_freeres_fn	00002032	0000000000580490	0000000000580490	00180490	2**4
	CONTENTS, ALLOC, LOAD, READONLY, CODE					
7	__libc_thread_freeres_fn	0000021b	00000000005824d0	00000000005824d0	001824d0	2**4
	CONTENTS, ALLOC, LOAD, READONLY, CODE					
8	.fini	00000009	00000000005826ec	00000000005826ec	001826ec	2**2
	CONTENTS, ALLOC, LOAD, READONLY, CODE					
9	.rodata	00044ac8	0000000000582700	0000000000582700	00182700	2**6
	CONTENTS, ALLOC, LOAD, READONLY, DATA					
10	__libc_subfreeres	000000c0	00000000005c71c8	00000000005c71c8	001c71c8	2**3
	CONTENTS, ALLOC, LOAD, READONLY, DATA					
11	.stapsdt.base	00000001	00000000005c7288	00000000005c7288	001c7288	2**0
	CONTENTS, ALLOC, LOAD, READONLY, DATA					
12	__libc_atexit	00000008	00000000005c7290	00000000005c7290	001c7290	2**3
	CONTENTS, ALLOC, LOAD, READONLY, DATA					
13	__libc_thread_subfreeres	00000018	00000000005c7298	00000000005c7298	001c7298	2**3
	CONTENTS, ALLOC, LOAD, READONLY, DATA					
14	.eh_frame	000141dc	00000000005c72b0	00000000005c72b0	001c72b0	2**3
	CONTENTS, ALLOC, LOAD, READONLY, DATA					
15	__aux_objdump_table	0000020b	00000000005c72b0	00000000005c72b0	001c72b0	2**3

# sections

tons of “sections”

not actually needed/used to run program

size, file offset, flags (code/data/etc.)

location in executable *and* in memory

some sections aren't stored (no “CONTENTS” flag)

just all zeroes

## selected sections

.text	program code
.bss	initially zero data (block started by symbol)
.data	other writeable data
.rodata	read-only data
.init/.fini	global constructors/destructors
.got/.plt	dynamic linking related
.eh_frame	try/catch related



# ELF example

`objdump -x /bin/busybox` (on my laptop)

`-x`: output all headers

```
/bin/busybox:      file format elf64-x86-64
/bin/busybox
architecture: i386:x86-64, flags 0x000000102:
EXEC_P, D_PAGED
start address 0x000000000000402170
```

Program Header:  
[...]

Sections:  
[...]

# ELF example

`objdump -x /bin/busybox` (on my laptop)

`-x`: output all headers

```
/bin/busybox:      file format elf64-x86-64
/bin/busybox
architecture: i386:x86-64, flags 0x000000102:
EXEC_P, D_PAGED
start address 0x000000000000402170
```

Program Header:  
[...]

Sections:  
[...]

# ELF example

`objdump -x /bin/busybox` (on my laptop)

`-x`: output all headers

```
/bin/busybox:      file format elf64-x86-64
/bin/busybox
architecture: i386:x86-64, flags 0x000000102:
EXEC_P, D_PAGED
start address 0x00000000000402170
```

Program Header:  
[...]

Sections:  
[...]

# a program header (1)

Program Header:

```
[...]  
LOAD off      0x0001000 vaddr 0x0401000 paddr 0x0401000 align 2**12  
      filesz 0x01b04ed memsz 0x01b04ed flags r-x  
[...]  
LOAD off      0x0207950 vaddr 0x0608950 paddr 0x0608950 align 2**12  
      filesz 0x0008f40 memsz 0x000c718 flags rw-
```

load 0x1bd04ed bytes (read+executable):

from 0x1000 bytes into the file  
to memory at 0x401000

load 0x8f40 bytes (read+write):

from 0x207950 bytes into the file  
to memory at 0x608950  
plus (0xc718-0x8f40) bytes of zeroes

# a program header (1)

Program Header:

```
[...]  
LOAD off      0x0001000 vaddr 0x0401000 paddr 0x0401000 align 2**12  
      filesz 0x01b04ed memsz 0x01b04ed flags r-x  
[...]  
LOAD off      0x0207950 vaddr 0x0608950 paddr 0x0608950 align 2**12  
      filesz 0x0008f40 memsz 0x000c718 flags rw-
```

load 0x1bd04ed bytes (read+executable):

from 0x1000 bytes into the file  
to memory at 0x401000

load 0x8f40 bytes (read+write):

from 0x207950 bytes into the file  
to memory at 0x608950  
plus (0xc718-0x8f40) bytes of zeroes

# a program header (1)

Program Header:

```
[...]  
LOAD off      0x0001000 vaddr 0x0401000 paddr 0x0401000 align 2**12  
      filesz 0x01b04ed memsz 0x01b04ed flags r-x  
[...]  
LOAD off      0x0207950 vaddr 0x0608950 paddr 0x0608950 align 2**12  
      filesz 0x0008f40 memsz 0x000c718 flags rw-
```

load 0x1bd04ed bytes (*read+executable*):

from 0x1000 bytes into the file  
to memory at 0x401000

load 0x8f40 bytes (read+write):

from 0x207950 bytes into the file  
to memory at 0x608950  
plus (0xc718-0x8f40) bytes of zeroes

# a program header (1)

Program Header:

```
[...]  
LOAD off      0x0001000 vaddr 0x0401000 paddr 0x0401000 align 2**12  
      filesz 0x01b04ed memsz 0x01b04ed flags r-x  
[...]  
LOAD off      0x0207950 vaddr 0x0608950 paddr 0x0608950 align 2**12  
      filesz 0x0008f40 memsz 0x000c718 flags rw-
```

load 0x1bd04ed bytes (read+executable):

from 0x1000 bytes into the file  
to memory at 0x401000

load 0x8f40 bytes (read+write):

from 0x207950 bytes into the file  
to memory at 0x608950

*plus (0xc718-0x8f40) bytes of zeroes*

## a program header (2)

Program Header:

[...]

```
    NOTE off      0x0000290 vaddr 0x0400290 paddr 0x0400290 align 2**2
      filesz 0x0000044 memsz 0x0000044 flags r--
    TLS off      0x0207950 vaddr 0x0608950 paddr 0x0608950 align 2**3
      filesz 0x0000030 memsz 0x0000092 flags r--
0x6474e553 off 0x0000270 vaddr 0x0400270 paddr 0x0400270 align 2**3
      filesz 0x0000020 memsz 0x0000020 flags r--
    STACK off    0x0000000 vaddr 0x0000000 paddr 0x0000000 align 2**4
      filesz 0x0000000 memsz 0x0000000 flags rw-
    RELRO off    0x0207950 vaddr 0x0608950 paddr 0x0608950 align 2**0
      filesz 0x00066b0 memsz 0x00066b0 flags r--
```

[...]

NOTE — comment


TLS — thread-local storage region (used via %fs)

0x6474e553 — 'GNU\_PROPERTY' — adtl linker/loader info


STACK — indicates stack is read/write



# ELF LOAD



# ELF LOAD



# dynamic library headers

```
/lib/x86_64-linux-gnu/libc.so.6:      file format elf64-x86-64
/lib/x86_64-linux-gnu/libc.so.6
architecture: i386:x86-64, flags 0x00000150:
HAS_SYMS, DYNAMIC, D_PAGED
start address 0x0000000000000271f0
```

## Program Header:

```
    PHDR off      0x00000000000000040 vaddr 0x00000000000000040 paddr 0x0000000000000000
        filesz 0x00000000000000310 memsz 0x00000000000000310 flags r--
  INTERP off      0x0000000000001c16a0 vaddr 0x0000000000001c16a0 paddr 0x0000000000001c16a0
        filesz 0x0000000000000001c memsz 0x0000000000000001c flags r--
    LOAD off      0x00000000000000000 vaddr 0x00000000000000000 paddr 0x00000000000000000
        filesz 0x00000000000024940 memsz 0x00000000000024940 flags r--
```

...

DYNAMIC — instead of EXEC_P
-----------------------------

# dynamic library headers

```
/lib/x86_64-linux-gnu/libc.so.6:      file format elf64-x86-64
/lib/x86_64-linux-gnu/libc.so.6
architecture: i386:x86-64, flags 0x00000150:
HAS_SYMS, DYNAMIC, D_PAGED
start address 0x0000000000000271f0
```

## Program Header:

```
    PHDR off      0x00000000000000040 vaddr 0x00000000000000040 paddr 0x0000000000000000
        filesz 0x00000000000000310 memsz 0x00000000000000310 flags r--
  INTERP off      0x000000000001c16a0 vaddr 0x000000000001c16a0 paddr 0x000000000001c16a0
        filesz 0x0000000000000001c memsz 0x0000000000000001c flags r--
    LOAD off      0x00000000000000000 vaddr 0x00000000000000000 paddr 0x00000000000000000
        filesz 0x00000000000024940 memsz 0x00000000000024940 flags r--
```

...

specifies loading starting at address 0

but dynamic linker will actually choose a different starting address

# position-independent executables

```
hello.exe:      file format elf64-x86-64
hello.exe
architecture: i386:x86-64, flags 0x00000150:
HAS_SYMS, DYNAMIC, D_PAGED
start address 0x00000000000001080
```

## Program Header:

```
    PHDR off      0x0000000000000040 vaddr 0x0000000000000040 paddr 0x0000000000000000
        filesz 0x00000000000002d8 memsz 0x00000000000002d8 flags r--
  INTERP off      0x0000000000000318 vaddr 0x0000000000000318 paddr 0x0000000000000003
        filesz 0x000000000000001c memsz 0x000000000000001c flags r--
    LOAD off      0x0000000000000000 vaddr 0x0000000000000000 paddr 0x0000000000000000
        filesz 0x00000000000005f8 memsz 0x00000000000005f8 flags r--
```

executable with headers like dynamic library

“position-independent executable”: can be loaded at any address

# position-independent executables

```
hello.exe:      file format elf64-x86-64
hello.exe
architecture: i386:x86-64, flags 0x00000150:
HAS_SYMS, DYNAMIC, D_PAGED
start address 0x00000000000001080
```

## Program Header:

PHDR	off	0x00000000000000040	vaddr	0x00000000000000040	paddr	0x0000000000000000
	filesz	0x000000000000002d8	memsz	0x000000000000002d8	flags	r--
INTERP	off	0x00000000000000318	vaddr	0x00000000000000318	paddr	0x0000000000000003
	filesz	0x0000000000000001c	memsz	0x0000000000000001c	flags	r--
LOAD	off	0x00000000000000000	vaddr	0x00000000000000000	paddr	0x0000000000000000
	filesz	0x000000000000005f8	memsz	0x000000000000005f8	flags	r--

# other executable formats

PE (Portable Executable) — Windows

Mach-O — MacOS X

broadly similar to ELF

differences:

- whether segment/section distinction exists
- how linking/debugging info represented
- how program start info represented

# simple executable startup

copy segments into memory

jump to start address



# executable startup code

Linux: executables don't start at `main`

why not?

- need to initialize `printf`, `cout`, `malloc`, etc. data structures
- `main` needs to return somewhere

compiler links in startup code

# linking

`callq printf`



`callq 0x458F0`

# static v. dynamic linking

static linking — linking *to create executable*

dynamic linking — linking *when executable is run*

# static v. dynamic linking

static linking — linking *to create executable*

dynamic linking — linking *when executable is run*

conceptually: no difference in how they work

reality — very different mechanisms

# linking data structures

symbol table: name  $\Rightarrow$  (section, offset)

example: `main:` in assembly adds symbol table entry for `main`

relocation table: offset  $\Rightarrow$  (name, kind)

example: `call printf` adds relocation for name `printf`  
kind depends on how instruction encodes address

# hello.s

```
.data
string: .asciz "Hello , World!"
.text
.globl main
main:
    movq $string, %rdi
    call puts
    ret
```

# hello.o (pre-static or dynamic linking)

## SYMBOL TABLE:

000000000000000000	l	d	.text	000000000000000000	.text
000000000000000000	l	d	.data	000000000000000000	.data
000000000000000000	l	d	.bss	000000000000000000	.bss
000000000000000000	l		.data	000000000000000000	string
000000000000000000	g		.text	000000000000000000	main
000000000000000000			*UND*	000000000000000000	puts

## RELOCATION RECORDS FOR [.text]:

OFFSET	TYPE	VALUE
000000000000000003	R_X86_64_32S	.data
000000000000000008	R_X86_64_PC32	puts-0x0000000000000004

# hello.o (pre-static or dynamic linking)

## SYMBOL TABLE:

000000000000000000	l	d	.text	000000000000000000	.text
000000000000000000	l	d	.data	000000000000000000	.data
000000000000000000	l	d	.bss	000000000000000000	.bss
000000000000000000	l		.data	000000000000000000	string
000000000000000000	g		.text	000000000000000000	main
000000000000000000			<i>*UND*</i>	<i>000000000000000000</i>	<i>puts</i>

## RELOCATION RECORDS FOR [.text]:

OFFSET	TYPE	VALUE
000000000000000003	R_X86_64_32S	.data
000000000000000008	R_X86_64_PC32	puts-0x0000000000000004

undefined symbol: look for puts elsewhere



# hello.o (pre-static or dynamic linking)

## SYMBOL TABLE:

000000000000000000	l	d	.text	000000000000000000	.text
000000000000000000	l	d	.data	000000000000000000	.data
000000000000000000	l	d	.bss	000000000000000000	.bss
000000000000000000	l		.data	000000000000000000	string
000000000000000000	g		.text	000000000000000000	main
000000000000000000			*UND*	000000000000000000	puts

## RELOCATION RECORDS FOR [.text]:

OFFSET	TYPE	VALUE
000000000000000003	R_X86_64_32S	.data
000000000000000008	R_X86_64_PC32	<i>puts</i> -0x00000000000000004

insert address of puts, format for call

# hello.o (pre-static or dynamic linking)

## SYMBOL TABLE:

000000000000000000	l	d	.text	000000000000000000	.text
000000000000000000	l	d	.data	000000000000000000	.data
000000000000000000	l	d	.bss	000000000000000000	.bss
000000000000000000	l		.data	000000000000000000	string
000000000000000000	g		.text	000000000000000000	main
000000000000000000			*UND*	000000000000000000	puts

## RELOCATION RECORDS FOR [.text]:

OFFSET	TYPE	VALUE
000000000000000003	R_X86_64_32S	<i>.data</i>
000000000000000008	R_X86_64_PC32	puts-0x0000000000000004

insert address of string, format for movq

# hello.o (pre-static or dynamic linking)

## SYMBOL TABLE:

000000000000000000	l	d	.text	000000000000000000	.text
000000000000000000	l	d	.data	000000000000000000	.data
000000000000000000	l	d	.bss	000000000000000000	.bss
000000000000000000	l		.data	000000000000000000	string
000000000000000000	g		.text	000000000000000000	main
000000000000000000			*UND*	000000000000000000	puts

## RELOCATION RECORDS FOR [.text]:

OFFSET	TYPE	VALUE
000000000000000003	<i>R_X86_64_32S</i>	.data
000000000000000008	<i>R_X86_64_PC32</i>	puts-0x0000000000000004

different ways to represent address

- 32S — signed 32-bit value
- PC32 — 32-bit difference from current address

# hello.o (pre-static or dynamic linking)

## SYMBOL TABLE:

000000000000000000	l	d	.text	000000000000000000	.text
000000000000000000	l	d	.data	000000000000000000	.data
000000000000000000	l	d	.bss	000000000000000000	.bss
000000000000000000	l		.data	000000000000000000	string
000000000000000000	<i>g</i>		.text	000000000000000000	main
000000000000000000			*UND*	000000000000000000	puts

## RELOCATION RECORDS FOR [.text]:

OFFSET	TYPE	VALUE
000000000000000003	R_X86_64_32S	.data
000000000000000008	R_X86_64_PC32	puts-0x00000000000000004

g: global — used by other files  
l: local

# hello.o (pre-static or dynamic linking)

## SYMBOL TABLE:

000000000000000000	l	d	.text	000000000000000000	.text
000000000000000000	l	d	.data	000000000000000000	.data
000000000000000000	l	d	.bss	000000000000000000	.bss
000000000000000000	l		.data	000000000000000000	string
000000000000000000	g		<i>.text</i>	<i>000000000000000000</i>	main
000000000000000000			*UND*	000000000000000000	puts

## RELOCATION RECORDS FOR [.text]:

OFFSET	TYPE	VALUE
000000000000000003	R_X86_64_32S	.data
000000000000000008	R_X86_64_PC32	puts-0x00000000000000004

.text segment beginning plus 0 bytes

# hello.o / statically linked, no PIE

hello.o:

Disassembly of section .text:

0000000000000000 <main>:

```
0:  48 c7 c7 00 00 00 00    mov     $0x0,%rdi
                               3: R_X86_64_32S .data
7:  e8 00 00 00 00          call    c <main+0xc>
                               8: R_X86_64_PLT32 puts-0x4
c:  c3                      ret
```

---

hello.exe (gcc -static -no-pie):

Disassembly of section .text:

...

00000000004016c3 <main>:

```
4016c3:►  48 c7 c7 b5 16 40 00 ► mov     $0x4016b5,%rdi
4016ca:►  e8 e1 a9 00 00      ► call   40c0b0 <_IO_puts>
4016cf:►  c3                  ► ret
```

...

# symbols in executable

SYMBOL TABLE:

```
00000000000000000000 l      df *ABS*  00000000000000000000 crt1.o
0000000000004002b4 l      0 .note.ABI-tag  000000000000000020 __ab
00000000000000000000 l      df *ABS*  00000000000000000000 assert.o
000000000000498530 l      0 .rodata          000000000000000013 errs
000000000000401100 l      F .text  00000000000000000f __assert_fai
```

...

by default, symbol information in statically-linked executable

...but not actually used to run it!

can be stripped (-s linker option or strip command) instead:

SYMBOL TABLE:

no symbols

## exercise: finding without symbols?

how can I find where functions are without symbols?

ideally: some *automated* way to do this?



## interlude: strace

strace — system call tracer

on Linux, some other Unices

OS X approx. equivalent: dtruss

Windows approx. equivalent: Process Monitor

indicates what system calls (operating system services) used by a program

# statically linked hello.exe

```
gcc -no-pie -static -o hello-static.exe hello.s
```

```
strace ./hello-static.exe:
```

```
execve("./hello-static.exe", [ "./hello-static.exe" ], [ /* 46 vars */ ]) = 0
uname(sysname="Linux", nodename="reiss-lenovo", ...) = 0
brk(NULL) = 0x20a5000
brk(0x20a61c0) = 0x20a61c0
arch_prctl(ARCH_SET_FS, 0x20a5880) = 0
readlink("/proc/self/exe", "/home/cr4bd/spring2017/cs4630/sl"... , 4096) = 62
brk(0x20c71c0) = 0x20c71c0
brk(0x20c8000) = 0x20c8000
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
fstat(1, st_mode=S_IFCHR|0620, st_rdev=makedev(136, 1), ...) = 0
write(1, "Hello, World!\n", 14) = 14
exit_group(14) = ?
+++ exited with 14 +++
```

# statically linked hello.exe

```
gcc -no-pie -static -o hello-static.exe hello.s
```

```
strace ./hello-static.exe:
```

```
execve("./hello-static.exe", ["./hello-static.exe"], [/* 46 vars */]) = 0
uname(sysname="Linux", nodename="reiss-lenovo", ...) = 0
brk(NULL) = 0x20a5000
brk(0x20a61c0) = 0x20a61c0
arch_prctl(ARCH_SET_FS, 0x20a5880) = 0
readlink("/proc/self/exe", "/home/cr4bd/spring2017/cs4630/sl"..., 4096) = 62
brk(0x20c71c0) = 0x20c71c0
brk(0x20c8000) = 0x20c8000
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
fstat(1, st_mode=S_IFCHR|0620, st_rdev=makedev(136, 1), ...) = 0
write(1, "Hello, World!\n", 14) = 14
exit_group(14) = ?
+++ exited with 14 +++
```

standard library startup

# statically linked hello.exe

```
gcc -no-pie -static -o hello-static.exe hello.s
```

```
strace ./hello-static.exe:
```

```
execve("./hello-static.exe", ["./hello-static.exe"], [/* 46 vars */]) = 0
uname(sysname="Linux", nodename="reiss-lenovo", ...) = 0
brk(NULL)                                = 0x20a5000
brk(0x20a61c0)                           = 0x20a61c0
arch_prctl(ARCH_SET_FS, 0x20a5880)       = 0
readlink("/proc/self/exe", "/home/cr4bd/spring2017/cs4630/sl"... , 4096) = 62
brk(0x20c71c0)                           = 0x20c71c0
brk(0x20c8000)                           = 0x20c8000
access("/etc/ld.so.nohwcap", F_OK)       = -1 ENOENT (No such file or directory)
fstat(1, st_mode=S_IFCHR|0620, st_rdev=makedev(136, 1), ...) = 0
write(1, "Hello, World!\n", 14)          = 14
exit_group(14)                           = ?
+++ exited with 14 +++
```

memory allocation

# statically linked hello.exe

```
gcc -no-pie -static -o hello-static.exe hello.s
```

```
strace ./hello-static.exe:
```

```
execve("./hello-static.exe", ["./hello-static.exe"], [/* 46 vars */]) = 0
uname(sysname="Linux", nodename="reiss-lenovo", ...) = 0
brk(NULL) = 0x20a5000
brk(0x20a61c0) = 0x20a61c0
arch_prctl(ARCH_SET_FS, 0x20a5880) = 0
readlink("/proc/self/exe", "/home/cr4bd/spring2017/cs4630/sl"... , 4096) = 62
brk(0x20c71c0) = 0x20c71c0
brk(0x20c8000) = 0x20c8000
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
fstat(1, st_mode=S_IFCHR|0620, st_rdev=makedev(136, 1), ...) = 0
write(1, "Hello, World!\n", 14) = 14
exit_group(14) = ?
+++ exited with 14 +++
```

implementation of puts

# statically linked hello.exe

```
gcc -no-pie -static -o hello-static.exe hello.s
```

```
strace ./hello-static.exe:
```

```
execve("./hello-static.exe", ["/hello-static.exe"], [/* 46 vars */]) = 0
uname(sysname="Linux", nodename="reiss-lenovo", ...) = 0
brk(NULL) = 0x20a5000
brk(0x20a61c0) = 0x20a61c0
arch_prctl(ARCH_SET_FS, 0x20a5880) = 0
readlink("/proc/self/exe", "/home/cr4bd/spring2017/cs4630/sl"... , 4096) = 62
brk(0x20c71c0) = 0x20c71c0
brk(0x20c8000) = 0x20c8000
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
fstat(1, st_mode=S_IFCHR|0620, st_rdev=makedev(136, 1), ...) = 0
write(1, "Hello, World!\n", 14) = 14
exit_group(14) = ?
+++ exited with 14 +++
```

standard library shutdown

# dynamically linked hello.exe

```
gcc -no-pie -o hello.exe hello.s
```

```
strace ./hello.exe:
```

```
execve("./hello.exe", [ "./hello.exe" ], [ /* 46 vars */ ]) = 0
...
mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7fdfeeb39000
access("/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory)
open("/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
fstat(3, st_mode=S_IFREG|0644, st_size=137808, ...) = 0
...
open("/lib/x86_64-linux-gnu/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
read(3, "\177ELF\2\1\1\3\0\0\0\0\0\0\0\3\0>\0\1\0\0\0P\t\2\0\0\0\0\0"... , 832) = 832
fstat(3, st_mode=S_IFREG|0755, st_size=1864888, ...) = 0
mmap(NULL, 3967392, PROT_READ|PROT_EXEC, ..., 3, 0) = 0x7fdfee54d000
mprotect(0x7fdfee70c000, 2097152, PROT_NONE) = 0
mmap(0x7fdfee90c000, 24576, PROT_READ|PROT_WRITE, ..., 3, 0x1bf000) = 0x7fdfee90c000
mmap(0x7fdfee912000, 14752, PROT_READ|PROT_WRITE, ..., -1, 0) = 0x7fdfee912000
close(3) = 0
...
write(1, "Hello, World!\n", 14) = 14
exit_group(14) = ?
+++ exited with 14 +++
```

# dynamically linked hello.exe

```
gcc -no-pie -o hello.exe hello.s
```

```
strace ./hello.exe:
```

```
execve("./hello.exe", [ "./hello.exe" ], [ /* 46 vars */ ]) = 0
...
mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7fdfeeb39000
access("/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory)
open("/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
fstat(3, st_mode=S_IFREG|0644, st_size=137808, ...) = 0
...
open("/lib/x86_64-linux-gnu/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
read(3, "\177ELF\2\1\1\3\0\0\0\0\0\0\0\0\3\0>\0\1\0\0\0P\t\2\0\0\0\0\0"... , 832) = 832
fstat(3, st_mode=S_IFREG|0755, st_size=1864888, ...) = 0
mmap(NULL, 3967392, PROT_READ|PROT_EXEC, ..., 3, 0) = 0x7fdfee54d000
mprotect(0x7fdfee70, 0, PROT_READ|PROT_WRITE|PROT_EXEC) = 0
mmap(0x7fdfee90c000, 1024, PROT_READ|PROT_WRITE|PROT_EXEC, 0, 0, 0) = 0x7fdfee90c000
mmap(0x7fdfee912000, 1024, PROT_READ|PROT_WRITE|PROT_EXEC, 0, 0, 0) = 0x7fdfee912000
close(3) = 0
...
write(1, "Hello, World!\n", 14) = 14
exit_group(14) = ?
+++ exited with 14 +++
```

the standard C library (includes puts)



# dynamically linked hello.exe

```
gcc -no-pie -o hello.exe hello.s
```

```
strace ./hello.exe:
```

```
execve("./hello.exe", [ "./hello.exe" ], [ /* 46 vars */ ]) = 0
```

```
...
```

```
mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7fdfeeb39000
```

```
access("/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory)
```

```
open("/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
```

```
fstat(3, st_mode=S_IFREG|0644, st_size=137808, ...) = 0
```

```
...
```

```
open("/lib/x86_64-linux-gnu/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
```

```
read(3, "\177ELF\2\1\1\3\0\0\0\0\0\0\0\3\0>\0\1\0\0\0P\t\2\0\0\0\0\0"... , 832) = 832
```

```
fstat(3, st_mode=S_IFREG|0755, st_size=1864888, ...) = 0
```

```
mmap(NULL, 3967392, PROT_READ|PROT_EXEC, ..., 3, 0) = 0x7fdfee54d000
```

```
mprotect(0x7fdfee70c000, 12288, PROT_READ|PROT_WRITE) = 0
```

```
mmap(0x7fdfee90c000, 12288, PROT_READ|PROT_WRITE, ..., 3, 0) = 0x7fdfee90c000
```

```
mmap(0x7fdfee912000, 12288, PROT_READ|PROT_WRITE, ..., 3, 0) = 0x7fdfee912000
```

```
close(3) = 0
```

```
...
```

```
write(1, "Hello, World!\n", 14) = 14
```

```
exit_group(14) = ?
```

```
+++ exited with 14 +++
```

memory allocation (different method)

0x7fdfee90c000  
fee912000

# dynamically linked hello.exe

```
gcc -no-pie -o hello.exe hello.s
```

```
strace ./hello.exe:
```

```
execve("./hello.exe", [ "./hello.exe" ], [ /* 46 vars */ ]) = 0
...
mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7fdfeeb39000
access("/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory)
open("/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
fstat(3, st_mode=S_IFREG|0644, st_size=137808, ...) = 0
...
open("/lib/x86_64-linux-gnu/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
read(3, "\177ELF\2\1\1\3\0\0\0\0\0\0\0\3\0>\0\1\0\0\0P\t\2\0\0\0\0\0"... , 832) = 832
fstat(3, st_mode=S_IFREG|0755, st_size=1864888, ...) = 0
mmap(NULL, 3967392, PROT_READ|PROT_EXEC, ..., 3, 0) = 0x7fdfee54d000
mprotect(0x7fdfee70c000, 245, PROT_READ|PROT_WRITE|PROT_EXEC) = 0
mmap(0x7fdfee90c000, 245, PROT_READ|PROT_WRITE|PROT_EXEC, ..., 3, 0) = 0x7fdfee90c000
mmap(0x7fdfee912000, 147, PROT_READ|PROT_WRITE|PROT_EXEC, ..., 3, 0) = 0x7fdfee912000
close(3) = 0
...
write(1, "Hello, World!\n", 14) = 14
exit_group(14) = ?
+++ exited with 14 +++
```

read standard C library header

# dynamically linked hello.exe

```
gcc -no-pie -o hello.exe hello.s
```

```
strace ./hello.exe:
```

```
execve("./hello.exe", [ "./hello.exe" ], [ /* 46 vars */ ]) = 0
...
mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7fdfeeb39000
access("/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory)
open("/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
fstat(3, st_mode=S_IFREG|0644, st_size=137808, ...) = 0
...
open("/lib/x86_64-linux-gnu/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
read(3, "\177ELF\2\1\1\3\0\0\0\0\0\0\0\0\3\0>\0\1\0\0\0P\t\2\0\0\0\0\0"... , 832) = 832
fstat(3, st_mode=S_IFREG|0755, st_size=1864888, ...) = 0
mmap(NULL, 3967392, PROT_READ|PROT_EXEC, ..., 3, 0) = 0x7fdfee54d000
mprotect(0x7fdfee700000, 1048576, PROT_READ|PROT_WRITE) = 0
mmap(0x7fdfee90c000, 1048576, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7fdfee90c000
mmap(0x7fdfee912000, 1048576, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7fdfee912000
close(3) = 0
...
write(1, "Hello, World!\n", 14) = 14
exit_group(14) = ?
+++ exited with 14 +++
```

load standard C library (3 = opened file)

# dynamically linked hello.exe

```
gcc -no-pie -o hello.exe hello.s
```

```
strace ./hello.exe:
```

```
execve("./hello.exe", [ "./hello.exe" ], [ /* 46 vars */ ]) = 0
...
mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7fdfeeb39000
access("/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory)
open("/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
fstat(3, st_mode=S_IFREG|0644, st_size=137808, ...) = 0
...
open("/lib/x86_64-linux-gnu/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
read(3, "\177ELF\2\1\1\3\0\0\0\0\0\0\0\3\0>\0\1\0\0\0P\t\2\0\0\0\0\0"... , 832) = 832
fstat(3, st_mode=S_IFREG|0755, st_size=1864888, ...) = 0
mmap(NULL, 3967392, PROT_READ|PROT_EXEC, ..., 3, 0) = 0x7fdfee54d000
mprotect(0x7fdfee54d000, 3967392, PROT_READ|PROT_WRITE|PROT_EXEC) = 0
mmap(0x7fdfee54d000, 3967392, PROT_READ|PROT_WRITE|PROT_EXEC, ..., 3, 0) = 0x7fdfee90c000
mmap(0x7fdfee90c000, 3967392, PROT_READ|PROT_WRITE|PROT_EXEC, ..., 3, 0) = 0x7fdfee90c000
close(3) = 0
...
write(1, "Hello, World!\n", 14) = 14
exit_group(14) = ?
+++ exited with 14 +++
```

allocate zero-initialized data segment for C library

## where's the linker

Where's the code that calls `open("../libc.so.6")`?

Could check `hello.exe` — it's not there!

## where's the linker

Where's the code that calls `open("../libc.so.6")`?

Could check `hello.exe` — it's not there!

instead: "interpreter" `/lib64/ld-linux-x86-64.so.2`

on Linux: contains loading code instead of core OS  
OS loads it instead of program

# objdump — the interpreter

excerpt from `objdump -sx hello.exe`:

Program Header:

...

```
INTERP off      0x0000238 vaddr 0x0400318 paddr 0x0400238 align 2**0  
      filesz 0x000001c memsz 0x000001c flags r--
```

...

Contents of section `.interp`:

```
400318 2f6c6962 36342f6c 642d6c69 6e75782d /lib64/ld-linux-  
400328 7838362d 36342e73 6f2e3200 x86-64.so.2.
```

# dynamic linking: what to load? (1)

excerpt from `objdump -sx hello.exe`:

Program Header:

```
...  
  DYNAMIC off      0x000000000000002e20 vaddr 0x000000000000403e20 paddr 0x000000000000403e20  
          filesz 0x000000000000001d0 memsz 0x000000000000001d0 flags rw-
```

Dynamic Section:

```
  NEEDED          libc.so.6  
  INIT            0x000000000000401000  
  ...  
  STRTAB          0x000000000000400420  
  ...
```

program header: identifies where dynamic linking info is

dynamic linking info: array of key-value pairs

- needed libraries

- constructor locations ('INIT')

- string table location



# dynamic linking: what to load? (2)

excerpt from `objdump -sx hello.exe`:

Program Header:

```
...  
DYNAMIC off      0x000000000000002e20 vaddr 0x000000000000403e20 paddr 0x000000000000403e  
      filesz 0x000000000000001d0 memsz 0x000000000000001d0 flags rw-  
...
```

Dynamic Section:

```
NEEDED          libc.so.6  
INIT            0x000000000000401000  
...  
STRTAB          0x000000000000400420  
...  
...  
403e20 01000000 00000000 01000000 00000000 .....  
403e30 0c000000 00000000 00104000 00000000 .....  
.....
```

---

type `0x1` = `"DT_NEEDED"` (from ELF manual)

value `0x1` = string table entry 1

---

type `0xC` = `"DT_INIT"`

value `0x401000`

# dynamic linking: what to load? (2)

excerpt from `objdump -sx hello.exe`:

Program Header:

```
...  
DYNAMIC off      0x000000000000002e20 vaddr 0x000000000000403e20 paddr 0x000000000000403e  
      filesz 0x000000000000001d0 memsz 0x000000000000001d0 flags rw-  
...
```

Dynamic Section:

```
NEEDED          libc.so.6  
INIT            0x000000000000401000  
...  
STRTAB          0x000000000000400420  
...  
...  
403e20 01000000 00000000 01000000 00000000 .....  
403e30 0c000000 00000000 00104000 00000000 .....  
.....
```

---

type 0x1 = "DT\_NEEDED" (from ELF manual)

value 0x1 = *string table entry 1*

---

type 0xC = "DT\_INIT"

value 0x401000

# dynamic linking: what to load? (2)

excerpt from `objdump -sx hello.exe`:

Program Header:

```
...  
DYNAMIC off      0x000000000000002e20 vaddr 0x000000000000403e20 paddr 0x000000000000403e  
      filesz 0x000000000000001d0 memsz 0x000000000000001d0 flags rw-  
...
```

Dynamic Section:

```
NEEDED          libc.so.6  
INIT           0x000000000000401000  
...  
STRTAB         0x000000000000400420  
...
```

```
...  
403e20 01000000 00000000 01000000 00000000 .....  
403e30 0c000000 00000000 00104000 00000000 .....  
-----
```

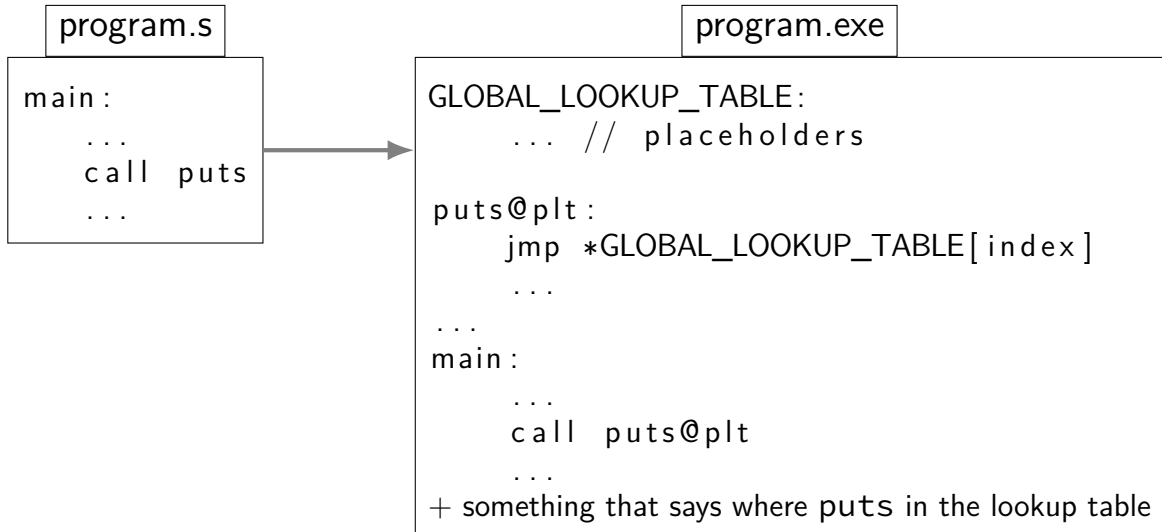
type 0x1 = "DT\_NEEDED" (from ELF manual)

value 0x1 = string table entry 1

type 0xC = "DT\_INIT"

value 0x401000

# adding linker stubs



# dynamic linking information

symbol table in libraries: list of functions/variables to find  
with their locations in the library

relocation records in programs: list of functions/variables  
with locations (probably in lookup table) to fill in

# dynamically linked puts (non-lazy)

## DYNAMIC RELOCATION RECORDS

OFFSET	TYPE	VALUE
0000000000403ff0	R_X86_64_GLOB_DAT	__libc_start_main@GLIBC_2.34
0000000000403ff8	R_X86_64_GLOB_DAT	__gmon_start__@Base
0000000000403fe8	R_X86_64_JUMP_SLOT	puts@GLIBC_2.2.5

...

## Text:

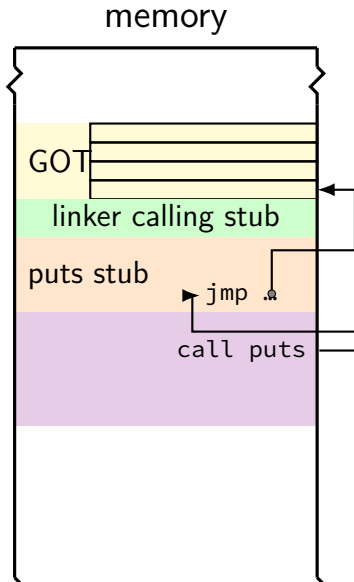
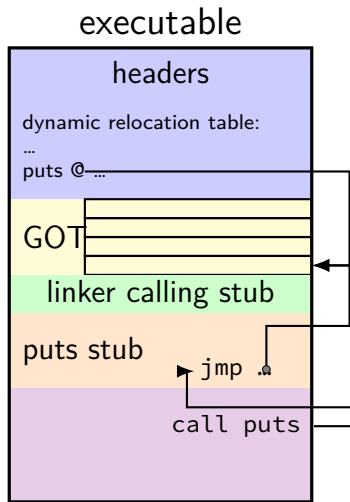
```
0000000000401030 <puts@plt>:  
    401030:      ff 25 b2 2f 00 00          jmp     *0x2fb2(%rip)           # 403fe8 <puts@plt>
```

stub reads pointer from 0x403fe8, jump to location

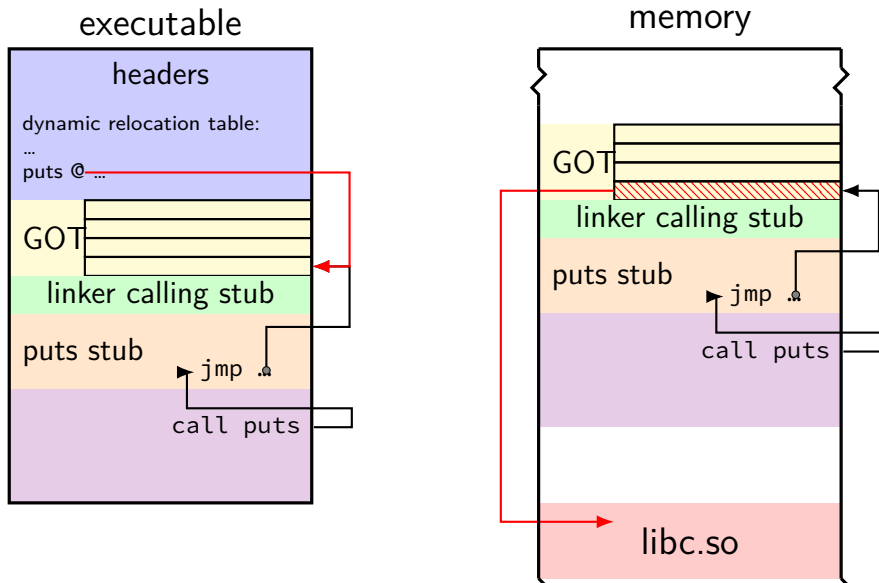
0x403fe8 part of 'global offset table' (GOT)

relocation table entry indicates where puts pointer goes

# dynamic puts (picture)

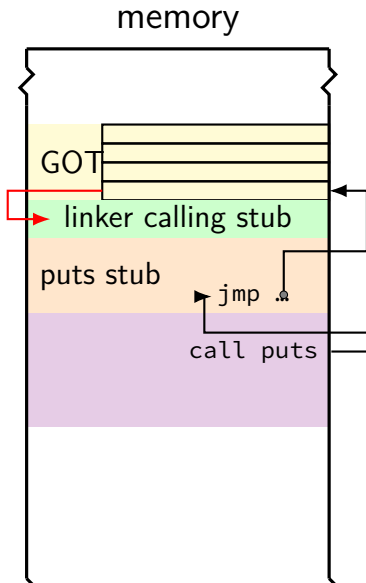
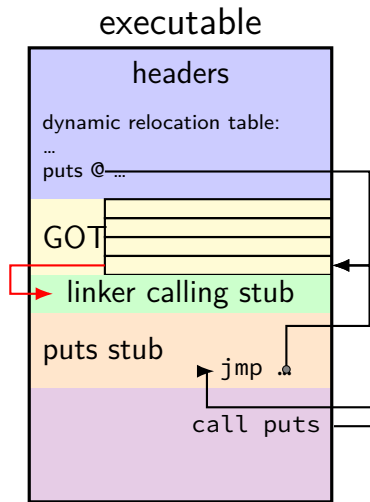


# dynamic puts (picture)

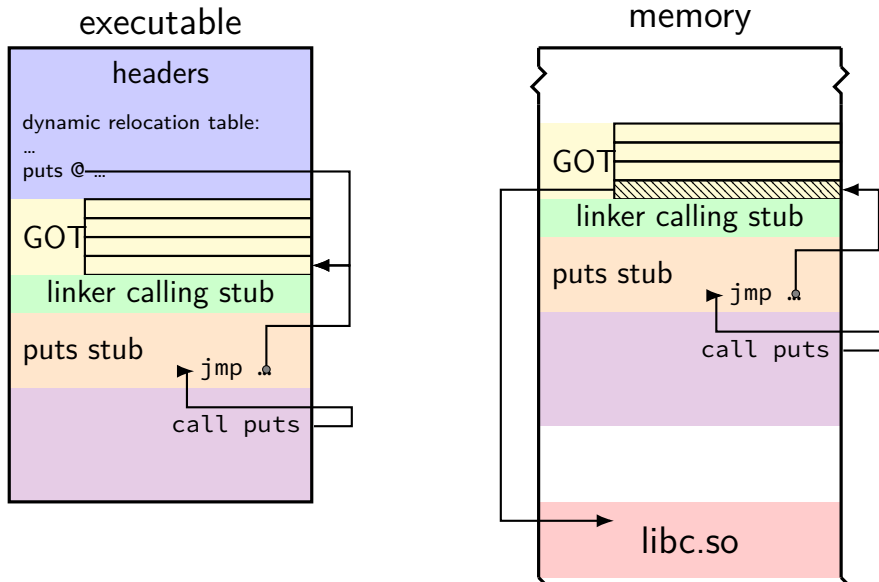




# dynamic puts (picture)



# dynamic puts (picture)



# lazy binding

```
0000000000401030 <puts@plt>:  
  401030:          ff 25 62 22 00 00      jmp     *0x2262(%rip)          # 403298 <puts@GLIBC_2.2  
...  
Contents of section .got.plt:  
  403280 a0304000 00000000 00000000 00000000  .0@.....  
  403290 00000000 00000000 36104000 00000000  .....6.@.....
```

initial contents of  $0x403298 = 0x401036$  (in .got.plt)  
not part of standard library????

# lazy binding

```
0000000000401030 <puts@plt>:
  401030:          ff 25 62 22 00 00          jmp     *0x2262(%rip)          # 403298 <puts@GLIBC_2.2
...
Contents of section .got.plt:
 403280 a0304000 00000000 00000000 00000000  .0@.....
 403290 00000000 00000000 36104000 00000000  ....6.@.....
```

initial contents of  $0x403298 = 0x401036$  (in .got.plt)  
not part of standard library????

code found at  $0x401036$  is routine to invoke dynamic linker code:

```
401020:          ff 35 62 22 00 00          push    0x2262(%rip)
# 403288 <_GLOBAL_OFFSET_TABLE_+0x8>
401026:          ff 25 64 22 00 00          jmp     *0x2264(%rip)
# 403290 <_GLOBAL_OFFSET_TABLE_+0x10>
...
401036:          68 00 00 00 00          push    $0x0
40103b:          e9 e0 ff ff ff          jmp     401020 <_init+0x20>
```

# lazy binding

with lazy binding turned on (not always done)

GOT loaded with address of linker routine hard-coded in executable

first call to puts:

- invoke dynamic linker routine pointed to by GOT

- linker routine fills in puts address in 0x404018

- then jumps to puts

second (and later) call to puts

- 0x404018 contains real address of puts, no indirection

# lazy binding pro/con

## advantages:

- faster program loading
- no overhead for unused code (often a lot of stuff)

## disadvantages:

- can move errors (missing functions, etc.) to runtime
- possibly more total overhead
- means global offset table needs to be writable?

# preview: exploits and dynamic linking

later we'll talk about memory error exploits  
buffer overflows, etc.

common goal: convert memory overwrite to running code  
bug that “just” allows overwriting memory somewhere  
easy to cause crashes...  
but not so easy to do something in particular

global offset table: function pointers in known location  
useful to overwrite in exploits