

address space layout randomization (ASLR)

assume: addresses don't leak

choose *random* addresses each time

for *everything*, not just the stack

enough possibilities that attacker won't "get lucky"

should prevent exploits — can't write GOT/shellcode location

Linux stack randomization (x86-64)

1. choose random number between 0 and 0x3F FFFF
2. stack starts at 0x7FFF FFFF FFFF - *random number* × 0x1000
randomization disabled? *random number* = 0



16 GB range!

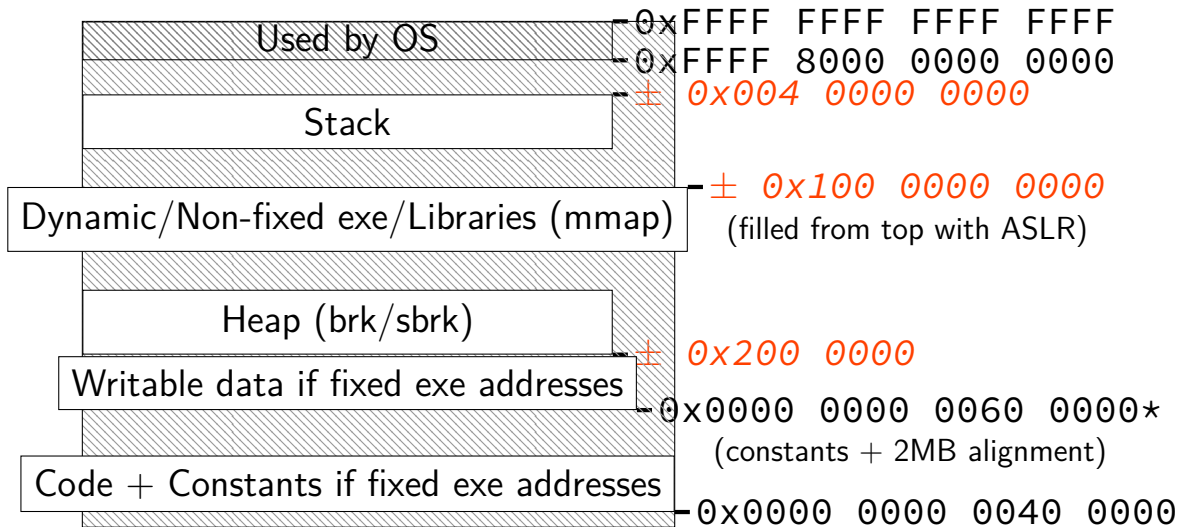
Linux stack randomization (x86-64)

1. choose random number between 0 and $0x3F\ FFFF$
2. stack starts at $0x7FFF\ FFFF\ FFFF - random\ number \times 0x1000$
randomization disabled? $random\ number = 0$

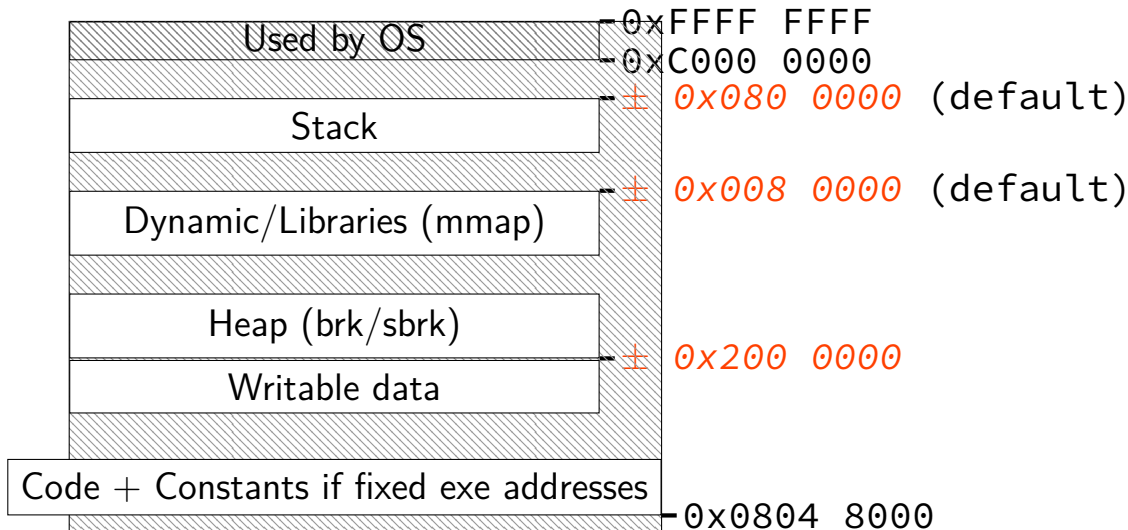


16 GB range!

program memory (x86-64 Linux; ASLR)



program memory (x86-32 Linux; ASLR)



how much guessing?

gaps change by multiples of page (4K)

lower 12 bits are *fixed*

64-bit: *huge* ranges — need millions of guesses

about *30 randomized bits* in addresses

32-bit: *smaller* ranges — hundreds of guesses

only about *8 randomized bits* in addresses

why? only 4 GB to work with!

can be configured higher — but larger gaps

why do we get multiple guesses?

why do we get multiple guesses?

wrong guess might not crash

wrong guess might not crash whole application

e.g. server that uses multiple processes

local programs we can repeatedly run

servers that are automatically restarted

dependencies between segments (1)

```
$ objdump -x foo.exe
```

```
...
```

LOAD	off	0x0000000000000000	vaddr	0x0000000000000000	paddr	0x000
	filesz	0x00000000000000620	memsz	0x00000000000000620	flags	r--
LOAD	off	0x00000000000001000	vaddr	0x00000000000001000	paddr	0x000
	filesz	0x00000000000000205	memsz	0x00000000000000205	flags	r-x
LOAD	off	0x00000000000002000	vaddr	0x00000000000002000	paddr	0x000
	filesz	0x00000000000000150	memsz	0x00000000000000150	flags	r--
LOAD	off	0x00000000000002db8	vaddr	0x00000000000003db8	paddr	0x000
	filesz	0x0000000000000025c	memsz	0x00000000000000260	flags	rw-

4 separately loaded segments: can we choose random addresses for each?

dependencies between segments (2)

```
00000000000001050 <__printf_chk@plt>:  
    1050:          f3 0f 1e fa                endbr64  
    1054:          f2 ff 25 75 2f 00 00      bnd jmpq *0x2f75(%rip)  
    105b:          0f 1f 44 00 00            nopl    0x0(%rax,%rax,1)
```

dependency from 2nd LOAD (0x1000-0x1205) to 4th LOAD
(0x3db8-0x4018)

uses relative addressing rather than linker filling in address

dependencies between segments (3)

```
00000000000001060 <main>:
    1060:      f3 0f 1e fa                endbr64
    1064:      50                        push    %rax
    1065:      8b 15 a5 2f 00 00          mov     0x2fa5(%rip),%edx
# 4010 <global>
    106b:      48 8d 35 92 0f 00 00      lea     0xf92(%rip),%rsi
# 2004 <_IO_stdin_used+0x4>
    1072:      31 c0                      xor     %eax,%eax
    1074:      bf 01 00 00 00            mov     $0x1,%edi
    1079:      e8 d2 ff ff ff            callq   1050 <__printf_chk@p
```

dependency from 2nd LOAD (0x1000-0x1205) to 3rd LOAD (0x2000-0x2150)

uses relative addressing rather than linker filling in address

why is this done?

Linux made a choice:

no editing code when loading programs, libraries

allows same code to be loaded in multiple processes

danger of leaking pointers

any stack pointer? know everything on the stack!

any pointer within executable? know everything in the executable!

any pointer to a particular library? know everything in library!

exercice: using a leak (1)

```
class Foo {  
    virtual const char *bar() { ... }  
};  
...  
Foo *f = new Foo;  
printf("%s\n", f);
```

Part 1: What address is most likely leaked by the above?

- A. the location of the Foo object allocated on the heap
- B. the location of the first entry in Foo's VTable"
- C. the location of the first instruction of Foo::Foo() (Foo's compiler-generated constructor)"
- D. the location of the stack pointer

exercise: using a leak (2)

```
class Foo {  
    virtual const char *bar() { ... }  
};  
...  
Foo *f = new Foo;  
char *p = new char[1024];  
printf("%s\n", f);
```

if leaked value was 0x822003 and in a debugger (with **different randomization**):

- stack pointer was 0x7ffff000

- Foo::bar's address was 0x400000

- f's address was 0x900000


- f's Vtable's address was 0x403000

- a "gadget" address from the main executable was 0x401034

- a "gadget" address from the C library was 0x2aaaa40034

- p's address was 0x901000

exes, libraries stay together



relocating: Windows

Windows will *edit code* to relocate
not everything uses a GOT-like lookup table

typically one fixed location per program/library **per boot**
same address used across all instances of program/library
still allows sharing memory

fixup once per program/library per boot
before ASLR: code could be pre-relocated

Windows + Visual Studio had 'full' ASLR by default since 2010

Windows ASLR limitation

same address in all programs — not very useful against local exploits

PIC: Linux, OS X

Linux, OS X: position-independent code

allows libraries code pages to be shared

...even if loaded at different addresses

avoids per-boot randomization of Windows, but...

exercise: avoiding absolute addresses

```
foo:                                lookupTable:
    movl    $3, %eax                .quad returnOne
    cmpq    $5, %rdi                .quad returnTwo
    ja      defaultCase              .quad returnOne
    jmp     *lookupTable(,%rdi,8)    .quad returnTwo
returnOne:                           .quad returnOne
    movl    $1, %eax                .quad returnOne
    ret
returnTwo:
    movl    $2, %eax
defaultCase:
    ret
```

exercise: rewrite this without absolute addresses

but fast

PIE jump-table

```
foo:
    movl    $3, %eax
    cmpq    $5, %rdi
    ja      retDefault
    leaq     jumpTable(%rip),%rax
    movslq   (%rax,%rdi,4),%rdx
    addq     %rdx, %rax
    jmp      *%rax
returnTwo:
    movl    $2, %eax
    ret
returnOne:
    movl    $1, %eax
defaultCase:
    ret
```

```
.section      .rodata
jumpTable:
    .long    returnOne-jumpTable
    .long    returnTwo-jumpTable
    .long    returnOne-jumpTable
    .long    returnTwo-jumpTable
    .long    returnOne-jumpTable
    .long    returnOne-jumpTable
```

PIE jump-table

```
foo:
    movl    $3, %eax
    cmpq    $5, %rdi
    ja      retDefault
    leaq     jumpTable(%rip),%rax
    movslq   (%rax,%rdi,4),%rdx
    addq     %rdx, %rax
    jmp      *%rax
returnTwo:
    movl    $2, %eax
    ret
returnOne:
    movl    $1, %eax
defaultCase:
    ret
```

```
.section      .rodata
jumpTable:
    .long   returnOne-jumpTable
    .long   returnTwo-jumpTable
    .long   returnOne-jumpTable
    .long   returnTwo-jumpTable
    .long   returnOne-jumpTable
    .long   returnOne-jumpTable
```

PIE jump-table

```
000000000000007ab <foo>:
b8 03 00 00 00      mov     $0x3,%eax
48 83 ff 05         cmp     $0x5,%rdi
77 1b              ja      7d0 <foo+0x25>
48 8d 05 ab 00 00 00 lea     0xab(%rip),%rax      # 868
48 63 14 b8         movslq  (%rax,%rdi,4),%rdx
48 01 d0           add     %rdx,%rax
ff e0             jmpq    *%rax
b8 02 00 00 00     mov     $0x2,%eax
c3               retq
b8 01 00 00 00     mov     $0x1,%eax
c3               retq
...
@ 868: -156 /* offset */
@ 870: -162
...
```

PIE jump-table

```
000000000000007ab <foo>:
b8 03 00 00 00      mov     $0x3,%eax
48 83 ff 05         cmp     $0x5,%rdi
77 1b              ja      7d0 <foo+0x25>
48 8d 05 ab 00 00 00 lea     0xab(%rip),%rax      # 868
48 63 14 b8         movslq  (%rax,%rdi,4),%rdx
48 01 d0           add     %rdx,%rax
ff e0             jmpq    *%rax
b8 02 00 00 00     mov     $0x2,%eax
c3               retq
b8 01 00 00 00     mov     $0x1,%eax
c3               retq
...
@ 868: -156 /* offset */
@ 870: -162
...
```


PIE jump-table

```
000000000000007ab <foo>:
b8 03 00 00 00      mov     $0x3,%eax
48 83 ff 05          cmp     $0x5,%rdi
77 1b                ja      7d0 <foo+0x25>
48 8d 05 ab 00 00 00 lea     0xab(%rip),%rax      # 868
48 63 14 b8          movslq  (%rax,%rdi,4),%rdx
48 01 d0             add     %rdx,%rax
ff e0               jmpq    *%rax
b8 02 00 00 00      mov     $0x2,%eax
c3                  retq
b8 01 00 00 00      mov     $0x1,%eax
c3                  retq
...
@ 868: -156 /* offset */
@ 870: -162
...
```

added cost

replace `jmp *jumpTable(,%rdi,8)`

with:

`lea` (get table address — with relative offset)

`movslq` (do table lookup of offset)

`add` (add to base)

`jmp` (to computed base)

32-bit x86 is worse

no relative addressing for mov, lea, ...

even changes “stubs” for printf:

// BEFORE: (fixed addresses)

08048310 <__printf_chk@plt>:

8048310: ff 25 10 a0 04 08 jmp *0x804a010

/ 0x804a010 == global offset table entry */*

// AFTER: (position-independent)

00000490 <__printf_chk@plt>:

490: ff a3 10 00 00 00 jmp *0x10(%ebx)

/ %ebx --- address of global offset table */*

/ needs to be set by caller */*

32-bit x86 is worse

no relative addressing for mov, lea, ...

even changes “stubs” for printf:

// BEFORE: (fixed addresses)

08048310 <__printf_chk@plt>:

```
8048310: ff 25 10 a0 04 08 jmp *0x804a010
/* 0x804a010 == global offset table entry */
```

// AFTER: (position-independent)

00000490 <__printf_chk@plt>:

```
490: ff a3 10 00 00 00 jmp *0x10(%ebx)
/* %ebx --- address of global offset table */
/* needs to be set by caller */
```

32-bit x86 is worse

no relative addressing for mov, lea, ...

even changes “stubs” for printf:

// BEFORE: (fixed addresses)

08048310 <__printf_chk@plt>:

8048310: ff 25 10 a0 04 08 jmp *0x804a010

/ 0x804a010 == global offset table entry */*

// AFTER: (position-independent)

00000490 <__printf_chk@plt>:

490: ff a3 10 00 00 00 jmp *0x10(%ebx)

/ %ebx --- address of global offset table */*

/ needs to be set by caller */*

PIE

position-independent executables (PIE)

no hardcoded addresses

alternative: *edit code (not global offset table) at load time*

Windows solution

GCC: `-pie -fPIE`

`-pie` is linking option

`-fPIE` is compilation option

related option: `-fPIC` (position independent code)

used to compile runtime-loaded libraries

hard-coded addresses? (64-bit)

```
int foo(long n) {  
    switch (n) {  
        case 0:  
        case 2:  
        case 4:  
        case 5:  
            return 1;  
        case 1:  
        case 3:  
            return 2;  
        default:  
            return 3;  
    }  
}
```

```
foo:  
    movl    $3, %eax  
    cmpq    $5, %rdi  
    ja      defaultCase  
    jmp     *lookupTable(,%rdi,8)  
    /* code for defaultCase, returnOne, ... */  
    .section .rodata  
lookupTable: /* read-only pointers: */  
    .quad   returnOne  
    .quad   returnTwo  
    .quad   returnOne  
    .quad   returnTwo  
    .quad   returnOne  
    .quad   returnOne
```

hard-coded addresses? (64-bit)

```
int foo(long n) {  
    switch (n) {  
        case 0:  
        case 2:  
        case 4:  
        case 5:  
            return 1;  
        case 1:  
        case 3:  
            return 2;  
        default:  
            return 3;  
    }  
}
```

```
400570 <foo>:  
b8 03 00 00 00      mov     $0x3,%eax  
48 83 ff 05         cmp     $0x5,%rdi  
                    /* jump to defaultCase: */  
77 12              ja      0x40058d  
                    /* lookup table jump: */  
ff 24 fd  
18 06 40 00         jmpq    *0x400618(,%rdi,8)  
...  
                    /* lookupTable @ 0x400618 */  
@ 400618: 0x400588 /* returnOne */  
@ 400620: 0x400582 /* returnTwo */  
@ 400628: 0x400588  
@ 400630: 0x400582
```


hard-coded addresses? (64-bit)

```
int foo(long n) {  
    switch (n) {  
        case 0:  
        case 2:  
        case 4:  
        case 5:  
            return 1;  
        case 1:  
        case 3:  
            return 2;  
        default:  
            return 3;  
    }  
}
```

```
400570 <foo>:  
b8 03 00 00 00      mov     $0x3,%eax  
48 83 ff 05         cmp     $0x5,%rdi  
                    /* jump to defaultCase: */  
77 12              ja      0x40058d  
                    /* lookup table jump: */  
ff 24 fd  
18 06 40 00         jmpq    *0x400618(,%rdi,8)  
...  
                    /* lookupTable @ 0x400618 */  
@ 400618: 0x400588 /* returnOne */  
@ 400620: 0x400582 /* returnTwo */  
@ 400628: 0x400588  
@ 400630: 0x400582
```


hard-coded addresses? (64-bit)

```
int foo(long n) {  
    switch (n) {  
        case 0:  
        case 2:  
        case 4:  
        case 5:  
            return 1;  
        case 1:  
        case 3:  
            return 2;  
        default:  
            return 3;  
    }  
}
```

```
400570 <foo>:  
b8 03 00 00 00      mov     $0x3,%eax  
48 83 ff 05         cmp     $0x5,%rdi  
                        /* jump to defaultCase: */  
77 12              ja      0x40058d  
                        /* lookup table jump: */  
ff 24 fd  
18 06 40 00         jmpq    *0x400618(,%rdi,8)  
...  
                        /* lookupTable @ 0x400618 */  
@ 400618: 0x400588 /* returnOne */  
@ 400620: 0x400582 /* returnTwo */  
@ 400628: 0x400588  
@ 400630: 0x400582
```

position independence cost (32-bit)

Overhead for -fPIE



position independence cost: Linux

geometric mean of SPECcpu2006 benchmarks on x86 Linux
with particular version of GCC, etc., etc.

64-bit: 2-3% (???)

“preliminary result”; couldn't find reliable published data

32-bit: 9-10%

depends on compiler, ...

position independence: deployment

common for a very long time in dynamic libraries

default for all executables in...

Microsoft Visual Studio 2010 and later
DYNAMICBASE linker option

OS since 10.7 (2011)

Fedora 23 (2015) and Red Hat Enterprise Linux 8 (2019) and later
default for “sensitive” programs earlier

Ubuntu 16.10 (2016) and later (for 64-bit), 17.10 (2017) and later
(for 32-bit)
default for “sensitive” programs earlier

backup slides