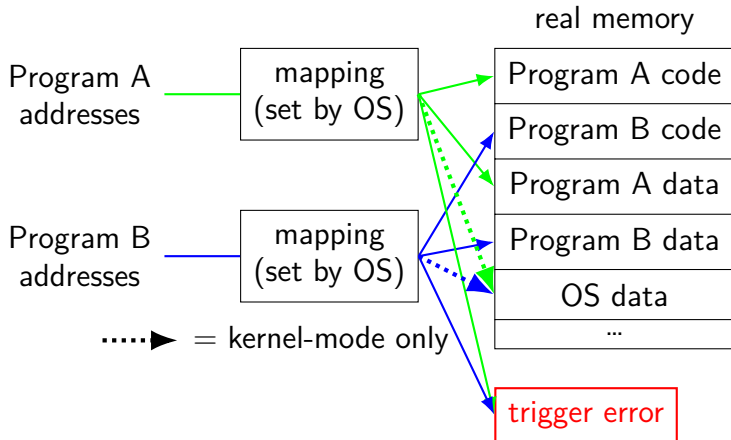


recall(?): virtual memory

illusion of *dedicated memory*



the mapping (set by OS)

program address range

0x0000 --- 0x0FFF

0x1000 --- 0x1FFF

...

0x40 0000 --- 0x40 0FFF

0x40 1000 --- 0x40 1FFF

0x40 2000 --- 0x40 2FFF

...

0x60 0000 --- 0x60 0FFF

0x60 1000 --- 0x60 1FFF

...

0x7FFF FF00 0000 — 0x7FFF FF00 0FFF

0x7FFF FF00 1000 — 0x7FFF FF00 1FFF

...

read?	write?
no	no
no	no

yes	no
yes	no
yes	no

yes	yes
yes	yes

yes	yes
yes	yes

real address

0x...
0x...
0x...

0x...
0x...

0x...
0x...

Virtual Memory

modern *hardware-supported* memory protection mechanism

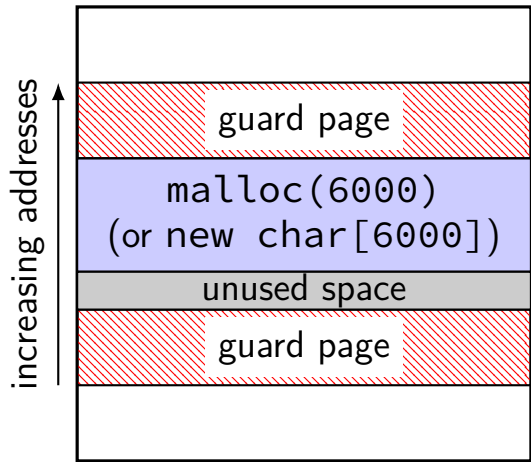
via *table*: OS decides *what memory program sees*
whether it's read-only or not

granularity of *pages* — typically 4KB

not in table — segfault (OS gets control)

malloc/new guard pages

the heap



guard pages

deliberate holes

accessing — segfault

call to OS to allocate (not very fast)

likely to 'waste' memory

guard around object? minimum 4KB object

guard pages for malloc/new

can implement malloc/new by placing guard pages around allocations

commonly done by real malloc/new's for *large allocations*

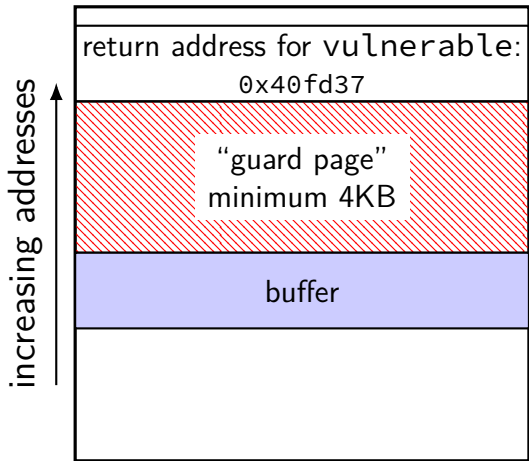
problem: minimum actual allocation 4KB

problem: substantially slower

example: “Electric Fence” allocator for Linux (early 1990s)

stack canary alternative

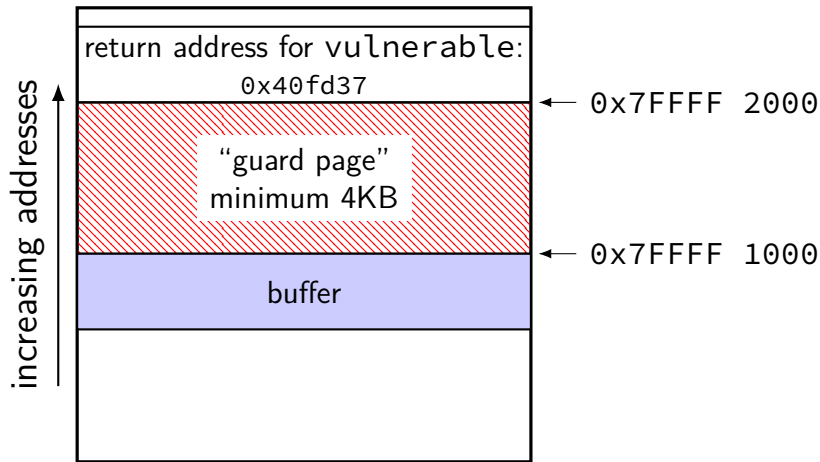
highest address (stack started here)



lowest address (stack grows here)

stack canary alternative

highest address (stack started here)

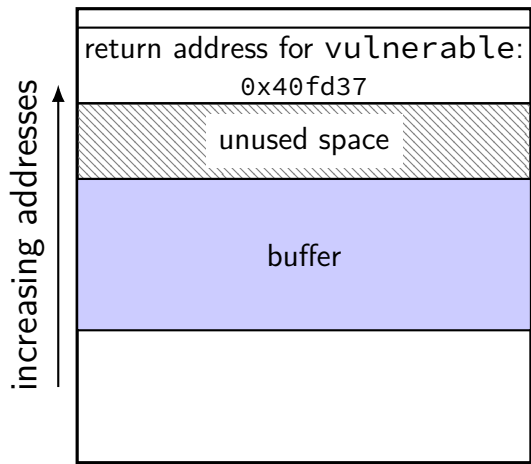


lowest address (stack grows here)

address	read	write
0x7FFFF2000-0x7FFFF2FFF	yes	yes
0x7FFFF1000-0x7FFFF1FFF	no	no
0x7FFFF0000-0x7FFFF0FFF	yes	yes

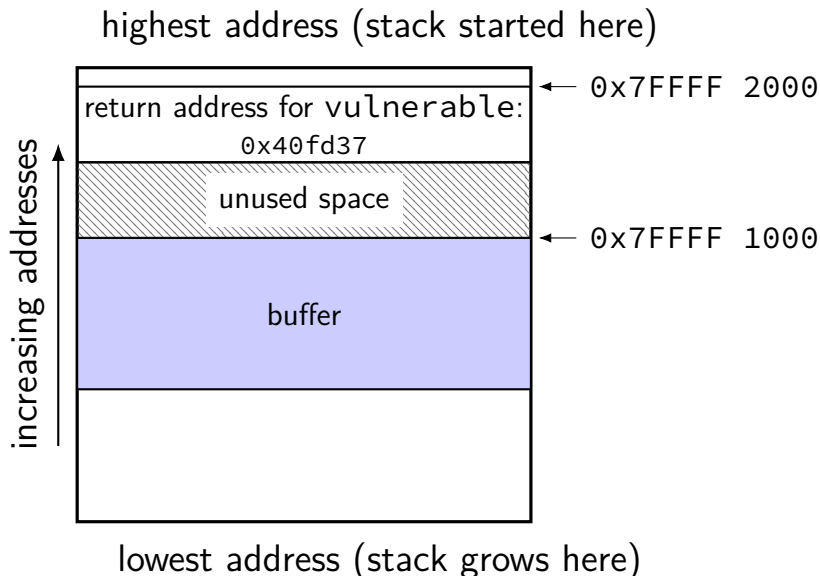
stack canary alternative 2

highest address (stack started here)



lowest address (stack grows here)

stack canary alternative 2



address	read	write
0x7FFFF2000-0x7FFFF2FFF	yes	yes
0x7FFFF1000-0x7FFFF1FFF	yes	<i>no</i>
0x7FFFF0000-0x7FFFF0FFF	yes	yes

exercise: guard page overhead

suppose heap allocations are:

- 100 000 objects of 100 bytes

- 1 000 objects of 1000 bytes

- 100 objects of approx. 10000 bytes

total allocation of approx 12 000 KB

assuming 4KB pages, estimate space overhead of using guard pages:

- for objects larger than 4096 bytes (1 page)

- for objects larger than 200 bytes

- for all objects

recall: function pointer targets

wanted to overwrite special pointer:

return addresses on stack

function pointers on in local variables

tables of function pointers used for inheritance

global offset table

last two: need to change infrequently

idea: make read-only

RELRO

RELocation **Read-Only**

Linux option: make dynamic linker structures read-only after startup

partial RELRO: everything but GOT pointers to library functions
notably includes C++ virtual function tables

full RELRO: everything including those pointers
requires disabling “lazy” linking
(could do without disabling — but slower (how much?) startup)

appears as ELF program header entry

a thought on permissions

if we can set memory non-writeable

how about non-executable?

we never want to execute things on the stack anyways, right?

write XOR execute

many names:

- W^X (write XOR execute)

- DEP (Data Execution Prevention)

- NX bit (No-eXecute) (hardware support)

- XD bit (eXecute Disable) (hardware support)

mark writeable memory as executable

how will users insert their machine code?

- can only code in application + libraries

- a problem, right?

hardware support for write XOR execute

everywhere today

not historically common

early x86: execute implied by read

NX support added with x86-64 and around 2000 for x86-32

deliberate use of writeable code

- “just-in-time” (JIT) compilers

 - fast virtual machine/language implementations

- some weird GCC features

- older “signals” on Linux

 - OS wrote machine code on stack for program to run

- couldn't even disable executable stacks without breaking applications

why doesn't $W \text{ xor } X$ solve the problem?

$W \text{ xor } X$ is “almost free”, keeps attacker from writing code?

problem: useful machine code is in program already

just need to find writable function pointer

saw special case: arc injection

happened to find useful code in existing application/library

turns out: almost always useful code

backup slides