



# stopping stack smashing?

how can you stop stack smashing?

# stopping stack smashing?

how can you stop stack smashing?

stop overrun — bounds-checking

stop return to attacker code

stop execution of attacker code

# exploit mitigations

idea: turn vulnerability to something less bad

e.g. crash instead of machine code execution

many of these targetted at buffer overflows

# mitigation agenda

we will look briefly at one mitigation — stack canaries

then look at exploits that don't care about it

then look at more flexible mitigations

then look at more flexible exploits

# mitigation priorities

effective? does it actually stop the attacker?

fast? how much does it hurt performance?

generic? does it require a recompile? rewriting software?

# address space layout randomization (ASLR)

assume: addresses don't leak

choose *random* addresses each time

for *everything*, not just the stack

*enough possibilities* that attacker won't "get lucky"

should prevent exploits — can't write GOT/shellcode location

# Linux stack randomization (x86-64)

1. choose random number between 0 and 0x3F FFFF
2. stack starts at 0x7FFF FFFF FFFF - *random number* × 0x1000  
randomization disabled? *random number* = 0



16 GB range!




# Linux stack randomization (x86-64)

1. choose random number between 0 and  $0x3F\ FFFF$
2. stack starts at  $0x7FFF\ FFFF\ FFFF - random\ number \times 0x1000$   
randomization disabled?  $random\ number = 0$

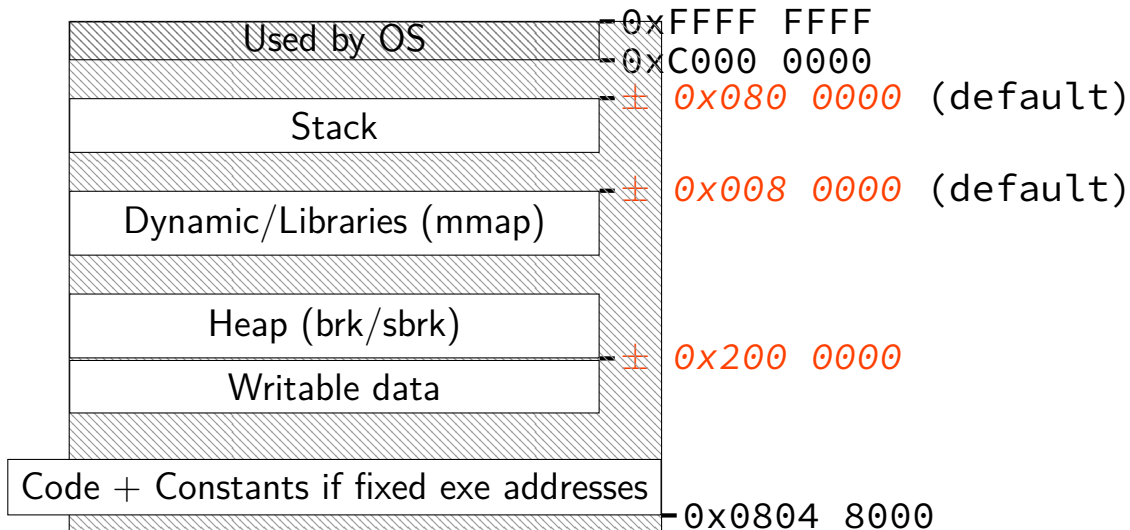


16 GB range!

# program memory (x86-64 Linux; ASLR)



# program memory (x86-32 Linux; ASLR)



# how much guessing?

gaps change by multiples of page (4K)

lower 12 bits are *fixed*

64-bit: *huge* ranges — need millions of guesses

about *30 randomized bits* in addresses

32-bit: *smaller* ranges — hundreds of guesses

only about *8 randomized bits* in addresses

why? only 4 GB to work with!

can be configured higher — but larger gaps

# why do we get multiple guesses?

why do we get multiple guesses?

wrong guess might not crash

wrong guess might not crash whole application

e.g. server that uses multiple processes

local programs we can repeatedly run

servers that are automatically restarted

# dependencies between segments (1)

```
$ objdump -x foo.exe
```

```
...
```

LOAD	off	0x0000000000000000	vaddr	0x0000000000000000	paddr	0x000
	filesz	0x00000000000000620	memsz	0x00000000000000620	flags	r--
LOAD	off	0x00000000000001000	vaddr	0x00000000000001000	paddr	0x000
	filesz	0x00000000000000205	memsz	0x00000000000000205	flags	r-x
LOAD	off	0x00000000000002000	vaddr	0x00000000000002000	paddr	0x000
	filesz	0x00000000000000150	memsz	0x00000000000000150	flags	r--
LOAD	off	0x00000000000002db8	vaddr	0x00000000000003db8	paddr	0x000
	filesz	0x0000000000000025c	memsz	0x00000000000000260	flags	rw-

4 separately loaded segments: can we choose random addresses for each?

## dependencies between segments (2)

```
00000000000001050 <__printf_chk@plt>:  
    1050:          f3 0f 1e fa                endbr64  
    1054:          f2 ff 25 75 2f 00 00      bnd jmpq *0x2f75(%rip)  
    105b:          0f 1f 44 00 00            nopl    0x0(%rax,%rax,1)
```

dependency from 2nd LOAD (0x1000-0x1205) to 4th LOAD  
(0x3db8-0x4018)

uses relative addressing rather than linker filling in address

## dependencies between segments (3)

```
00000000000001060 <main>:
    1060:      f3 0f 1e fa                endbr64
    1064:      50                        push    %rax
    1065:      8b 15 a5 2f 00 00          mov     0x2fa5(%rip),%edx
# 4010 <global>
    106b:      48 8d 35 92 0f 00 00      lea     0xf92(%rip),%rsi
# 2004 <_IO_stdin_used+0x4>
    1072:      31 c0                      xor     %eax,%eax
    1074:      bf 01 00 00 00            mov     $0x1,%edi
    1079:      e8 d2 ff ff ff            callq   1050 <__printf_chk@p
```

dependency from 2nd LOAD (0x1000-0x1205) to 3rd LOAD (0x2000-0x2150)

uses relative addressing rather than linker filling in address



# why is this done?

Linux made a choice:

no editing code when loading programs, libraries

allows same code to be loaded in multiple processes

# danger of leaking pointers

any stack pointer? know everything on the stack!

any pointer within executable? know everything in the executable!

any pointer to a particular library? know everything in library!

## exercice: using a leak (1)

```
class Foo {  
    virtual const char *bar() { ... }  
};  
...  
Foo *f = new Foo;  
printf("%s\n", f);
```

Part 1: What address is most likely leaked by the above?

- A. the location of the Foo object allocated on the heap
- B. the location of the first entry in Foo's VTable"
- C. the location of the first instruction of Foo::Foo() (Foo's compiler-generated constructor)"
- D. the location of the stack pointer

## exercise: using a leak (2)

```
class Foo {  
    virtual const char *bar() { ... }  
};  
...  
Foo *f = new Foo;  
char *p = new char[1024];  
printf("%s\n", f);
```

if leaked value was 0x822003 and in a debugger (with **different randomization**):

- stack pointer was 0x7ffff000

- Foo::bar's address was 0x400000

- f's address was 0x900000

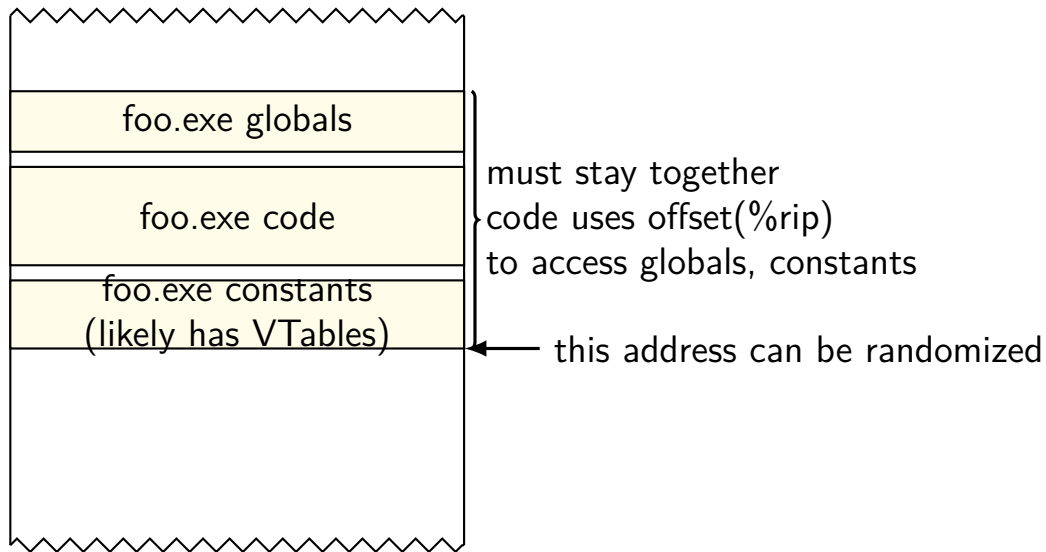
- f's Vtable's address was 0x403000

- a "gadget" address from the main executable was 0x401034

- a "gadget" address from the C library was 0x2aaaa40034

- p's address was 0x901000

## exes, libraries stay together



# relocating: Windows

Windows will *edit code* to relocate  
not everything uses a GOT-like lookup table

typically one fixed location per program/library **per boot**  
same address used across all instances of program/library  
still allows sharing memory

fixup once per program/library per boot  
before ASLR: code could be pre-relocated

Windows + Visual Studio had 'full' ASLR by default since 2010

# Windows ASLR limitation

same address in all programs — not very useful against local exploits

# PIC: Linux, OS X

Linux, OS X: position-independent code

allows libraries code pages to be shared

...even if loaded at different addresses

avoids per-boot randomization of Windows, but...



# exercise: avoiding absolute addresses

```
foo:                                lookupTable:
    movl    $3, %eax                .quad returnOne
    cmpq    $5, %rdi                .quad returnTwo
    ja      defaultCase              .quad returnOne
    jmp     *lookupTable(,%rdi,8)    .quad returnTwo
returnOne:                           .quad returnOne
    movl    $1, %eax                .quad returnOne
    ret                                .quad returnOne
returnTwo:
    movl    $2, %eax
defaultCase:
    ret
```

exercise: rewrite this without absolute addresses

but fast

# PIE jump-table

```
foo:
    movl    $3, %eax
    cmpq    $5, %rdi
    ja      retDefault
    leaq     jumpTable(%rip),%rax
    movslq   (%rax,%rdi,4),%rdx
    addq     %rdx, %rax
    jmp      *%rax
returnTwo:
    movl    $2, %eax
    ret
returnOne:
    movl    $1, %eax
defaultCase:
    ret
```

```
.section      .rodata
jumpTable:
    .long    returnOne-jumpTable
    .long    returnTwo-jumpTable
    .long    returnOne-jumpTable
    .long    returnTwo-jumpTable
    .long    returnOne-jumpTable
    .long    returnOne-jumpTable
```

# PIE jump-table

```
foo:
    movl    $3, %eax
    cmpq    $5, %rdi
    ja      retDefault
    leaq     jumpTable(%rip),%rax
    movslq   (%rax,%rdi,4),%rdx
    addq     %rdx, %rax
    jmp      *%rax
returnTwo:
    movl    $2, %eax
    ret
returnOne:
    movl    $1, %eax
defaultCase:
    ret
```

```
.section      .rodata
jumpTable:
    .long    returnOne-jumpTable
    .long    returnTwo-jumpTable
    .long    returnOne-jumpTable
    .long    returnTwo-jumpTable
    .long    returnOne-jumpTable
    .long    returnOne-jumpTable
```

# PIE jump-table

```
000000000000007ab <foo>:
b8 03 00 00 00      mov     $0x3,%eax
48 83 ff 05         cmp     $0x5,%rdi
77 1b              ja      7d0 <foo+0x25>
48 8d 05 ab 00 00 00 lea     0xab(%rip),%rax      # 868
48 63 14 b8         movslq  (%rax,%rdi,4),%rdx
48 01 d0           add     %rdx,%rax
ff e0             jmpq    *%rax
b8 02 00 00 00     mov     $0x2,%eax
c3               retq
b8 01 00 00 00     mov     $0x1,%eax
c3               retq
...
@ 868: -156 /* offset */
@ 870: -162
...
```

# PIE jump-table

```
000000000000007ab <foo>:
b8 03 00 00 00      mov     $0x3,%eax
48 83 ff 05         cmp     $0x5,%rdi
77 1b              ja      7d0 <foo+0x25>
48 8d 05 ab 00 00 00 lea     0xab(%rip),%rax      # 868
48 63 14 b8         movslq  (%rax,%rdi,4),%rdx
48 01 d0           add     %rdx,%rax
ff e0             jmpq    *%rax
b8 02 00 00 00     mov     $0x2,%eax
c3               retq
b8 01 00 00 00     mov     $0x1,%eax
c3               retq
...
@ 868: -156 /* offset */
@ 870: -162
...
```

# PIE jump-table

```
000000000000007ab <foo>:
b8 03 00 00 00      mov     $0x3,%eax
48 83 ff 05         cmp     $0x5,%rdi
77 1b              ja      7d0 <foo+0x25>
48 8d 05 ab 00 00 00 lea     0xab(%rip),%rax      # 868
48 63 14 b8         movslq  (%rax,%rdi,4),%rdx
48 01 d0           add     %rdx,%rax
ff e0             jmpq    *%rax
b8 02 00 00 00     mov     $0x2,%eax
c3               retq
b8 01 00 00 00     mov     $0x1,%eax
c3               retq
...
@ 868: -156 /* offset */
@ 870: -162
...
```

## added cost

replace `jmp *jumpTable(,%rdi,8)`

with:

`lea` (get table address — with relative offset)

`movslq` (do table lookup of offset)

`add` (add to base)

`jmp` (to computed base)

## 32-bit x86 is worse

no relative addressing for mov, lea, ...

even changes “stubs” for printf:

*// BEFORE: (fixed addresses)*

08048310 <\_\_printf\_chk@plt>:

8048310: ff 25 10 a0 04 08 jmp \*0x804a010

*/\* 0x804a010 == global offset table entry \*/*

*// AFTER: (position-independent)*

00000490 <\_\_printf\_chk@plt>:

490: ff a3 10 00 00 00 jmp \*0x10(%ebx)

*/\* %ebx --- address of global offset table \*/*

*/\* needs to be set by caller \*/*



## 32-bit x86 is worse

no relative addressing for mov, lea, ...

even changes “stubs” for printf:

*// BEFORE: (fixed addresses)*

08048310 <\_\_printf\_chk@plt>:

```
8048310: ff 25 10 a0 04 08 jmp *0x804a010
/* 0x804a010 == global offset table entry */
```

*// AFTER: (position-independent)*

00000490 <\_\_printf\_chk@plt>:

```
490: ff a3 10 00 00 00 jmp *0x10(%ebx)
/* %ebx --- address of global offset table */
/* needs to be set by caller */
```

## 32-bit x86 is worse

no relative addressing for mov, lea, ...

even changes “stubs” for printf:

*// BEFORE: (fixed addresses)*

08048310 <\_\_printf\_chk@plt>:

8048310: ff 25 10 a0 04 08 jmp \*0x804a010

*/\* 0x804a010 == global offset table entry \*/*

*// AFTER: (position-independent)*

00000490 <\_\_printf\_chk@plt>:

490: ff a3 10 00 00 00 jmp \*0x10(%ebx)

*/\* %ebx --- address of global offset table \*/*

*/\* needs to be set by caller \*/*

# PIE

position-independent executables (PIE)

no hardcoded addresses

alternative: *edit code (not global offset table) at load time*

Windows solution

GCC: `-pie -fPIE`

`-pie` is linking option

`-fPIE` is compilation option

related option: `-fPIC` (position independent code)

used to compile runtime-loaded libraries

# hard-coded addresses? (64-bit)

```
int foo(long n) {  
    switch (n) {  
        case 0:  
        case 2:  
        case 4:  
        case 5:  
            return 1;  
        case 1:  
        case 3:  
            return 2;  
        default:  
            return 3;  
    }  
}
```

```
foo:  
    movl    $3, %eax  
    cmpq    $5, %rdi  
    ja      defaultCase  
    jmp     *lookupTable(,%rdi,8)  
    /* code for defaultCase, returnOne, ... */  
    .section .rodata  
lookupTable: /* read-only pointers: */  
    .quad   returnOne  
    .quad   returnTwo  
    .quad   returnOne  
    .quad   returnTwo  
    .quad   returnOne  
    .quad   returnOne
```

# hard-coded addresses? (64-bit)

```
int foo(long n) {  
    switch (n) {  
        case 0:  
        case 2:  
        case 4:  
        case 5:  
            return 1;  
        case 1:  
        case 3:  
            return 2;  
        default:  
            return 3;  
    }  
}
```

```
400570 <foo>:  
b8 03 00 00 00      mov     $0x3,%eax  
48 83 ff 05         cmp     $0x5,%rdi  
                    /* jump to defaultCase: */  
77 12              ja      0x40058d  
                    /* lookup table jump: */  
ff 24 fd  
18 06 40 00         jmpq    *0x400618(,%rdi,8)  
...  
    /* lookupTable @ 0x400618 */  
@ 400618: 0x400588 /* returnOne */  
@ 400620: 0x400582 /* returnTwo */  
@ 400628: 0x400588  
@ 400630: 0x400582
```

# hard-coded addresses? (64-bit)

```
int foo(long n) {  
    switch (n) {  
        case 0:  
        case 2:  
        case 4:  
        case 5:  
            return 1;  
        case 1:  
        case 3:  
            return 2;  
        default:  
            return 3;  
    }  
}
```

```
400570 <foo>:  
b8 03 00 00 00      mov     $0x3,%eax  
48 83 ff 05         cmp     $0x5,%rdi  
                    /* jump to defaultCase: */  
77 12              ja      0x40058d  
                    /* lookup table jump: */  
ff 24 fd  
18 06 40 00         jmpq    *0x400618(,%rdi,8)  
...  
/* lookupTable @ 0x400618 */  
@ 400618: 0x400588 /* returnOne */  
@ 400620: 0x400582 /* returnTwo */  
@ 400628: 0x400588  
@ 400630: 0x400582
```


# hard-coded addresses? (64-bit)

```
int foo(long n) {  
    switch (n) {  
        case 0:  
        case 2:  
        case 4:  
        case 5:  
            return 1;  
        case 1:  
        case 3:  
            return 2;  
        default:  
            return 3;  
    }  
}
```

```
400570 <foo>:  
b8 03 00 00 00      mov     $0x3,%eax  
48 83 ff 05      cmp     $0x5,%rdi  
                /* jump to defaultCase: */  
77 12              ja      0x40058d  
                /* lookup table jump: */  
ff 24 fd  
18 06 40 00      jmpq    *0x400618(,%rdi,8)  
...  
                /* lookupTable @ 0x400618 */  
@ 400618: 0x400588 /* returnOne */  
@ 400620: 0x400582 /* returnTwo */  
@ 400628: 0x400588  
@ 400630: 0x400582
```

# position independence cost (32-bit)

Overhead for -fPIE





# position independence cost: Linux

geometric mean of SPECcpu2006 benchmarks on x86 Linux  
with particular version of GCC, etc., etc.

64-bit: 2-3% (???)

“preliminary result”; couldn't find reliable published data

32-bit: 9-10%

depends on compiler, ...

# position independence: deployment

common for a very long time in dynamic libraries

default for all executables in...

Microsoft Visual Studio 2010 and later  
DYNAMICBASE linker option

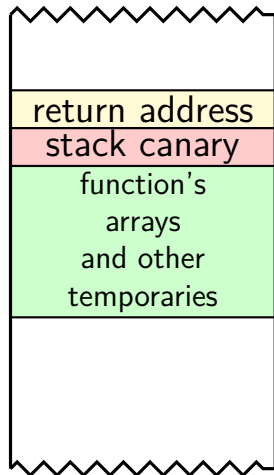
OS since 10.7 (2011)

Fedora 23 (2015) and Red Hat Enterprise Linux 8 (2019) and later  
default for “sensitive” programs earlier

Ubuntu 16.10 (2016) and later (for 64-bit), 17.10 (2017) and later  
(for 32-bit)  
default for “sensitive” programs earlier

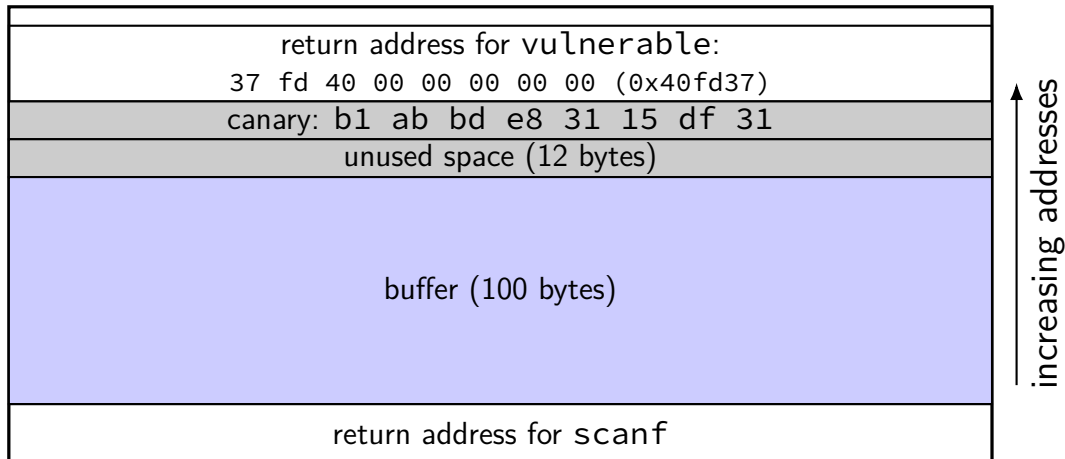
# compiler generated code

```
    pushq %rbx
    sub $0x20,%rsp
    /* copy value from thread-local storage */
    mov $0x28,%ebx
    mov %fs:(%rbx),%rax
    /* onto the stack */
    mov %rax,0x18(%rsp)
    /* clear register holding value */
    xor %eax, %eax
    ...
    ...
    /* copy value back from stack */
    mov 0x18(%rsp),%rax
    /* xor to compare */
    xor %fs:(%rbx),%rax
    /* if result non-zero, do not return */
    jne call_stack_chk_fail
    ret
call_stack_chk_fail:
```



# stack canary

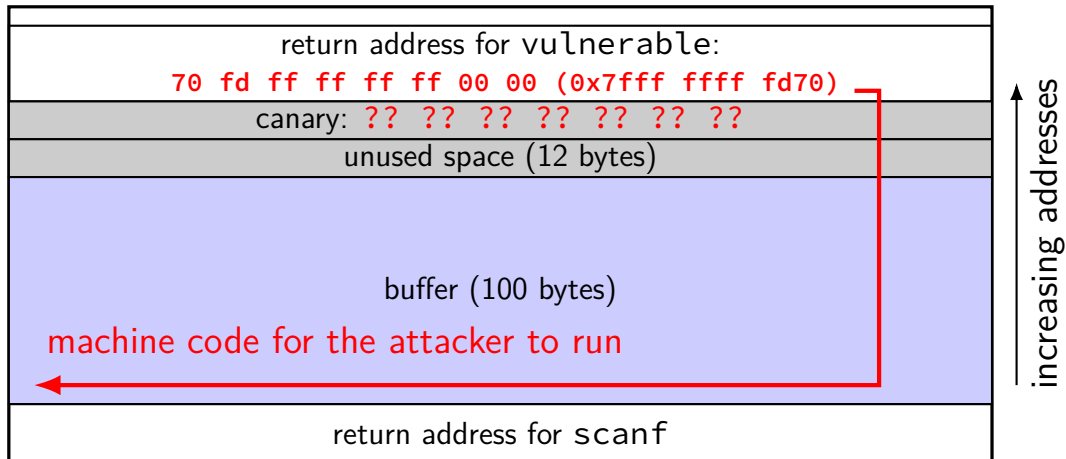
highest address (stack started here)



lowest address (stack grows here)

# stack canary

highest address (stack started here)



lowest address (stack grows here)

# stack canary hopes

overwrite return address  $\implies$  overwrite canary

canary is secret

# good choices of canary

*random* — guessing should not be practical  
not always — sometimes static or only  $2^{15}$  possible

GNU libc: canary contains:

leading `\0` (string terminator)  
    `printf %s` won't print it  
    copying a C-style string won't write it

a newline  
    read line functions can't input it

`\xFF`  
    hard to input?

# stack canaries implementation

“StackGuard” — 1998 paper proposing strategy

GCC: command-line options

- fstack-protector
- fstack-protector-strong
- fstack-protector-all

one of these often default

three differ in how many functions are ‘protected’

Microsoft C/C++ compiler: /GS

on by default



# stack canary overheads

less than 1% runtime if added to “risky” functions  
functions with character arrays, etc.

large overhead if added to all functions  
StackGuard paper: 5–20%?

similar space overheads

(for typical applications)  
could be much worse: tons of ‘risky’ function calls

# stack canaries pro/con

pro: no change to calling convention

pro: recompile only — no extra work

con: can't protect existing executable/library files (without recompile)

con: doesn't protect against many ways of exploiting buffer overflows

con: vulnerable to information leaks

# stack canaries pro/con

pro: no change to calling convention

pro: recompile only — no extra work

con: can't protect existing executable/library files (without recompile)

con: *doesn't protect against many ways of exploiting buffer overflows*

con: vulnerable to information leaks

# stack canaries pro/con

pro: no change to calling convention

pro: recompile only — no extra work

con: can't protect existing executable/library files (without recompile)

con: doesn't protect against many ways of exploiting buffer overflows

con: *vulnerable to information leaks*

# stack canary summary

stack canary — simplest of many *mitigations*

key idea: detect corruption of return address

assumption: if return address changed, so is adjacent token

assumption: attacker can't learn true value of token

often possible with memory bug

later: workarounds to break these assumptions

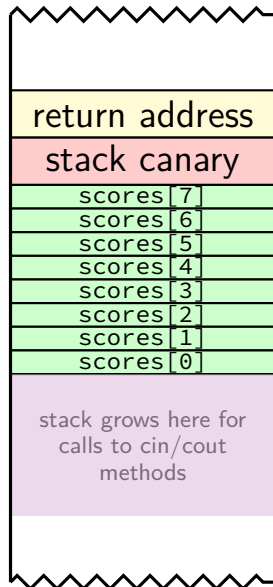
# stack canary hopes

*overwrite return address  $\implies$  overwrite canary*

canary is secret

# non-contiguous overwrites


```
void vulnerable() {  
    int scores[8]; bool done = false;  
    while (!done) {  
        cout << "Edit_which_score?(0_to_7)_";  
        int i;  
        cin >> i;  
        /* Oops!  
           sizeof(scores) is 8 * sizeof(int) */  
        if (i < 0 || i >= sizeof(scores))  
            continue;  
        cout << "Set_to_what_value?" << endl;  
        cin >> scores[i];  
        ...  
    }  
    ...  
}
```



# exercise: non-contiguous overwrites

```
void vulnerable() {  
    int scores[8]; bool done = false;  
    while (!done) {  
        cout << "Edit_which_score?(0_to_7)_";  
        int i;  
        cin >> i;  
        /* Oops!  
           sizeof(scores) is 4 * sizeof(int) */  
        if (i < 0 || i >= sizeof(scores))  
            continue;  
        cout << "Set_to_what_value?" << endl;  
        cin >> scores[i];  
    }  
}
```

exercise: to set return address to 0x123456789,  
set what scores to what values?





# stack canary hopes

overwrite return address  $\implies$  overwrite canary

*canary is secret*

# information disclosure (1a)

```
string command;  
void vulnerable() {  
    int value;  
    for (;;) {  
        cin >> command;  
        if (command == "set") {  
            cin >> value;  
        } else if (command == "get") {  
            cout << value << endl;  
        } else if ...  
    }  
}
```

“get” command: can read *uninitialized value*

example: when I compiled this, value was stored on the stack

# information disclosure (1b)

```
void vulnerable() {  
    int value;  
    ...  
    } else if (command == "get") {  
        cout << value << endl;  
    }  
    ...  
}
```

```
void leak() {  
    int secrets[] = {  
        12345678, 23456789, 34567890,  
        45678901, 56789012, 67890123,  
    };  
    cout << (void*) secrets << endl;  
    ...  
}
```

running this program  
(input in bold):

**get**  
67890123

## information disclosure (2)

```
void process() {  
    char buffer[8] = "\0\0\0\0\0\0\0\0";  
    char c = '_';  
    for (int i = 0; c != '\n' && i < 8; ++i) {  
        c = getchar();  
        buffer[i] = c;  
    }  
    printf("You_input%s\n", buffer);  
}
```

input aaaaaaaaa

output You input aaaaaaaaa(*whatever was on stack*)

## information disclosure (3)

```
struct foo {  
    char buffer[8];  
    long *numbers;  
};  
  
void process(struct foo* thing) {  
    ...  
    scanf("%s", thing->buffer);  
    ...  
    printf("first_number: %ld\n", thing->numbers[0]);  
}
```

input: aaaaaaaaaa(*address of canary*)

address on stack *or* where canary is read from in thread-local storage

## recall: ASLR

easier mentioned ASLR (address space layout randomization)

for stack: choose secret starting address for stack

info disclosure bugs are a big problem for this!

## exercise

```
struct point {      struct point *p;
    int x, y, z;; ...
};                  if (command == "get") {
                    /* 'p' could be uninitialized */
                    printf("%d,%d,%d\n", p->x, p->y, p->z);
                    } ...
                    ...
```

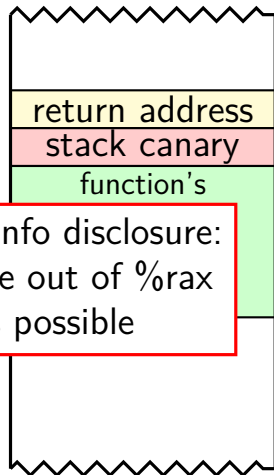
Which initial value for p (“left over” from prior use of register, etc.) would be most useful for figuring out the address of the stack pointer?

- A. p is an invalid pointer and accessing it will crash the program
- B. p points to space on the stack that is currently unallocated, but last contained an input buffer
- C. p points to a struct allocated on the heap

# compiler generated code

```
    pushq %rbx
    sub $0x20,%rsp
    /* copy value from thread-local storage */
    mov $0x28,%ebx
    mov %fs:(%rbx),%rax
    /* onto the stack */
    mov %rax,0x18(%rsp)
    /* clear register holding value */
    xor %eax,%eax
    ...
    ...
    /* copy value back from stack */
    mov 0x18(%rsp),%rax
    /* xor to compare */
    xor %fs:(%rbx),%rax
    /* if result non-zero, do not return */
    jne call_stack_chk_fail
    ret
call_stack_chk_fail:
```

trying to avoid info disclosure:  
get canary value out of %rax  
as soon as possible





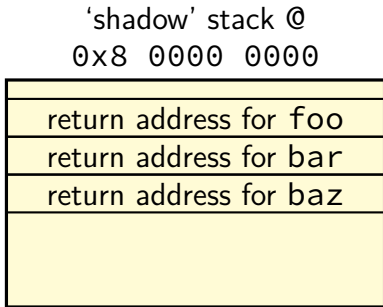
# intuition: shadow stacks

problem with stack: easy to leak address/values because used for lots of data

goal: keep sensitive data in **separate region**  
easier to keep address secret?

can use this for (stronger?) alternative to stack canaries

# shadow stacks



← shadow  
stack pointer

← stack pointer

# implementing shadow stacks

bigger changes to compiler than canaries

more overhead to call/return from function

most commonly: store return address twice

# shadow stacks on x86-64 (1)

idea 1: dedicate %r15 as shadow stack pointer,  
copy RA to shadow stack pointer in function prologue

function:

```
movq (%rsp), %rax      // RAX ← return address
addq $-8, %r15         // R15 ← R15 - 8
movq %rax, (%r15)      // M[R15] ← RAX
...
movq (%rsp), %rdx      // RDX ← return address
cmpq %rdx, (%r15)
jne CRASH_THE_PROGRAM  // if RDX != M[R15] goto
add $8, %r15           // R15 ← R15 - 8
ret
```

## shadow stacks on x86-64 (2)

idea 2: dedicate %r15 as shadow stack pointer,  
avoid normal call/return instruction

```
addq $-8, %r15
leaq after_call(%rip), %rax
movq %rax, (%r15)
jmp function
```

after\_call:

function:

```
...
addq $8, %r15           // R15 ← R15 + 8
jmp *-8(%r15)           // jmp M[R15-8]
```

# Android/AArch64 shadow stacks (1)

via <https://clang.llvm.org/docs/ShadowCallStack.html> (see also [https://security.googleblog.com/2019/10/protecting-against-code-reuse-in-linux\\_30.html](https://security.googleblog.com/2019/10/protecting-against-code-reuse-in-linux_30.html))

dedicate register x18 to shadow stack pointer

x30 = return address (after ARM's call instruction (bl))

ARM call instruction saves return address in register...

without

```
stp    x29, x30, [sp, #-16]!  
mov    x29, sp  
bl     bar  
add    w0, w0, #1  
ldp    x29, x30, [sp], #16  
ret
```

with shadow stack

```
str     x30, [x18], #8  
stp    x29, x30, [sp, #-16]!  
mov    x29, sp  
bl     bar  
add    w0, w0, #1  
ldp    x29, x30, [sp], #16  
ldr     x30, [x18, #-8]!  
ret
```

# Intel CET shadow stacks

future Intel processor extension adds shadow stacks  
“Control-flow Enforcement Technology”

new shadow stack pointer

CALL/RET: push/pop from BOTH stacks

shadow stack pages are marked as read-only in page table  
cannot be written through normal instructions  
extra bit identifying as shadow stack (not “normal” read-only page)

# preventing shadow stack writes?

ARM64 scheme: prevent writes if

- shadow stack pointer is never leaked (dedicated register)

- shadow stack random location can't be guessed (or queried otherwise)

Intel CET: prevent writes unless

- OS (privileged/kernel mode) instructions to setup shadow stack used

can we prevent writes without relying on avoiding info leaks...  
and without special hardware support?

- well, yes, but ...



# some early stack canary benchmarks

from Chiueh and Hsu, “RAD: A Compile-Time Solution to Buffer Overflow Attacks” (2001)

Program size	Program tested	User time	System time	Real time
11991 lines	Original ctags	0.57	0.05	0.62
	MineZone RAD-protected ctags	0.58	0.05	0.63
	Read-Only RAD-protected ctags	8.16	19.17	27.3

**Table 3 Macro-benchmark results of ctags**

Program size	Program tested	User time	System time	Real time
4500 lines	Original gcc	3.53	0.19	3.72
	Mine Zone RAD-protected gcc	4.67	0.2	4.87
	Read-Only RAD-protected gcc	20.46	50.43	70.8

**Table 4 Macro-benchmark results of gcc**

# automatic shadow stacks?

if we change how CALL/RET works...

...maybe we can add shadow stack support to existing programs?

either with hardware support, or  
in software with emulation techniques?

well, there's a problem...

# the problem in C++

```
void Foo() {  
    try {  
        ... Bar() ...  
    } except (std::runtime_error &error) {  
        ...  
    }  
}  
  
void Bar() {  
    ... Quux() ...  
}  
void Quux() {  
    ...  
    throw std::runtime_error("...");  
    ...  
}
```

# the problem in C

```
jmp_buf env;  
const char *error;  
void Foo() {  
    if (0 == setjmp(env)) {  
        Bar();  
    } else {  
        ...  
    }  
}
```

```
void Bar() {  
    ... Quux() ...  
}
```

```
void Quux() {  
    ...  
    error = "...";  
    longjmp(env, 1);  
    ...  
}
```

# dealing with non-local returns


exceptions and setjmp/longjmp deliberately skip return calls

one solution: “direct” shadow stack

fixed (possibly secret) offset from normal stack

shadow stack only stores return addresses

space in between return addresses left as nulls



# what do shadow stacks stop?


combined with a information leak that can dump arbitrary bytes of memory,

which of these exploits would shadow stacks stop...

- A. using format string exploit to point stack return address to the 'system' function
- B. using format string exploit to point VTable to the 'system' function
- C. using an unchecked string copy that goes over the end of a stack buffer into the return address and pointing the return address to the 'system' function
- D. using a buffer overflow that overwrites a saved stack pointer value to cause return to use a different address
- E. using pointer subterfuge to overwrite the GOT entry for 'printf' to point to the 'system' function

# recall(?): virtual memory

illusion of *dedicated memory*



# the mapping (set by OS)

program address range

0x0000 --- 0x0FFF

0x1000 --- 0x1FFF

...

0x40 0000 --- 0x40 0FFF

0x40 1000 --- 0x40 1FFF

0x40 2000 --- 0x40 2FFF

...

0x60 0000 --- 0x60 0FFF

0x60 1000 --- 0x60 1FFF

...

0x7FFF FF00 0000 — 0x7FFF FF00 0FFF

0x7FFF FF00 1000 — 0x7FFF FF00 1FFF

...

read?	write?
no	no
no	no

yes	no
yes	no
yes	no

yes	yes
yes	yes

yes	yes
yes	yes

real address
---
---

0x...
0x...
0x...

0x...
0x...

0x...
0x...



# Virtual Memory

modern *hardware-supported* memory protection mechanism

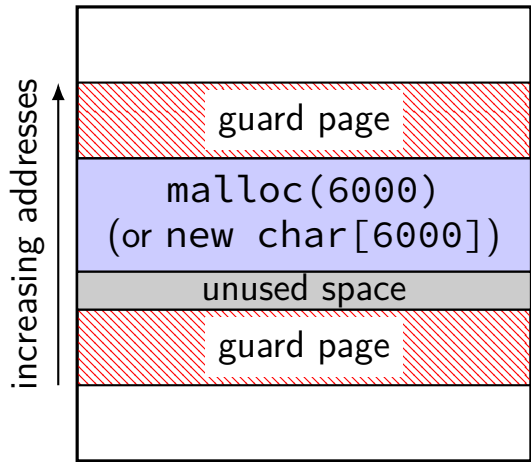
via *table*: OS decides *what memory program sees*  
whether it's read-only or not

granularity of *pages* — typically 4KB

not in table — segfault (OS gets control)

# malloc/new guard pages

the heap



# guard pages

deliberate holes

accessing — segfault

call to OS to allocate (not very fast)

likely to 'waste' memory

guard around object? minimum 4KB object

# guard pages for malloc/new

can implement malloc/new by placing guard pages around allocations

commonly done by real malloc/new's for *large allocations*

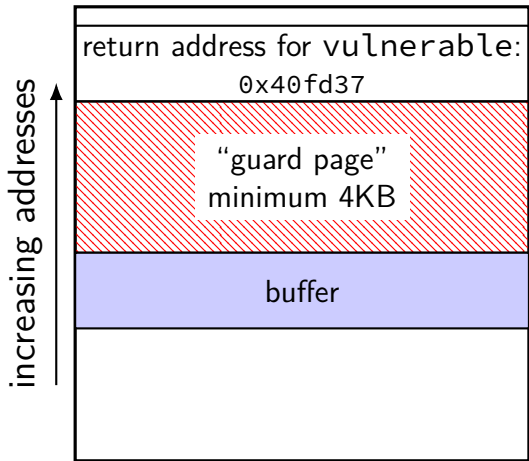
problem: minimum actual allocation 4KB

problem: substantially slower

example: “Electric Fence” allocator for Linux (early 1990s)

# stack canary alternative

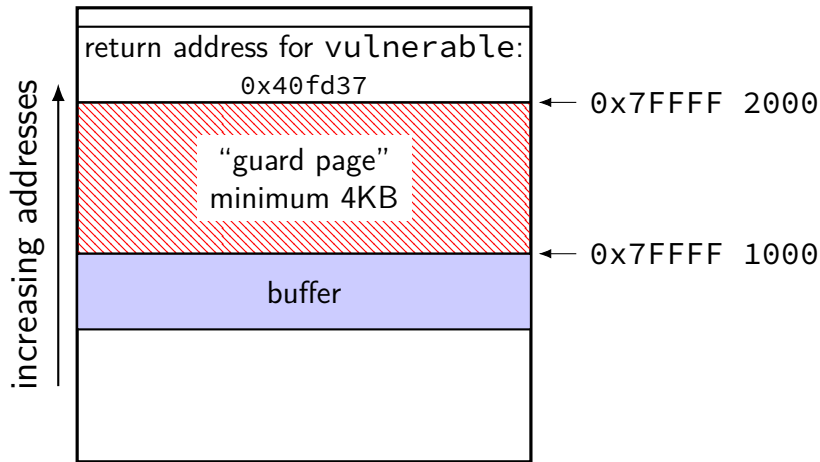
highest address (stack started here)



lowest address (stack grows here)

# stack canary alternative

highest address (stack started here)

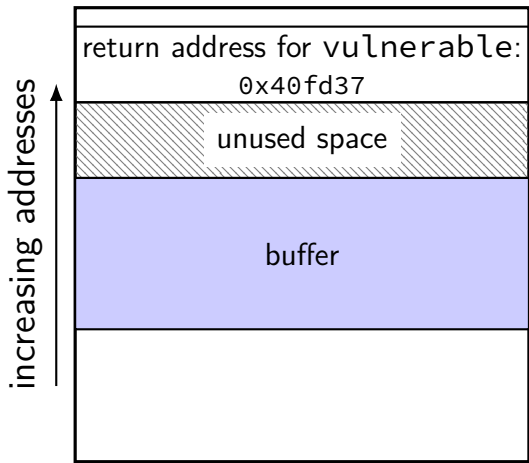


lowest address (stack grows here)

address	read	write
0x7FFFF2000-0x7FFFF2FFF	yes	yes
0x7FFFF1000-0x7FFFF1FFF	no	no
0x7FFFF0000-0x7FFFF0FFF	yes	yes

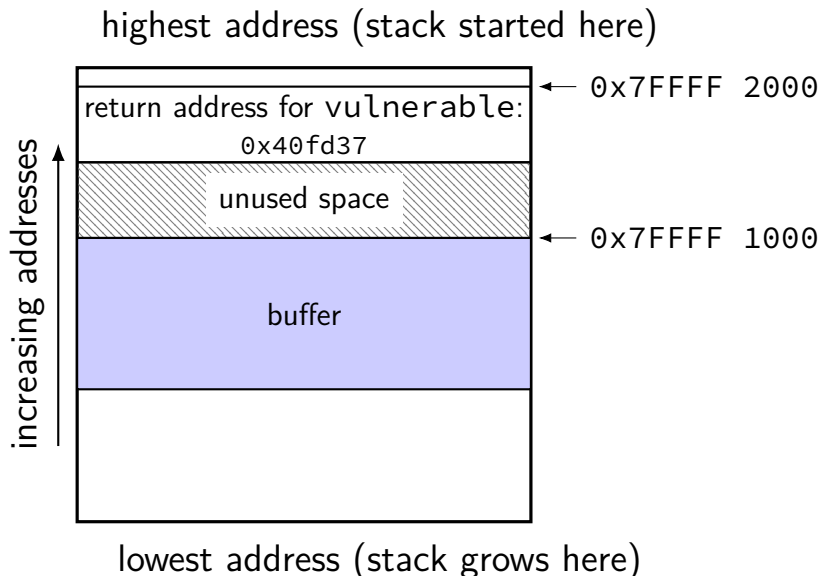
## stack canary alternative 2

highest address (stack started here)



lowest address (stack grows here)

# stack canary alternative 2



address	read	write
0x7FFFF2000-0x7FFFF2FFF	yes	yes
0x7FFFF1000-0x7FFFF1FFF	yes	<i>no</i>
0x7FFFF0000-0x7FFFF0FFF	yes	yes



## exercise: guard page overhead

suppose heap allocations are:

- 100 000 objects of 100 bytes

- 1 000 objects of 1000 bytes

- 100 objects of approx. 10000 bytes

total allocation of approx 12 000 KB

assuming 4KB pages, estimate space overhead of using guard pages:

- for objects larger than 4096 bytes (1 page)

- for objects larger than 200 bytes

- for all objects

# recall: function pointer targets

wanted to overwrite special pointer:

return addresses on stack

function pointers on in local variables

tables of function pointers used for inheritance

global offset table

last two: need to change infrequently

idea: make read-only

# RELRO

## **RE**Location **Read-Only**

Linux option: make dynamic linker structures read-only after startup

partial RELRO: everything but GOT pointers to library functions  
notably includes C++ virtual function tables

full RELRO: everything including those pointers  
requires disabling “lazy” linking  
(could do without disabling — but slower (how much?) startup)

appears as ELF program header entry

## a thought on permissions

if we can set memory non-writeable

how about non-executable?

we never want to execute things on the stack anyways, right?

# write XOR execute

many names:

- W<sup>X</sup> (write XOR execute)

- DEP (Data Execution Prevention)

- NX bit (No-eXecute) (hardware support)

- XD bit (eXecute Disable) (hardware support)

mark writeable memory as executable

how will users insert their machine code?

- can only code in application + libraries

- a problem, right?

# hardware support for write XOR execute

everywhere today

not historically common

early x86: execute implied by read

NX support added with x86-64 and around 2000 for x86-32

# deliberate use of writeable code

- “just-in-time” (JIT) compilers

  - fast virtual machine/language implementations

- some weird GCC features

- older “signals” on Linux

  - OS wrote machine code on stack for program to run

- couldn't even disable executable stacks without breaking applications

# why doesn't $W \text{ xor } X$ solve the problem?

$W \text{ xor } X$  is “almost free”, keeps attacker from writing code?

problem: useful machine code is in program already

just need to find writable function pointer

saw special case: arc injection

happened to find useful code in existing application/library

turns out: almost always useful code



# backup slides

## recall: relocation

```
.data
string: .asciz "Hello, World!"
.text
main:
    movq $string, %rdi /* NOT PC/RIP-relative mov */
```

generates: (objdump --disassemble --reloc)

```
0: 48 c7 c7 00 00 00 00    mov     $0x0,%rdi
3: R_X86_64_32S .data
```

*relocation record* says how to fix 0x0 in mov

3: location in machine code

R\_X86\_64\_32S: 32-bit signed integer

.data: address to insert